

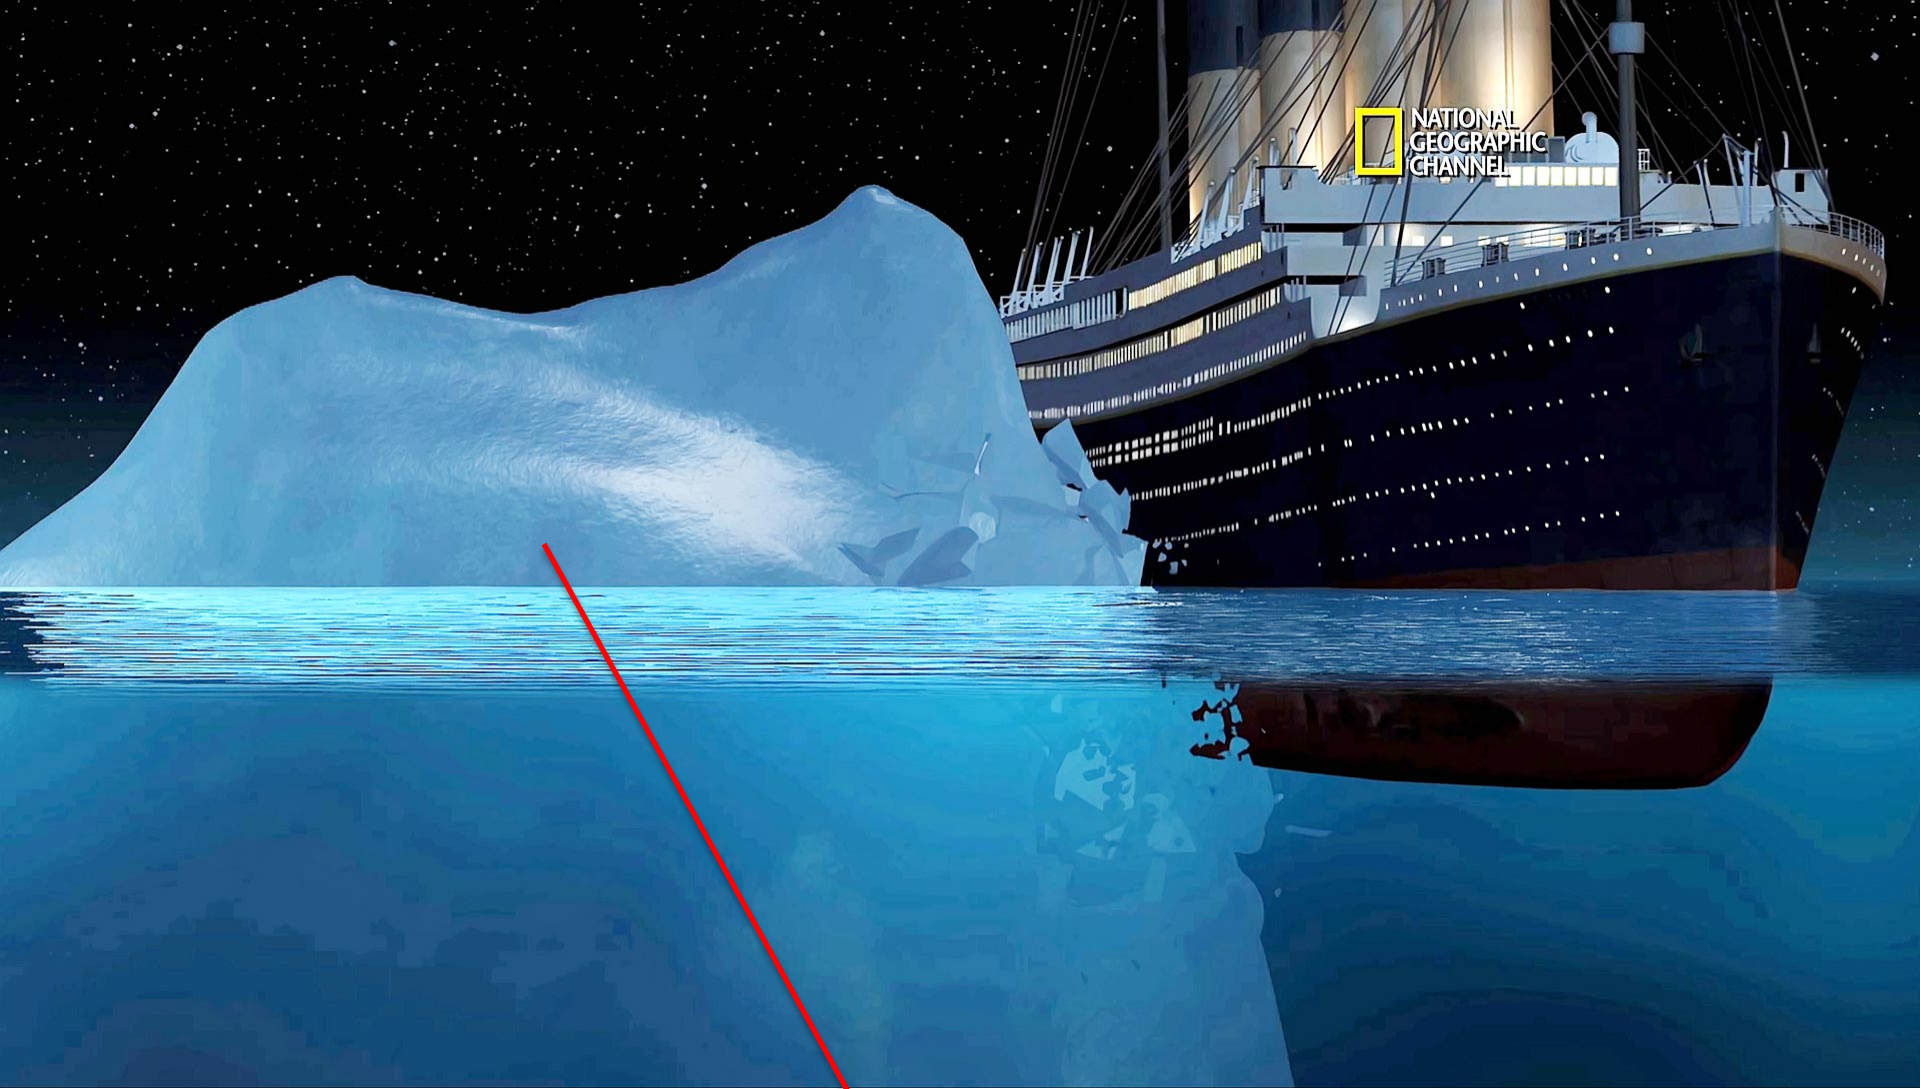
A romantic scene from the movie Titanic, featuring Jack and Rose about to kiss. The background is a blurred cityscape at night.

Chapter 5. Data Cube Technology

Meng Jiang

CS412 Summer 2017:

Introduction to Data Mining

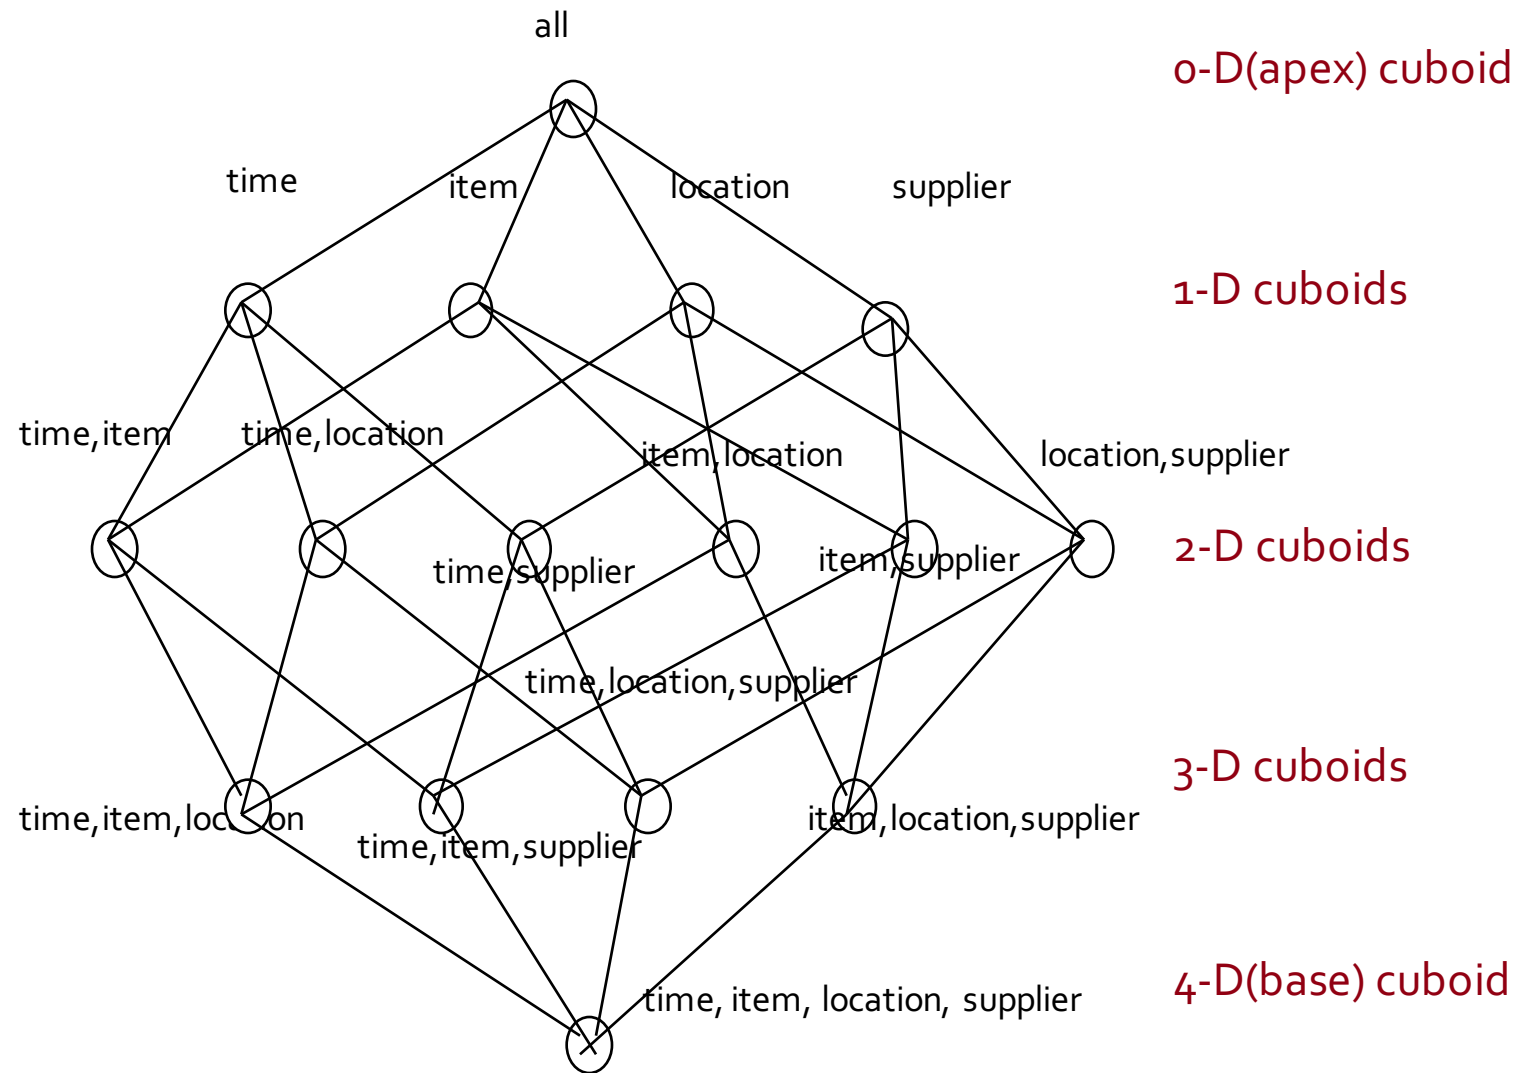


What the hell is that?!!

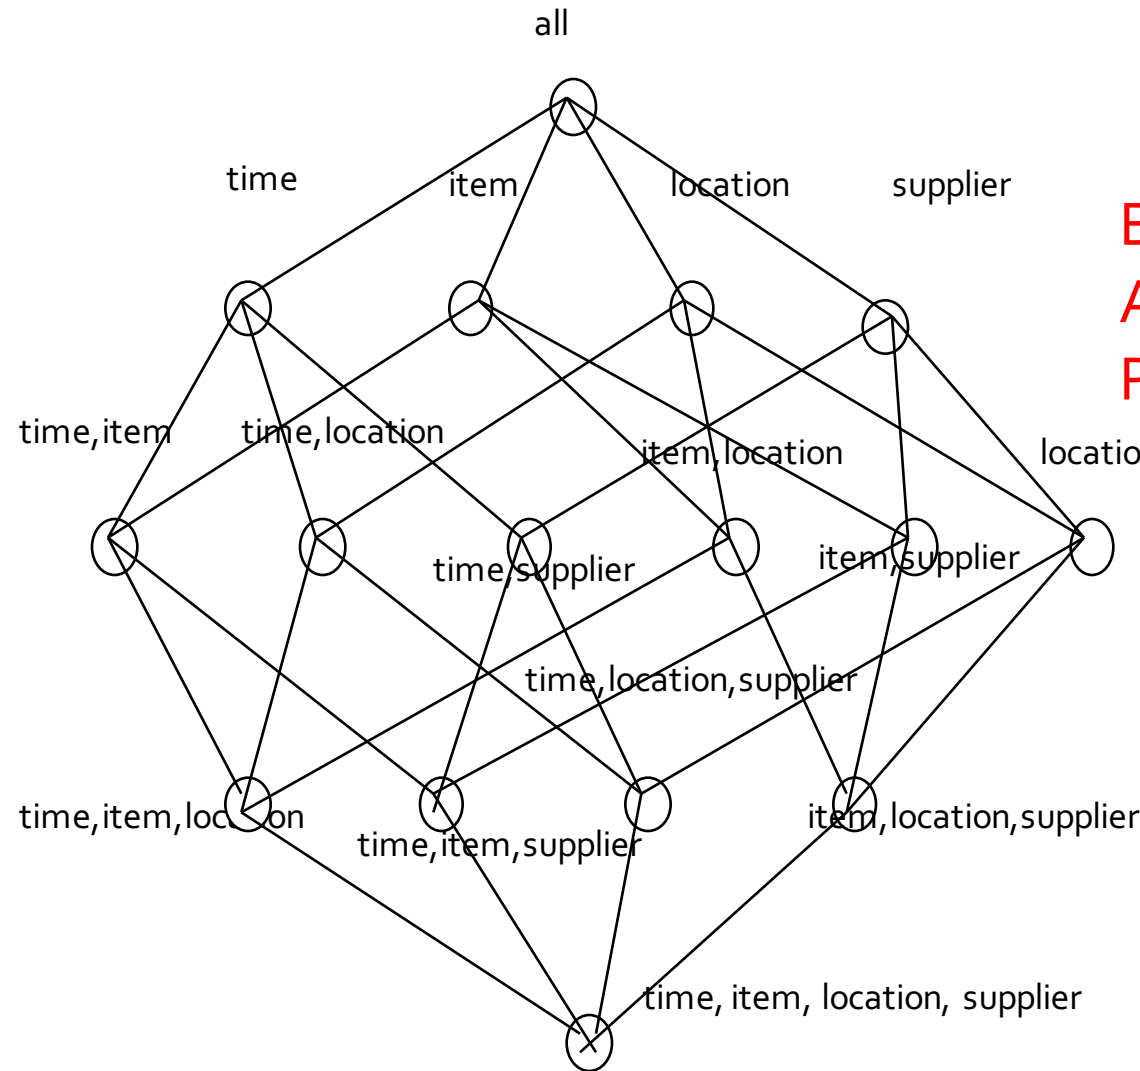
Data Cube Technology

- **Data Cube Computation: Basic Concepts**
- Data Cube Computation Methods
- Multidimensional Data Analysis in Cube Space

Cuboids in Data Cube



Cells in Cuboid



Base vs. aggregate cells
Ancestor vs. descendant cells
Parent vs. child cells

(*,*,*)

(*, milk, *, *)

(*, milk, Urbana, *)

(*, milk, Chicago, *)

(9/15, milk, Urbana, *)

(9/15, milk, Urbana, Dairy_land)

Cube Materialization: Full Cube vs. Iceberg Cube

- Full cube vs. iceberg cube
compute cube sales **iceberg** as
select month, city, customer group, count(*)
from salesInfo
cube by month, city, customer group
having **count(*) >= min support**
- Compute *only* the cells whose measure satisfies
the iceberg condition
- Only a small portion of cells may be “above the
water” in a **sparse cube**
- Ex.: Show only those cells whose count is no less
than 100



Why Iceberg Cube?

- Advantages of computing iceberg cubes
 - No need to save nor show those cells whose value is below the threshold (iceberg condition)
 - Efficient methods may even avoid computing the un-needed, intermediate cells
 - Avoid explosive growth
- Example: A cube with 100 dimensions
 - Suppose it contains only 2 base cells: $\{(a_1, a_2, a_3, \dots, a_{100}), (a_1, a_2, b_3, \dots, b_{100})\}$
 - How many aggregate cells if “having count ≥ 1 ”?

Why Iceberg Cube?

- Advantages of computing iceberg cubes
 - No need to save nor show those cells whose value is below the threshold (iceberg condition)
 - Efficient methods may even avoid computing the un-needed, intermediate cells
 - Avoid explosive growth
- Example: A cube with 100 dimensions
 - Suppose it contains only 2 base cells: $\{(a_1, a_2, a_3, \dots, a_{100}), (a_1, a_2, b_3, \dots, b_{100})\}$
 - How many aggregate cells if “having count ≥ 1 ”?
 - Answer: $(2^{101} - 2) - 4$

Why Iceberg Cube?

- Advantages of computing iceberg cubes
 - No need to save nor show those cells whose value is below the threshold (iceberg condition)
 - Efficient methods may even avoid computing the un-needed, intermediate cells
 - Avoid explosive growth
- Example: A cube with 100 dimensions
 - Suppose it contains only 2 base cells: $\{(a_1, a_2, a_3, \dots, a_{100}), (a_1, a_2, b_3, \dots, b_{100})\}$
 - How many aggregate cells if “having count ≥ 1 ”?
 - Answer: $(2^{101} - 2) - 4$
 - What about the iceberg cells, (i.e., with condition: “having count ≥ 2 ”)?

Why Iceberg Cube?

- Advantages of computing iceberg cubes
 - No need to save nor show those cells whose value is below the threshold (iceberg condition)
 - Efficient methods may even avoid computing the un-needed, intermediate cells
 - Avoid explosive growth
- Example: A cube with 100 dimensions
 - Suppose it contains only 2 base cells: $\{(a_1, a_2, a_3, \dots, a_{100}), (a_1, a_2, b_3, \dots, b_{100})\}$
 - How many aggregate cells if “having count ≥ 1 ”?
 - Answer: $(2^{101} - 2) - 4$
 - What about the iceberg cells, (i.e., with condition: “having count ≥ 2 ”)?
 - Answer: 4

Is Iceberg Cube Good Enough?

Closed Cube & Cube Shell

- Let cube P have only 2 base cells: $\{(a_1, a_2, a_3 \dots, a_{100}):10, (a_1, a_2, b_3, \dots, b_{100}):10\}$
 - How many cells will the iceberg cube contain if “having count(*) ≥ 10 ”?
 - Answer: $2^{101} - 4$ (still too big!)
- **Close cube:**
 - A cell c is **closed** if there exists no cell d , such that d is a descendant of c , and d has the same measure value as c
 - Ex. The same cube P has only 3 closed cells:
 - $\{(a_1, a_2, *, \dots, *): 20, (a_1, a_2, a_3 \dots, a_{100}): 10, (a_1, a_2, b_3, \dots, b_{100}): 10\}$
 - A **closed cube** is a cube consisting of only closed cells
- **Cube Shell:** The **cuboids** involving only a **small # of dimensions**, e.g., 2
 - Idea: Only compute cube shells, other dimension combinations can be computed on the fly

Data Cube Technology

- Data Cube Computation: Basic Concepts
- **Data Cube Computation Methods**
- Multidimensional Data Analysis in Cube Space

Roadmap for Efficient Computation

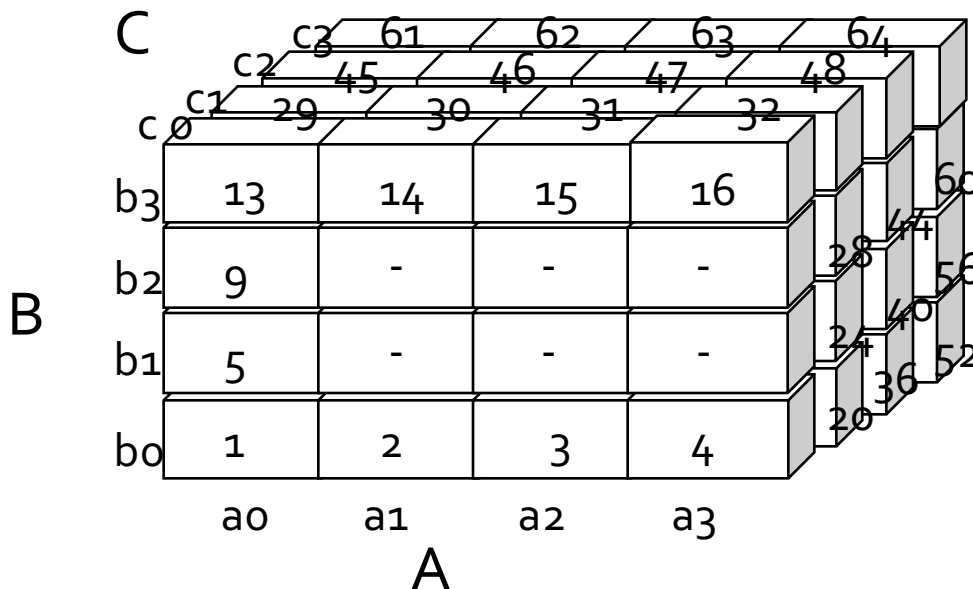
- **General computation heuristics** (Agarwal et al.'96)
- Computing full/iceberg cubes: 3 methodologies
 - Bottom-Up:
 - **Multi-way array aggregation** (Zhao, Deshpande & Naughton, SIGMOD'97)
 - Top-down:
 - **BUC** (Beyer & Ramakrishnan, SIGMOD'99)
 - Integrating Top-Down and Bottom-Up:
 - **Star-cubing algorithm** (Xin, Han, Li & Wah: VLDB'03)
- High-dimensional OLAP:
 - **A shell-fragment approach** (Li, et al. VLDB'04)
- Computing alternative kinds of cubes:
 - Partial cube, closed cube, approximate cube,

Efficient Data Cube Computation: General Heuristics

- Sorting, hashing, and grouping operations are applied to the dimension attributes in order to reorder and cluster related tuples
- Aggregates may be computed from previously computed aggregates, rather than from the base fact table
 - Smallest-child: computing a cuboid **from the smallest**, previously computed cuboid
 - Cache-results: caching results of a cuboid from which other cuboids are computed to **reduce disk I/Os**
 - Amortize-scans: computing **as many as possible** cuboids at the same time to **amortize disk reads**
 - Share-sorts: **sharing sorting costs** across multiple cuboids when **sort-based method** is used
 - Share-partitions: **sharing the partitioning cost** across multiple cuboids when **hash-based algorithms** are used

Cube Computation: Multi-Way Array Aggregation (MOLAP)

- Bottom-up: Partition a huge *sparse* array into *chunks* (a small subcube which fits in memory) and aggregation.
- Data addressing: Compressed *sparse array addressing* (chunk_id, offset)
- Compute aggregates in “multiway” by visiting cube cells in the order which **minimizes** the # of times to visit each cell, and **reduces** memory access and storage cost




What is the best traversing order to do multi-way aggregation?

Multi-way Array Aggregation (3-D to 2-D)

- How to minimize the memory requirement and reduced I/Os?
 - Keep the **smallest** plane in **main memory**
 - Fetch and compute **only one chunk** at a time for the **largest** plane
 - The planes should be **sorted** and computed according to their **size** in ascending

One **chunk** of BC plane:
 $(400/4) * (4000/4)$
 = 100,000

Min memory size:
 $156,000 < \dots$
 $< 64,000,000$




Entire AB plane:
 $40 * 400$
 = 16,000

A: 40 (department),
 B: 400 (item),
 C: 4000 (time)

One **column** of AC plane:
 $40 * (4000/4) = 40,000$

Multi-Way Array Aggregation

- Array-based “**bottom-up**” algorithm (from ABC to AB,...)
- Using multi-dimensional **chunks**
- Simultaneous aggregation on multiple dimensions
- Intermediate aggregate values are re-used for computing ancestor cuboids
- Cannot do *Apriori* pruning: No iceberg optimization
- Comments on the method
 - Efficient for computing the full cube for **a small number of dimensions**
 - If there are a large number of dimensions, “top-down” computation and iceberg cube computation methods (e.g., BUC) should be used



Cube Computation: Computing in Reverse Order

- BUC (Beyer & Ramakrishnan, SIGMOD'99)

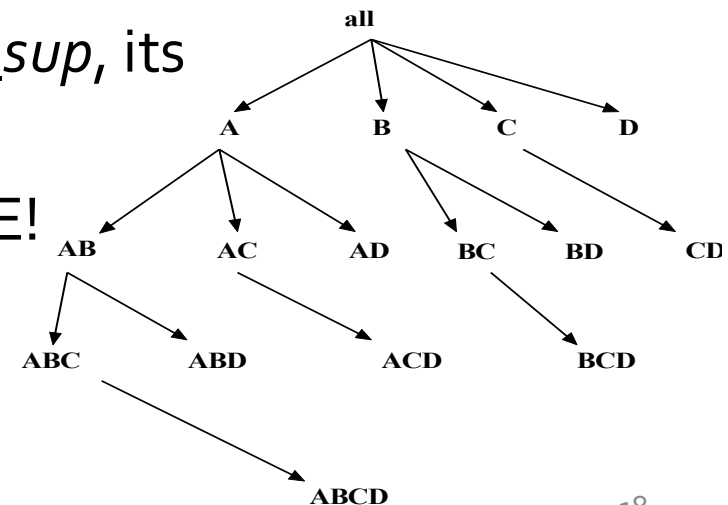
BUC: acronym of Bottom-Up (cube) Computation

(Note: It is “**top-down**” in our view since we put Apex cuboid on the top!)

- Divides dimensions into partitions and facilitates **iceberg pruning (it works now!)**


- If a partition does not satisfy *min_sup*, its **descendants** can be pruned
- If *min_sup* = 1 \Rightarrow compute full CUBE!

- No simultaneous aggregation



BUC: Partitioning and Aggregating


- Usually, entire data set cannot fit in main memory
- Sort *distinct* values
 - partition into blocks that fit
- Continue processing
- Optimizations
 - Partitioning
 - External Sorting, Hashing, Counting Sort
 - Ordering dimensions to encourage pruning
 - Cardinality, Skew, Correlation
 - Collapsing duplicates
 - Cannot do holistic aggregates anymore!



High-Dimensional OLAP?

— The Curse of Dimensionality

- High-DOLAP: Needed in many applications
 - Science and engineering analysis
 - Bio-data analysis: thousands of genes
 - Statistical surveys: hundreds of variables
- None of the previous cubing method can handle high dimensionality!
 - Iceberg cube and compressed cubes: only delay the inevitable explosion
 - Full materialization: still significant overhead in accessing results on disk
- **A shell-fragment approach:** X. Li, J. Han, and H. Gonzalez, High-Dimensional OLAP: A Minimal Cubing Approach, VLDB'04



A curse of dimensionality: A database of 600,000 tuples. Each dimension has cardinality of 100 and *zipf* of 2.

Fast High-Dimensional OLAP with Minimal Cubing

- Observation: OLAP occurs only on a small subset of dimensions at a time
- Semi-Online Computational Model
 - Partition the set of dimensions into **shell fragments**
 - Compute data cubes for each shell fragment while retaining **inverted indices** or **value-list indices**
 - Given the pre-computed **fragment cubes**, dynamically compute cube cells of the high-dimensional data cube *online*
- Major idea: Tradeoff between the amount of **pre-computation** and the speed of **online computation**
 - Reducing computing high-dimensional cube into **pre-computing** a set of **lower dimensional** cubes
 - **Online re-construction** of original high-dimensional space
 - Lossless reduction

Computing a 5-D Cube with 2-Shell Fragments

- Example: Let the cube aggregation function be **count**

<i>tid</i>	A	B	C	D	E
1	a1	b1	c1	d1	e1
2	a1	b2	c1	d2	e1
3	a1	b2	c1	d1	e2
4	a2	b1	c1	d1	e2
5	a2	b1	c1	d1	e3

Attribute Value	TID List	List Size
a1	1 2 3	3
a2	4 5	2
b1	1 4 5	3
b2	2 3	2
c1	1 2 3 4 5	5
d1	1 3 4 5	4
d2	2	1
e1	1 2	2
e2	3 4	2
e3	5	1

- Divide the 5-D table into 2 shell fragments:
 - (A, B, C) and (D, E)
- Build traditional invert index (TID) or RID list

Shell Fragment Cubes: Ideas

- Generalize the 1-D inverted indices to **multi-dimensional inverted indices** in the data cube sense
- Compute **all cuboids for data cubes ABC and DE** while retaining the inverted indices

Attribute Value	TID List	List Size
a1	1 2 3	3
a2	4 5	2
b1	1 4 5	3
b2	2 3	2
c1	1 2 3 4 5	5
d1	1 3 4 5	4
d2	2	1
e1	1 2	2
e2	3 4	2
e3	5	1

Cell	Intersection	TID List	List Size
a1 b1	$1\ 2\ 3 \cap 1\ 4\ 5$	1	1
a1 b2	$1\ 2\ 3 \cap 2\ 3$	2 3	2
a2 b1	$4\ 5 \cap 1\ 4\ 5$	4 5	2
a2 b2	$4\ 5 \cap 2\ 3$	ϕ	0

Cell	Intersection	TID List	List Size
d1 e1	?	?	?
d1 e2	?	?	?
d1 e3	?	?	?
d2 e1	?	?	?

Answer: http://hanj.cs.illinois.edu/pdf/vldbo4_hdolap.pdf

Shell Fragment Cubes: Size and Design

- Given a database of T tuples, **D dimensions**, and F shell fragment size, the fragment cubes' space requirement is:
 - For **$F < 5$** , the growth is sub-linear $O\left(T \left\lceil \frac{D}{F} \right\rceil (2^F - 1)\right)$
- Shell fragments do not have to be disjoint
- Fragment groupings can be arbitrary to allow for maximum online performance
 - Known common combinations (e.g., <city, state>) should be grouped together
- Shell fragment sizes can be adjusted for optimal balance between offline and online computation

Data Cube Technology

- Data Cube Computation: Basic Concepts
- Data Cube Computation Methods
- **Multidimensional Data Analysis in Cube Space**

Data Mining in Cube Space

- Data cube greatly increases the analysis bandwidth
- Four ways to **interact OLAP-styled analysis and data mining**
 - Using cube space to define **data space** for mining
 - Using OLAP queries to generate **features and targets** for mining, e.g., multi-feature cube
 - Using data-mining models as **building blocks** in a multi-step mining process, e.g., prediction cube
 - Using data-cube computation techniques to **speed up** repeated model construction
 - Cube-space data mining may require building a model for each candidate data space
 - Sharing computation across model-construction for different candidates may lead to efficient mining

Complex Aggregation at Multiple Granularities: Multi-Feature Cubes

- Multi-feature cubes (Ross, et al. 1998): Compute complex queries involving multiple dependent aggregates at multiple granularities
- Ex. Grouping by all subsets of {item, region, month}, find the **maximum price** in 2010 for each group, and the **total sales** among all maximum price tuples

```
select item, region, month, max(price), sum(R.sales)
from purchases
where year = 2010
cube by item, region, month: R
such that R.price = max(price)
```

- Continuing the last example, **among the max price tuples**, find the **min and max shelf life**, and find the fraction of the total sales due to tuple that have **min shelf life** within the set of all max price tuples

Discovery-Driven Exploration of Data Cubes

- Discovery-driven exploration of huge cube space (Sarawagi, et al.'98)
 - Effective navigation of large OLAP data cubes
 - Pre-compute measures indicating exceptions, guide user in the data analysis, at all levels of aggregation
 - Exception: significantly different from the value anticipated, based on a statistical model
 - Visual cues such as background color are used to reflect the degree of exception of each cell
- Kinds of exceptions
 - **SelfExp**: represents the surprise value of the cell relative to other cells at same level of aggregation
 - **InExp**: represents the degree of surprise somewhere beneath this cell if we drill down from the cell
 - **PathExp**: represents the degree of surprise for each drill-down path from the cell

Exceptions: SelfExp, InExp, PathExp

Product	(All)
Region	(All)

SelfExp

Avg.Sales	Month												
Region	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec	
Total		2%	0%	2%	2%	4%	3%	0%	-8%	0%	-3%	4%	



InExp

Region	(All)
--------	-------

Avg.Sales	Month												
Product	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec	
Birch-B		10%	-7%	3%	-4%	15%	-12%	-3%	1%	42%	-14%	-10%	
Chery-S		1%	1%	4%	3%	5%	5%	-9%	-12%	1%	-5%	5%	
Cola		-1%	2%	3%	4%	9%	4%	1%	-11%	-8%	-2%	7%	
Cream-S		3%	1%	6%	3%	3%	8%	-3%	-12%	-2%	1%	10%	
Diet-B		1%	1%	-1%	2%	1%	2%	0%	-6%	-1%	-4%	2%	
Diet-C		3%	2%	5%	2%	4%	7%	-7%	-12%	-2%	-2%	8%	
Diet-S		2%	-1%	0%	0%	4%	2%	4%	-9%	5%	-3%	0%	
Grape-S		1%	1%	0%	4%	5%	1%	3%	-9%	-1%	-8%	4%	
Jolt-C		-1%	-4%	2%	2%	0%	-4%	2%	6%	-2%	0%	0%	
Kiwi-S		2%	1%	4%	1%	-1%	3%	-1%	-4%	4%	0%	1%	
Old-B		4%	-1%	0%	1%	5%	2%	7%	-10%	3%	-3%	1%	
Orang-S		1%	1%	3%	4%	2%	1%	-1%	-1%	-6%	-4%	9%	
Sasprla		-1%	2%	1%	3%	-3%	5%	-10%	-2%	-1%	1%	5%	

Product	Diet-S
---------	--------

Avg.Sales	Month												
Region	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec	
C		0%	-2%	0%	1%	4%	1%	5%	-6%	2%	-2%	-2%	
E		0%	2%	-8%	7%	0%	5%	-40%	10%	-33%	2%	8%	
S		0%	-1%	3%	-2%	2%	-2%	19%	-1%	12%	-1%	0%	
W		5%	1%	0%	-2%	6%	6%	2%	-17%	9%	-7%	2%	

Market	(All)
Product	Cola

Avg.Sales	Month												
Region	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec	
C		3%	1%	4%	1%	4%	10%	-11%	-14%	-3%	5%	11%	
E		-3%	3%	4%	4%	13%	2%	0%	-10%	-13%	-3%	8%	
S		2%	-1%	1%	9%	6%	3%	21%	-15%	1%	-5%	4%	
W		-2%	2%	2%	4%	12%	1%	1%	-9%	-11%	-4%	6%	

Summary

- Data Cube Computation: Preliminary Concepts
- Data Cube Computation Methods
 - Multi-Way Array Aggregation
 - BUC
 - High-Dimensional OLAP with Shell-Fragments
- Multidimensional Data Analysis in Cube Space
 - Multi-feature Cubes
 - Discovery-Driven Exploration of Data Cubes

References

- S. Agarwal, R. Agrawal, P. M. Deshpande, A. Gupta, J. F. Naughton, R. Ramakrishnan, and S. Sarawagi. On the computation of multidimensional aggregates. VLDB'96
- K. Beyer and R. Ramakrishnan. Bottom-Up Computation of Sparse and Iceberg CUBEs.. SIGMOD'99
- J. Han, J. Pei, G. Dong, K. Wang. Efficient Computation of Iceberg Cubes With Complex Measures. SIGMOD'01
- L. V. S. Lakshmanan, J. Pei, and J. Han, Quotient Cube: How to Summarize the Semantics of a Data Cube, VLDB'02
- X. Li, J. Han, and H. Gonzalez, High-Dimensional OLAP: A Minimal Cubing Approach, VLDB'04
- X. Li, J. Han, Z. Yin, J.-G. Lee, Y. Sun, "Sampling Cube: A Framework for Statistical OLAP over Sampling Data", SIGMOD'08
- K. Ross and D. Srivastava. Fast computation of sparse datacubes. VLDB'97
- D. Xin, J. Han, X. Li, B. W. Wah, Star-Cubing: Computing Iceberg Cubes by Top-Down and Bottom-Up Integration, VLDB'03
- Y. Zhao, P. M. Deshpande, and J. F. Naughton. An array-based algorithm for simultaneous multidimensional aggregates. SIGMOD'97
- D. Burdick, P. Deshpande, T. S. Jayram, R. Ramakrishnan, and S. Vaithyanathan. OLAP over uncertain and imprecise data. VLDB'05

References (cont.)

- R. Agrawal, A. Gupta, and S. Sarawagi. Modeling multidimensional databases. ICDE'97
- B.-C. Chen, L. Chen, Y. Lin, and R. Ramakrishnan. Prediction cubes. VLDB'05
- B.-C. Chen, R. Ramakrishnan, J.W. Shavlik, and P. Tamma. Bellwether analysis: Predicting global aggregates from local regions. VLDB'06
- Y. Chen, G. Dong, J. Han, B. W. Wah, and J. Wang, Multi-Dimensional Regression Analysis of Time-Series Data Streams, VLDB'02
- R. Fagin, R. V. Guha, R. Kumar, J. Novak, D. Sivakumar, and A. Tomkins. Multi-structural databases. PODS'05
- J. Han. Towards on-line analytical mining in large databases. SIGMOD Record, 27:97–107, 1998
- T. Imielinski, L. Khachiyan, and A. Abdulghani. Cubegrades: Generalizing association rules. Data Mining & Knowledge Discovery, 6:219–258, 2002.
- R. Ramakrishnan and B.-C. Chen. Exploratory mining in cube space. Data Mining and Knowledge Discovery, 15:29–54, 2007.
- K. A. Ross, D. Srivastava, and D. Chatziantoniou. Complex aggregation at multiple granularities. EDBT'98
- S. Sarawagi, R. Agrawal, and N. Megiddo. Discovery-driven exploration of OLAP data cubes. EDBT'98
- G. Sathe and S. Sarawagi. Intelligent Rollups in Multidimensional OLAP Data. VLDB'01