

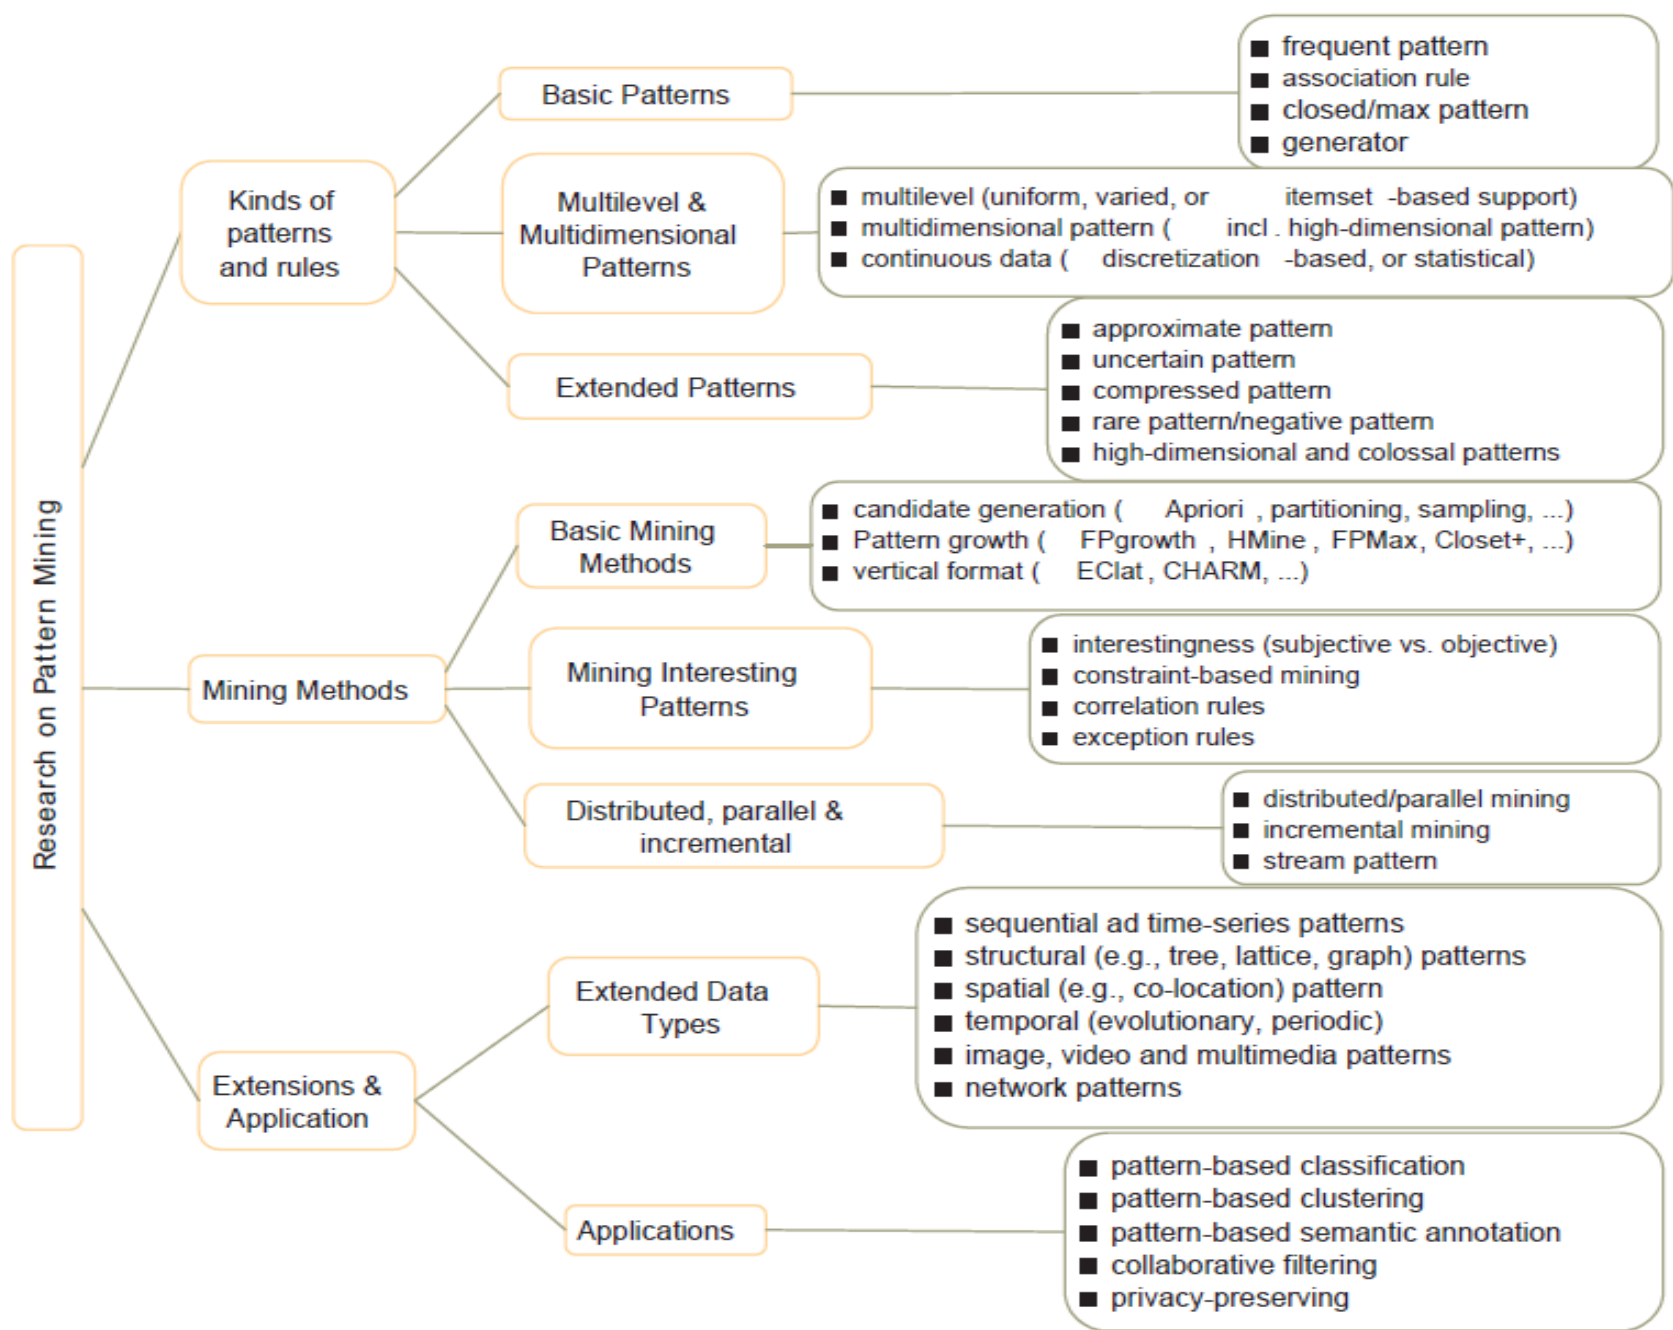
Chapter 7. Advanced Frequent Pattern Mining

Meng Jiang

CS412 Summer 2017:

Introduction to Data Mining

Research on Pattern Mining: A Road Map



Advanced Frequent Pattern Mining

- **Mining Diverse Patterns**
- Constraint-Based Frequent Pattern Mining
- Mining High-Dimensional Data and Colossal Patterns
- Sequential Pattern Mining
- Graph Pattern Mining

Mining Diverse Patterns

- Mining Multiple-Level Associations
- Mining Multi-Dimensional Associations
- Mining Quantitative Associations
- Mining Negative Correlations
- Mining Compressed and Redundancy-Aware Patterns

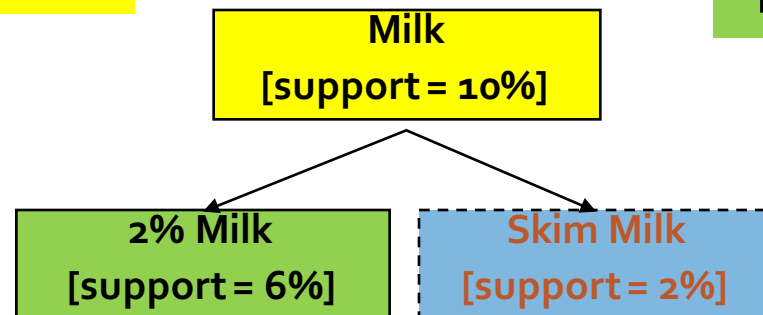
Mining Multiple-Level Frequent Patterns

- Items often form hierarchies
 - Ex.: Dairyland 2% milk; Wonder wheat bread
- How to set min-support thresholds?
 - Uniform min-support across multiple levels (reasonable?)
 - Level-reduced min-support: Items at the lower level are expected to have lower support
- Efficient mining: Shared multi-level mining
 - Use the lowest min-support to pass down the set of candidates

Uniform support

Level 1
min_sup = 5%

Level 2
min_sup = 5%



Reduced support

Level 1
min_sup = 5%

Level 2
min_sup = 1%

Redundancy Filtering at Mining Multi-Level Associations

- Multi-level association mining may generate many redundant rules
- Redundancy filtering: Some rules may be redundant due to “ancestor” relationships between items
 - (Suppose the 2% milk sold is about $\frac{1}{4}$ of milk sold in gallons)
 - milk \Rightarrow wheat bread [support = 8%, confidence = 70%] (1)
 - 2% milk \Rightarrow wheat bread [support = 2%, confidence = 72%] (2)
- A rule is *redundant* if its support is close to the “expected” value, according to its “ancestor” rule, and it has a similar confidence as its “ancestor”
 - Rule (1) is an ancestor of rule (2), which one to prune?

Customized Min-Supports for Different Kinds of Items

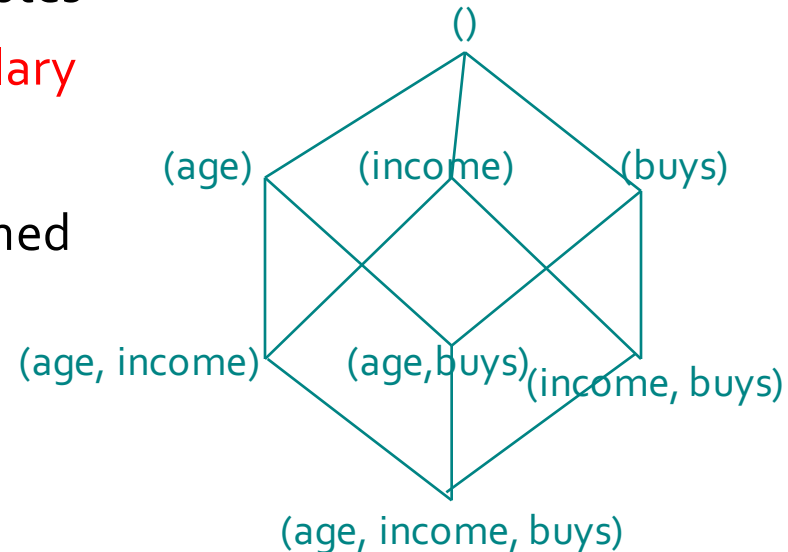
- We have used the same min-support threshold for all the items or item sets to be mined in each association mining
- In reality, some items (e.g., diamond, watch, ...) are valuable but less frequent
- It is necessary to have customized min-support settings for different kinds of items
- One Method: Use **group-based “individualized” min-support**
 - E.g., {diamond, watch}: 0.05%; {bread, milk}: 5%; ...
 - How to mine such rules efficiently?
 - Existing scalable mining algorithms can be easily extended to cover such cases

Mining Multi-Dimensional Associations

- Single-dimensional rules (e.g., items are all in “product” dimension)
 - $\text{buys}(X, \text{“milk”}) \Rightarrow \text{buys}(X, \text{“bread”})$
- Multi-dimensional rules (i.e., items in ≥ 2 dimensions or predicates)
 - Inter-dimension association rules (*no repeated predicates*)
 - $\text{age}(X, \text{“18-25”}) \wedge \text{occupation}(X, \text{“student”}) \Rightarrow \text{buys}(X, \text{“coke”})$
 - Hybrid-dimension association rules (*repeated predicates*)
 - $\text{age}(X, \text{“18-25”}) \wedge \text{buys}(X, \text{“popcorn”}) \Rightarrow \text{buys}(X, \text{“coke”})$
- Attributes can be categorical or numerical
 - Categorical Attributes (e.g., *profession, product*: no ordering among values): Data cube for inter-dimension association
 - Quantitative Attributes: Numeric, implicit ordering among values—discretization, clustering, and gradient approaches

Mining Quantitative Associations

- Mining associations with numerical attributes
 - Ex.: Numerical attributes: **age** and **salary**
- Methods
 - Static discretization based on predefined concept hierarchies
 - Data cube-based aggregation
 - Dynamic discretization based on data distribution
 - Clustering: Distance-based association
 - First one-dimensional clustering, then association
 - Deviation analysis:
 - Gender = female \Rightarrow Wage: mean=\$7/hr (overall mean = \$9)



Mining Extraordinary Phenomena in Quantitative Association Mining

- Mining extraordinary (i.e., interesting) phenomena
 - Ex.: $\text{Gender} = \text{female} \Rightarrow \text{Wage: mean} = \$7/\text{hr}$ (overall mean = \$9)
 - LHS: a subset of the population
 - RHS: an extraordinary behavior of this subset
- The rule is accepted only if a statistical test (e.g., Z-test) confirms the inference with high confidence
- Subrule: Highlights the extraordinary behavior of a subset of the population of the super rule
 - Ex.: $(\text{Gender} = \text{female}) \wedge (\text{South} = \text{yes}) \Rightarrow \text{mean wage} = \$6.3/\text{hr}$
- Rule condition can be categorical or numerical (quantitative rules)
 - Ex.: $\text{Education in } [14-18] \text{ (yrs)} \Rightarrow \text{mean wage} = \$11.64/\text{hr}$
- Efficient methods have been developed for mining such extraordinary rules (e.g., Aumann and Lindell@KDD'99)

Rare Patterns vs. Negative Patterns

- Rare patterns
 - Very low support but interesting (e.g., buying Rolex watches)
 - How to mine them? Setting individualized, group-based min-support thresholds for different groups of items
- Negative patterns
 - Negatively correlated: Unlikely to happen together
 - Ex.: Since it is unlikely that the same customer buys both a **Ford Expedition** (an SUV car) and a **Ford Fusion** (a hybrid car), buying a **Ford Expedition** and buying a **Ford Fusion** are likely negatively correlated patterns
 - How to define negative patterns?

Defining Negative Correlated Patterns

- A support-based definition
 - If itemsets A and B are both frequent but rarely occur together, i.e., $\text{sup}(A \cup B) \ll \text{sup}(A) \times \text{sup}(B)$
 - Then A and B are negatively correlated
- Is this a good definition for large transaction datasets?
- Ex.: Suppose a store sold two needle packages A and B 100 times each, but only one transaction contained both A and B
 - When there are in total 200 transactions, we have
 - $s(A \cup B) = 0.005, s(A) \times s(B) = 0.25, s(A \cup B) \ll s(A) \times s(B)$
 - But when there are 10^5 transactions, we have
 - $s(A \cup B) = 1/10^5, s(A) \times s(B) = 1/10^3 \times 1/10^3, s(A \cup B) > s(A) \times s(B)$
 - What is the problem? — Null transactions: The support-based definition is not null-invariant!

Does this remind you the definition of *lift*?

Defining Negative Correlation: Need Null-Invariance in Definition

- A good definition on negative correlation should take care of the null-invariance problem
 - Whether two itemsets A and B are negatively correlated should not be influenced by the number of null-transactions
- A Kulczynski measure-based definition
 - If itemsets A and B are frequent but $(P(A|B) + P(B|A))/2 < \epsilon$, where ϵ is a negative pattern threshold, then A and B are negatively correlated
- For the same needle package problem:
 - No matter there are in total 200 or 10^5 transactions
 - If $\epsilon = 0.01$, we have $(P(A|B) + P(B|A))/2 = (0.01 + 0.01)/2 < \epsilon$

Mining Compressed Patterns

- Why mining compressed patterns?
 - Too many scattered patterns but not so meaningful

- Pattern distance measure

$$Dist(P_1, P_2) = 1 - \frac{|T(P_1) \cap T(P_2)|}{|T(P_1) \cup T(P_2)|}$$


- δ -clustering: For each pattern P, find all patterns which can be expressed by P and whose distance to P is within δ (δ -cover)
- All patterns in the cluster can be represented by P
- Method for efficient, direct mining of compressed frequent patterns (e.g., D. Xin, J. Han, X. Yan, H. Cheng, "On Compressing Frequent Patterns", Knowledge and Data Engineering, 60:5-29, 2007)

Pat-ID	Item-Sets	Support
P ₁	{38,16,18,12}	205227
P ₂	{38,16,18,12,17}	205211
P ₃	{39,38,16,18,12,17}	101758
P ₄	{39,16,18,12,17}	161563
P ₅	{39,16,18,12}	161576


- Closed patterns
 - P₁, P₂, P₃, P₄, P₅
 - Emphasizes too much on support
 - There is no compression
- Max-patterns
 - P₃: information loss
- Desired output (a good balance):
 - P₂, P₃, P₄

Redundancy-Aware Top-k Patterns


- Desired patterns: high significance & low redundancy




(a) a set of patterns



(b) redundancy-aware top- k



(c) traditional top- k



(d) summarization

- Method: Use MMS (Maximal Marginal Significance) for measuring the combined significance of a pattern set
- Xin et al., Extracting Redundancy-Aware Top-K Patterns, KDD'o6

Advanced Frequent Pattern Mining

- Mining Diverse Patterns
- **Constraint-Based Frequent Pattern Mining**
- Mining High-Dimensional Data and Colossal Patterns
- Sequential Pattern Mining
- Graph Pattern Mining

Why Constraint-Based Mining?

- Finding **all** the patterns in a dataset **autonomously**? — unrealistic!
 - Too many patterns but not necessarily user-interested!
- Pattern mining should be an **interactive** process
 - User directs what to be mined using a **data mining query language** (or a graphical user interface)
- Constraint-based mining
 - User flexibility: provides **constraints** on what to be mined
 - Optimization: explores such constraints for efficient mining
 - **Constraint-based mining**: Constraint-pushing, similar to push selection first in DB query processing

Constraints in General Data Mining

A data mining query can be in the form of a meta-rule or with the following language primitives.

- Knowledge type constraint:
 - Ex.: classification, association, clustering, outlier finding,
- Data constraint — using SQL-like queries
 - Ex.: find products sold together in NY stores this year
- Dimension/level constraint
 - Ex.: in relevance to region, price, brand, customer category
- Rule (or pattern) constraint
 - Ex.: small sales (price < \$10) triggers big sales (sum > \$200)
- Interestingness constraint
 - Ex.: strong rules: $\text{min_sup} \geq 0.02$, $\text{min_conf} \geq 0.6$, $\text{min_correlation} \geq 0.7$

Meta-Rule Guided Mining

- A meta-rule can contain partially instantiated predicates & constants
 - $P_1(X, Y) \wedge P_2(X, W) \Rightarrow \text{buys}(X, \text{"iPad"})$
- The resulting mined rule can be
 - $\text{age}(X, \text{"15-25"}) \wedge \text{profession}(X, \text{"student"}) \Rightarrow \text{buys}(X, \text{"iPad"})$
- In general, (meta) rules can be in the form of
 - $P_1 \wedge P_2 \wedge \dots \wedge P_l \Rightarrow Q_1 \wedge Q_2 \wedge \dots \wedge Q_r$
- Method to find meta-rules
 - Find frequent ($l + r$) predicates (based on *min-support*)
 - Push constants deeply when possible into the mining process
 - Using constraint-push techniques introduced in this lecture
 - Also, push *min_conf*, *min_correlation*, and other measures as early as possible (measures acting as constraints)

Different Kinds of Constraints Lead to Different Pruning Strategies

- Constraints can be categorized as
 - Pattern space pruning constraints vs. data space pruning constraints
- Pattern space pruning constraints
 - Anti-monotonic: If constraint c is violated, its further mining can be terminated
 - Monotonic: If c is satisfied, no need to check c again
 - Succinct: if the constraint c can be enforced by directly manipulating the data
 - Convertible: c can be converted to monotonic or anti-monotonic if items can be properly ordered in processing
- Data space pruning constraints
 - Data succinct: Data space can be pruned at the initial pattern mining process
 - Data anti-monotonic: If a transaction t does not satisfy c , then t can be pruned to reduce data processing effort

Pattern Space Pruning with Pattern Anti-Monotonicity

- Constraint c is anti-monotone
 - If an itemset S violates constraint c , so does any of its superset
 - That is, mining on itemset S can be terminated
- Ex. 1: $c_1: \text{sum}(S.\text{price}) \leq v$ is anti-monotone
- Ex. 2: $c_2: \text{range}(S.\text{profit}) \leq 15$ is anti-monotone
 - Itemset ab violates c_2 ($\text{range}(ab) = 40$)
 - So does every superset of ab
- Ex. 3. $c_3: \text{sum}(S.\text{Price}) \geq v$ is not anti-monotone
- Ex. 4. Is $c_4: \text{support}(S) \geq \sigma$ anti-monotone?
 - Yes! Apriori pruning is essentially pruning with an anti-monotonic constraint!

TID	Transaction	Item	Profit
10	a, b, c, d, f, h	a	40
20	b, c, d, f, g, h	b	0
30	b, c, d, f, g	c	-20
40	a, c, e, f, g	d	-15
		e	-30
		f	-10
		g	20
		h	5

min_sup = 2

price(item) > 0

Pattern Monotonicity and Its Roles

- A constraint c is monotone: if an itemset S satisfies the constraint c , so does any of its superset
 - That is, we do not need to check c in subsequent mining
- Ex. 1: c_1 : $\text{sum}(S.\text{Price}) \geq v$ is monotone
- Ex. 2: c_2 : $\text{min}(S.\text{Price}) \leq v$ is monotone
- Ex. 3: c_3 : $\text{range}(S.\text{profit}) \geq 15$ is monotone
 - Itemset ab satisfies c_3
 - So does every superset of ab

TID	Transaction	Item	Profit
10	a, b, c, d, f, h	a	40
20	b, c, d, f, g, h	b	0
30	b, c, d, f, g	c	-20
40	a, c, e, f, g	d	-15
		e	-30
		f	-10
		g	20
		h	5

$\text{min_sup} = 2$

$\text{price}(\text{item}) > 0$

Data Space Pruning with Data Anti-Monotonicity

- A constraint c is **data anti-monotone**: In the mining process, if a data entry t cannot satisfy a pattern p under c , t cannot satisfy p 's superset either
 - Data space pruning: Data entry t can be pruned
- Ex. 1: $c_1: \text{sum}(S.\text{Profit}) \geq v$ is **data anti-monotone**
 - Let constraint c_1 be: $\text{sum}\{S.\text{Profit}\} \geq 25$
 - $T_{30}: \{b, c, d, f, g\}$ can be removed since none of their combinations can make an S whose sum of the profit is ≥ 25
- Ex. 2: $c_2: \text{min}(S.\text{Price}) \leq v$ is **data anti-monotone**
 - Consider $v = 5$ but every item in transaction T_{50} has a price higher than 10
- Ex. 3: $c_3: \text{range}(S.\text{Profit}) \geq 25$ is **data anti-monotone**

TID	Transaction	Item	Profit
10	a, b, c, d, f, h	a	40
20	b, c, d, f, g, h	b	0
30	b, c, d, f, g	c	-20
40	a, c, e, f, g	d	-15
		e	-30
		f	-10
		g	20
		h	5

min_sup = 2

price(item) > 0

Data Space Pruning Should Be Explored Recursively

Example. $c_3: \text{range}(S.\text{Profit}) > 25$

- We check b's projected database
 - But item "a" is infrequent ($\text{sup} = 1$)
- After removing "a (40)" from T_{10}
 - T_{10} cannot satisfy c_3 any more
 - since "b (0)" and "c (-20), d (-15), f (-10), h (5)"
 - By removing T_{10} , we can also prune "h" in T_{20}

b's-proj. DB		TID	Transaction	Item	Profit
TID	Transaction	10	a, b, c, d, f, h	a	40
10	a, c, d, f, h	20	b, c, d, f, g, h	b	0
20	c, d, f, g, h	30	b, c, d, f, g	c	-20
30	c, d, f, g	40	a, c, e, f, g	d	-15
				e	-30
				f	-10
				g	20
				h	5

$\text{min_sup} = 2$

$\text{price}(\text{item}) > 0$

Constraint:
 $\text{range}\{S.\text{profit}\} > 25$

b's-proj. DB

TID	Transaction
10	a, c, d, f, h
20	c, d, f, g, h
30	c, d, f, g

Recursive
Data
Pruning

b's FP-tree

single branch: cdfg: 2

Only a single branch "cdfg: 2"
to be mined in b's projected DB

- Note: c_3 prunes T_{10} effectively only after "a" is pruned (by min-sup) in b's projected DB

Succinctness: Pruning Both Data and Pattern Spaces

- Succinctness: if the constraint c can be enforced by directly manipulating the data
- Ex. 1: To find those patterns without item i
 - Remove i from DB and then mine (pattern space pruning)
- Ex. 2: To find those patterns containing item i
 - Mine only i -projected DB (data space pruning)
- Ex. 3: $c_3: \min(S.\text{Price}) \leq v$ is succinct
 - Start with only items whose price $\leq v$ (pattern space pruning) and remove transactions with high-price items only (data space pruning)
- Ex. 4: $c_4: \sum(S.\text{Price}) \geq v$ is not succinct
 - It cannot be determined beforehand since sum of the price of itemset S keeps increasing

Convertible Constraints: Ordering Data in Transactions

- Convert tough constraints into (anti-)monotone by proper ordering of items in transactions
- Examine c_1 : $\text{avg}(S.\text{profit}) > 20$
 - Order items in value-descending order
 - $\langle a, g, f, h, b, d, c, e \rangle$
 - An itemset ab violates c_1 ($\text{avg}(ab) = 20$)
 - So does ab^* (i.e., ab -projected DB)
 - C_1 : anti-monotone if patterns grow in the right order!
- Can item-reordering work for Apriori?
 - Does not work for level-wise candidate generation!
 - $\text{avg}(agf) = 23.3 > 20$, but $\text{avg}(gfh) = 15 < 20$

		Item	Profit
		a	40
		b	0
		c	-20
		d	-15
		e	-30
		f	10
	min_sup = 2	g	20
	price(item)>0	h	5

TID	Transaction
10	a, b, c, d, f, h
20	b, c, d, f, g, h
30	b, c, d, f, g
40	a, c, e, f, g

26

How to Handle Multiple Constraints?

- It is beneficial to use multiple constraints in pattern mining
- But different constraints may require potentially conflicting item-ordering
 - If there exists an order R making both c_1 and c_2 convertible, try to sort items in the order that benefits pruning most
 - If there exists conflict ordering between c_1 and c_2
 - Try to sort data and enforce one constraint first (which one?)
 - Then enforce the other when mining the projected databases
- Ex. c_1 : $\text{avg}(S.\text{profit}) > 20$, and c_2 : $\text{avg}(S.\text{price}) < 50$
 - Sorted in profit descending order and use c_1 first (assuming c_1 has more pruning power)
 - For each project DB, sort trans. in price ascending order and use c_2 at mining

Advanced Frequent Pattern Mining

- Mining Diverse Patterns
- Constraint-Based Frequent Pattern Mining
- **Mining High-Dimensional Data and Colossal Patterns**
- Sequential Pattern Mining
- Graph Pattern Mining

Mining Long Patterns: Challenges

- Mining long patterns is needed in bioinformatics, social network analysis, software engineering, ...
 - But the methods introduced so far mine only short patterns (e.g., length < 10)
- Challenges of mining long patterns
 - The curse of “downward closure” property of frequent patterns
 - Any sub-pattern of a frequent pattern is frequent
 - If $\{a_1, a_2, \dots, a_{100}\}$ is frequent, then $\{a_1\}, \{a_2\}, \dots, \{a_{100}\}, \{a_1, a_2\}, \{a_1, a_3\}, \dots, \{a_1, a_{100}\}, \{a_1, a_2, a_3\}, \dots$ are all frequent! There are about 2^{100} such frequent itemsets!
 - No matter searching in breadth-first (e.g., Apriori) or depth-first (e.g., FPgrowth), **if we still adopt the “small to large” step-by-step growing paradigm**, we have to examine so many patterns, which leads to combinatorial explosion!

Colossal Patterns: A Motivating Example

- Let min-support $\sigma = 20$
- # of closed/maximal patterns of size 20: about $\binom{40}{20}$
- But there is only one pattern with size close to 40 (*i.e.*, long or colossal)
 - $\alpha = \{41, 42, \dots, 79\}$ of size 39
- Q: How to find it without generating an exponential number of size-20 patterns?

$T_1 = 2\ 3\ 4\ \dots\ 39\ 40$

$T_2 = 1\ 3\ 4\ \dots\ 39\ 40$

: .

: .

$T_{40} = 1\ 2\ 3\ 4\ \dots\ 39$

$T_{41} = 41\ 42\ 43\ \dots\ 79$

$T_{42} = 41\ 42\ 43\ \dots\ 79$

: .

: .


$T_{60} = 41\ 42\ 43\ \dots\ 79$

The existing fastest mining algorithms (*e.g.*, FPClose, LCM) fail to complete running

A new algorithm, *Pattern-Fusion*, outputs this colossal pattern in seconds

What Is Pattern-Fusion?

- Not strive for completeness (why?)
- Jump out of the swamp of the mid-sized intermediate “results”
- Strive for mining almost complete and representative colossal patterns: identify “short-cuts” and take “leaps”
- Key observation
 - The larger the pattern or the more distinct the pattern, the greater chance it will be generated from small ones
- Philosophy: Collection of small patterns hints at the larger patterns
- Pattern fusion strategy (“not crawl but jump”): Fuse small patterns together in one step to generate new pattern candidates of significant sizes



Observation: Colossal Patterns and Core Patterns

- Suppose dataset D contains 4 colossal patterns (below) plus many small patterns
 - $\{a_1, a_2, \dots, a_{50}\}: 40, \{a_3, a_6, \dots, a_{99}\}: 60, \{a_5, a_{10}, \dots, a_{95}\}: 80, \{a_{10}, a_{20}, \dots, a_{100}\}: 100$
- If you check the pattern pool of size-3, you may likely find
 - $\{a_2, a_4, a_{45}\}: \sim 40, \{a_3, a_{34}, a_{39}\}: \sim 40, \dots, \{a_5, a_{15}, a_{85}\}: \sim 80, \dots, \{a_{20}, a_{40}, a_{85}\}: \sim 80, \dots$
- If you merge the patterns with similar support, you may obtain candidates of much bigger size and easily validate whether they are true patterns
- *Core patterns* of a colossal pattern α : A set of subpatterns of α that cluster around α by sharing a similar support
- A colossal pattern has far more core patterns than a small-sized pattern
- A random draw from a complete set of pattern of size c would be more likely to pick a core pattern (or its descendant) of a colossal pattern
- A colossal pattern can be generated by merging a set of core patterns

Robustness of Colossal Patterns

- Core Patterns: For a frequent pattern α , a subpattern β is a τ -core pattern of α if β shares a similar support set with α , i.e.,
$$\frac{|D_\alpha|}{|D_\beta|} \geq \tau \quad 0 < \tau \leq 1$$
 where τ is called the core ratio
- (d, τ) -robustness: A pattern α is (d, τ) -robust if d is the maximum number of items that can be removed from α for the resulting pattern to remain a τ -core pattern of α
- For a (d, τ) -robust pattern α , it has $\Omega(2^d)$ core patterns
- Robustness of Colossal Patterns: A colossal pattern tends to have much more core patterns than small patterns
- Such core patterns can be clustered together to form “dense balls” based on pattern distance defined by

A random draw in the pattern space will hit somewhere in the ball with high probability

$$Dist(\alpha, \beta) = 1 - \frac{|D_\alpha \cap D_\beta|}{|D_\alpha \cup D_\beta|}$$

The Pattern-Fusion Algorithm

- Initialization (Creating initial pool): Use an existing algorithm to mine all frequent patterns up to a small size, e.g., 3
- Iteration (Iterative Pattern Fusion):
 - At each iteration, K seed patterns are randomly picked from the current pattern pool
 - For each seed pattern thus picked, we find all the patterns within a bounding ball centered at the seed pattern
 - All these patterns found are fused together to generate a set of super-patterns
 - All the super-patterns thus generated form a new pool for the next iteration
- Termination: when the current pool contains no more than K patterns at the beginning of an iteration


Experimental Results on Data Set: ALL

- ALL: A popular gene expression clinical data set on ALL-AML leukemia, with 38 transactions, each with 866 columns. There are 1,736 items in total.
 - When minimum support is high (e.g., 30), Pattern-Fusion gets all the largest colossal patterns with size greater than 85

Pattern Size	110	107	102	91	86	84	83
The complete set	1	1	1	1	1	2	6
Pattern-Fusion	1	1	1	1	1	1	4

Pattern Size	82	77	76	75	74	73	71
The complete set	1	2	1	1	1	2	1
Pattern-Fusion	0	2	0	1	1	1	1

Mining colossal patterns on a Leukemia dataset



Algorithm runtime comparison on another dataset

Advanced Frequent Pattern Mining

- Mining Diverse Patterns
- Constraint-Based Frequent Pattern Mining
- Mining High-Dimensional Data and Colossal Patterns
- **Sequential Pattern Mining**
- Graph Pattern Mining

Sequence Databases and Sequential Patterns

- Sequential pattern mining has broad applications
 - Customer shopping sequences
 - Purchase a laptop first, then a digital camera, and then a smartphone, within 6 months
 - Medical treatments, natural disasters (e.g., earthquakes), science & engineering processes, stocks and markets, ...
 - Weblog click streams, calling patterns, ...
 - Software engineering: Program execution sequences, ...
 - Biological sequences: DNA, protein, ...
- Transaction DB, sequence DB vs. time-series DB
- Gapped vs. non-gapped sequential patterns
 - Shopping sequences, clicking streams vs. biological sequences

Sequential Pattern and Sequential Pattern Mining

- Sequential pattern mining: Given a set of sequences, find the complete set of frequent subsequences (i.e., satisfying the min_sup threshold)

A sequence database

SID	Sequence
10	<a(<u>ab</u> c)(a <u>c</u>)d(cf)>
20	<(ad)c(bc)(ae)>
30	<(ef)(<u>ab</u>)(df) <u>c</u> b>
40	<eg(af)cbc>

A sequence: <(ef)(ab)(df)c b>

- An element may contain a set of items (also called events)
- Items within an element are unordered and we list them alphabetically

<a(bc)dc> is a subsequence of <a(abc)(ac)d(cf)>

- Given support threshold min_sup = 2, <(ab)c> is a sequential pattern

Sequential Pattern Mining Algorithms

- Algorithm requirement: **Efficient, scalable, finding complete set, incorporating various kinds of user-specific constraints**
- The Apriori property still holds: If a subsequence s_1 is infrequent, none of s_1 's super-sequences can be frequent
- Representative algorithms
 - **GSP** (Generalized Sequential Patterns): Srikant & Agrawal @ EDBT'96)
 - Vertical format-based mining: **SPADE** (Zaki@Machine Learning'00)
 - Pattern-growth methods: **PrefixSpan** (Pei, et al. @TKDE'04)
- Mining closed sequential patterns: **CloSpan** (Yan, et al. @SDM'03)
- Constraint-based sequential pattern mining

GSP: Apriori-Based Sequential Pattern Mining

- Initial candidates: All singleton sequences
 - <a>, , <c>, <d>, <e>, <f>, <g>, <h>
- Scan DB once, count support for each candidate
- Generate length-2 candidate sequences

SID	Sequence
10	<(bd)cb(ac)>
20	<(bf)(ce)b(fg)>
30	<(ah)(bf)abf>
40	<(be)(ce)d>
50	<a(bd)bcb(ade)>

$min_sup = 2$

Cand.	sup
<a>	3
	5
<c>	4
<d>	3
<e>	3
<f>	2
<g>	1
<h>	1

	<a>		<c>	<d>	<e>	<f>
<a>	<aa>	<ab>	<ac>	<ad>	<ae>	<af>
	<ba>	<bb>	<bc>	<bd>	<be>	<bf>
<c>	<ca>	<cb>	<cc>	<cd>	<ce>	<cf>
<d>	<da>	<db>	<dc>	<dd>	<de>	<df>
<e>	<ea>	<eb>	<ec>	<ed>	<ee>	<ef>
<f>	<fa>	<fb>	<fc>	<fd>	<fe>	<ff>

	<a>		<c>	<d>	<e>	<f>
<a>		<(ab)>	<(ac)>	<(ad)>	<(ae)>	<(af)>
			<(bc)>	<(bd)>	<(be)>	<(bf)>
<c>				<(cd)>	<(ce)>	<(cf)>
<d>					<(de)>	<(df)>
<e>						<(ef)>
<f>						

Length-2 candidates:
 $36 + 15 = 51$
 Without Apriori pruning:
 $8 * 8 + 8 * 7 / 2 = 9$
 2 candidates

GSP
 (Generalized Sequential Patterns):
 Srikant & Agrawal @ EDBT'96

GSP Mining and Pruning

- Repeat (for each level (i.e., length- k))
 - Scan DB to find length- k frequent sequences
 - Generate length- $(k+1)$ candidate sequences from length- k frequent sequences using Apriori
 - set $k = k+1$
- Until no frequent sequence or no candidate can be found

Sequential Pattern Mining in Vertical Data Format: The SPADE Algorithm

- A sequence database is mapped to: <SID, EID>
- Grow the subsequences (patterns) one item at a time by Apriori candidate generation

SID	Sequence
1	<a(<u>abc</u>)(a <u>c</u>)d(cf)>
2	<(ad)c(bc)(ae)>
3	<(ef)(<u>ab</u>)(df) <u>cb</u> >
4	<eg(af)cb <u>c</u> >

min_sup = 2

SID	EID	Items
1	1	a
1	2	abc
1	3	ac
1	4	d
1	5	cf
2	1	ad
2	2	c
2	3	bc
2	4	ae
3	1	ef
3	2	ab
3	3	df
3	4	c
3	5	b
4	1	e
4	2	g
4	3	af
4	4	c
4	5	b
4	6	c

a		b		...
SID	EID	SID	EID	...
1	1	1	2	
1	2	2	3	
1	3	3	2	
2	1	3	5	
2	4	4	5	
3	2			
4	3			

ab			ba			...
SID	EID (a)	EID(b)	SID	EID (b)	EID(a)	...
1	1	2	1	2	3	
2	1	3	2	3	4	
3	2	5				
4	3	5				

aba				...
SID	EID (a)	EID(b)	EID(a)	...
1	1	2	3	
2	1	3	4	

Ref: SPADE (Sequential Pattern
Discovery using Equivalent Class)
[M. Zaki 2001]

PrefixSpan: A Pattern-Growth Approach

- Prefix and suffix
 - Given $\langle a(abc)(ac)d(cf) \rangle$
 - **Prefixes:** $\langle a \rangle$, $\langle aa \rangle$, $\langle a(ab) \rangle$, $\langle a(abc) \rangle$, ...
 - **Suffix:** Prefixes-based projection
- PrefixSpan Mining: Prefix Projections
 - Step 1: Find length-1 sequential patterns
 - $\langle a \rangle$, $\langle b \rangle$, $\langle c \rangle$, $\langle d \rangle$, $\langle e \rangle$, $\langle f \rangle$
 - Step 2: Divide search space and mine each projected DB
 - $\langle a \rangle$ -projected DB,
 - $\langle b \rangle$ -projected DB,
 - ...
 - $\langle f \rangle$ -projected DB, ...

SID	Sequence
10	$\langle a(\underline{a}bc)(a\underline{c})d(cf) \rangle$
20	$\langle (ad)c(bc)(ae) \rangle$
30	$\langle (ef)(\underline{a}b)(df)\underline{c}b \rangle$
40	$\langle eg(af)cbc \rangle$


Prefix	<u>Suffix</u> (Projection)
$\langle a \rangle$	$\langle (abc)(ac)d(cf) \rangle$
$\langle aa \rangle$	$\langle (_bc)(ac)d(cf) \rangle$
$\langle ab \rangle$	$\langle (_c)(ac)d(cf) \rangle$

PrefixSpan (Prefix-projected Sequential pattern mining) Pei, et al. @TKDE'o4


CloSpan: Mining Closed Sequential Patterns

- A **closed sequential pattern** s : There exists no superpattern s' such that $s' \supset s$, and s' and s have the same support
- Which ones are closed? $\langle abc \rangle: 20$, $\langle abcd \rangle: 20$, $\langle abcde \rangle: 15$
- Why directly mine closed sequential patterns?
 - Reduce # of (redundant) patterns
 - Attain the same expressive power
- Property P_1 : If $s \supset s_1$, s is closed iff two project DBs have the same size
- Explore **Backward Subpattern** and **Backward Superpattern** pruning to prune redundant search space
- Greatly enhances efficiency (Yan, et al., SDM'03)

Backward subpattern pruning



Backward superpattern pruning



Constraint-Based Sequential-Pattern Mining

- Share many similarities with constraint-based itemset mining
- **Anti-monotonic:** If S violates c , the super-sequences of S also violate c
 - $\text{sum}(S.\text{price}) < 150; \min(S.\text{value}) > 10$
- **Monotonic:** If S satisfies c , the super-sequences of S also do so
 - $\text{element_count}(S) > 5; S \supseteq \{\text{PC}, \text{digital_camera}\}$
- **Data anti-monotonic:** If a sequence s_1 with respect to S violates c_3 , s_1 can be removed
 - $c_3: \text{sum}(S.\text{price}) \geq v$
- **Succinct:** Enforce constraint c by explicitly manipulating data
 - $S \supseteq \{\text{i-phone}, \text{MacAir}\}$
- **Convertible:** Projection based on the sorted value not sequence order
 - $\text{value_avg}(S) < 25; \text{profit_sum}(S) > 160$
 - $\text{max}(S)/\text{avg}(S) < 2; \text{median}(S) - \text{min}(S) > 5$

Timing-Based Constraints in Seq.-Pattern Mining

- **Order constraint:** Some items must happen before the other
 - {algebra, geometry} \rightarrow {calculus} (where “ \rightarrow ” indicates ordering)
 - Anti-monotonic: Constraint-violating sub-patterns pruned
- **Min-gap/max-gap constraint:** Confines two elements in a pattern
 - E.g., mingap = 1, maxgap = 4
 - Succinct: Enforced directly during pattern growth
- **Max-span constraint:** Maximum allowed time difference between the 1st and the last elements in the pattern
 - E.g., maxspan (S) = 60 (days)
 - Succinct: Enforced directly when the 1st element is determined
- **Window size constraint:** Events in an element do not have to occur at the same time: Enforce max allowed time difference
 - E.g., window-size = 2: Various ways to merge events into elements

Episodes and Episode Pattern Mining


- Episodes and regular expressions: Alternative to seq. patterns
 - Serial episodes: $A \rightarrow B$
 - Parallel episodes: $A \mid B$ **Indicating partial order relationships**
 - Regular expressions: $(A|B)C^*(D \rightarrow E)$
- Methods for episode pattern mining
 - Variations of Apriori/GSP-like algorithms
 - Projection-based pattern growth
 - Q_1 : Can you work out the details?
 - Q_2 : What are the differences between mining episodes and constraint-based pattern mining?

Advanced Frequent Pattern Mining

- Mining Diverse Patterns
- Constraint-Based Frequent Pattern Mining
- Mining High-Dimensional Data and Colossal Patterns
- Sequential Pattern Mining
- **Graph Pattern Mining**


Frequent (Sub)Graph Patterns

- Given a labeled graph dataset $D = \{G_1, G_2, \dots, G_n\}$, the supporting graph set of a subgraph g is $D_g = \{G_i \mid g \subseteq G_i, G_i \in D\}$.
 - $\text{support}(g) = |D_g| / |D|$
- A (sub)graph g is **frequent** if $\text{support}(g) \geq \text{min_sup}$ Ex.: Chemical structures
- Alternative:
 - Mining frequent subgraph patterns from a single large graph or network



min_sup = 2

Frequent Graph Patterns



Applications of Graph Pattern Mining


- Bioinformatics
 - Gene networks, protein interactions, metabolic pathways
- Chem-informatics: Mining chemical compound structures
- Social networks, web communities, tweets, ...
- Cell phone networks, computer networks, ...
- Web graphs, XML structures, semantic Web, information networks
- Software engineering: program execution flow analysis
- Building blocks for graph classification, clustering, compression, comparison, and correlation analysis
- Graph indexing and graph similarity search

Graph Pattern Mining Algorithms: Different Methodologies

- Generation of candidate subgraphs
 - Apriori vs. pattern growth (e.g., FSG vs. gSpan)
- Search order
 - Breadth vs. depth
- Elimination of duplicate subgraphs
 - Passive vs. active (e.g., gSpan (Yan&Han'02))
- Support calculation
 - Store embeddings (e.g., GASTON (Nijssen&Kok'04, FFSM (Huan, et al.'03), MoFa (Borgelt and Berthold ICDM'02))
- Order of pattern discovery
 - Path \rightarrow tree \rightarrow graph (e.g., GASTON (Nijssen&Kok'04))


Apriori-Based Approach

- The Apriori property (anti-monotonicity): A size- k subgraph is frequent if and only if all of its subgraphs are frequent
- A candidate size- $(k+1)$ edge/vertex subgraph is generated if its corresponding two k -edge/vertex subgraphs are frequent
- Iterative mining process:
 - Candidate-generation \rightarrow candidate pruning \rightarrow support counting \rightarrow candidate elimination



Candidate Generation: Vertex Growing vs. Edge Growing


- Methodology: breadth-search, Apriori joining two size- k graphs
 - Many possibilities at generating size- $(k+1)$ candidate graphs



- Generating new graphs with one more vertex
 - AGM (Inokuchi, et al., PKDD'00)
- Generating new graphs with one more edge
 - FSG (Kuramochi and Karypis, ICDM'01)
- Performance shows via edge growing is more efficient


Pattern-Growth Approach

- Depth-first growth of subgraphs from k -edge to $(k+1)$ -edge, then $(k+2)$ -edge subgraphs
- Major challenge
 - Generating many duplicate subgraphs
- Major idea to solve the problem
 - Define an order to generate subgraphs
 - DFS spanning tree: Flatten a graph into a sequence using depth-first search
 - gSpan (Yan & Han: ICDM'02)




gSPAN: Graph Pattern Growth in Order

- **Right-most path extension** in subgraph pattern growth
 - Right-most path: The path from root to the right-most leaf (choose the vertex w. the smallest index at each step)
 - Reduce generation of duplicate subgraphs
- **Completeness:** The Enumeration of graphs using right-most path extension is complete
- DFS Code: Flatten a graph into a sequence using depth-first search



Why Mining Closed Graph Patterns?

- Challenge: An n -edge frequent graph may have 2^n subgraphs
- Motivation: Explore *closed frequent subgraphs* to handle graph pattern explosion problem
- A frequent graph G is *closed* if there exists no supergraph of G that carries the same support as G




If this subgraph is *closed* in the graph dataset, it implies that none of its frequent super-graphs carries the same support

- *Lossless compression*: Does not contain non-closed graphs, but still ensures that the mining result is complete
- Algorithm CloseGraph: Mines closed graph patterns directly


CloseGraph: Directly Mining Closed Graph Patterns

- CloseGraph: Mining closed graph patterns by extending gSpan
- Suppose G and G_1 are frequent, and G is a subgraph of G_1
- If **in any part of the graph in the dataset where G occurs, G_1 also occurs**, then we need not grow G (except some special, subtle cases), since none of G 's children will be closed except those of G_1




Experiment and Performance Comparison


- The AIDS antiviral screen compound dataset from NCI/NIH
- The dataset contains 43,905 chemical compounds
- Discovered Patterns: The smaller minimum support, the bigger and more interesting subgraph patterns discovered



of Patterns: Frequent vs. Closed



Runtime: Frequent vs. Closed



References: Mining Diverse Patterns

- R. Srikant and R. Agrawal, "Mining generalized association rules", VLDB'95
- Y. Aumann and Y. Lindell, "A Statistical Theory for Quantitative Association Rules", KDD'99
- K. Wang, Y. He, J. Han, "Pushing Support Constraints Into Association Rules Mining", IEEE Trans. Knowledge and Data Eng. 15(3): 642-658, 2003
- D. Xin, J. Han, X. Yan and H. Cheng, "On Compressing Frequent Patterns", Knowledge and Data Engineering, 60(1): 5-29, 2007
- D. Xin, H. Cheng, X. Yan, and J. Han, "Extracting Redundancy-Aware Top-K Patterns", KDD'o6
- J. Han, H. Cheng, D. Xin, and X. Yan, "Frequent Pattern Mining: Current Status and Future Directions", Data Mining and Knowledge Discovery, 15(1): 55-86, 2007
- F. Zhu, X. Yan, J. Han, P. S. Yu, and H. Cheng, "Mining Colossal Frequent Patterns by Core Pattern Fusion", ICDE'o7

References: Constraint-Based Frequent Pattern Mining

- R. Srikant, Q. Vu, and R. Agrawal, "Mining association rules with item constraints", KDD'97
- R. Ng, L.V.S. Lakshmanan, J. Han & A. Pang, "Exploratory mining and pruning optimizations of constrained association rules", SIGMOD'98
- G. Grahne, L. Lakshmanan, and X. Wang, "Efficient mining of constrained correlated sets", ICDE'00
- J. Pei, J. Han, and L. V. S. Lakshmanan, "Mining Frequent Itemsets with Convertible Constraints", ICDE'01
- J. Pei, J. Han, and W. Wang, "Mining Sequential Patterns with Constraints in Large Databases", CIKM'02
- F. Bonchi, F. Giannotti, A. Mazzanti, and D. Pedreschi, "ExAnte: Anticipated Data Reduction in Constrained Pattern Mining", PKDD'03
- F. Zhu, X. Yan, J. Han, and P. S. Yu, "gPrune: A Constraint Pushing Framework for Graph Pattern Mining", PAKDD'07

References: Sequential Pattern Mining

- R. Srikant and R. Agrawal, "Mining sequential patterns: Generalizations and performance improvements", EDBT'96
- M. Zaki, "SPADE: An Efficient Algorithm for Mining Frequent Sequences", Machine Learning, 2001
- J. Pei, J. Han, B. Mortazavi-Asl, J. Wang, H. Pinto, Q. Chen, U. Dayal, and M.-C. Hsu, "Mining Sequential Patterns by Pattern-Growth: The PrefixSpan Approach", IEEE TKDE, 16(10), 2004
- X. Yan, J. Han, and R. Afshar, "CloSpan: Mining Closed Sequential Patterns in Large Datasets", SDM'03
- J. Pei, J. Han, and W. Wang, "Constraint-based sequential pattern mining: the pattern-growth methods", J. Int. Inf. Sys., 28(2), 2007
- M. N. Garofalakis, R. Rastogi, K. Shim: Mining Sequential Patterns with Regular Expression Constraints. IEEE Trans. Knowl. Data Eng. 14(3), 2002
- H. Mannila, H. Toivonen, and A. I. Verkamo, "Discovery of frequent episodes in event sequences", Data Mining and Knowledge Discovery, 1997

References: Graph Pattern Mining

- C. Borgelt and M. R. Berthold, Mining molecular fragments: Finding relevant substructures of molecules, ICDM'02
- J. Huan, W. Wang, and J. Prins. Efficient mining of frequent subgraph in the presence of isomorphism, ICDM'03
- A. Inokuchi, T. Washio, and H. Motoda. An apriori-based algorithm for mining frequent substructures from graph data, PKDD'00
- M. Kuramochi and G. Karypis. Frequent subgraph discovery, ICDM'01
- S. Nijssen and J. Kok. A Quickstart in Frequent Structure Mining can Make a Difference. KDD'04
- N. Vanetik, E. Gudes, and S. E. Shimony. Computing frequent graph patterns from semistructured data, ICDM'02
- X. Yan and J. Han, gSpan: Graph-Based Substructure Pattern Mining, ICDM'02
- X. Yan and J. Han, CloseGraph: Mining Closed Frequent Graph Patterns, KDD'03
- X. Yan, P. S. Yu, J. Han, Graph Indexing: A Frequent Structure-based Approach, SIGMOD'04
- X. Yan, P. S. Yu, and J. Han, Substructure Similarity Search in Graph Databases, SIGMOD'05