

Tutorial: Data-Driven Approaches towards Malicious Behavior Modeling



Meng Jiang
University of Notre Dame



Srijan Kumar
Stanford University



Christos Faloutsos
Carnegie Mellon
University



V.S. Subrahmanian
University of Maryland,
College Park

Acknowledgement



Outline

Introduction

Feature-based algorithms

Bots

Sockpuppets

Vandals

Hoaxes

Spectral-based algorithms

Visualization: “spokes”, “blocks”, “staircases”

Camouflage

Theoretical guarantee

Density-based algorithms

Ill-gotten Likes

Synchronized Behaviors

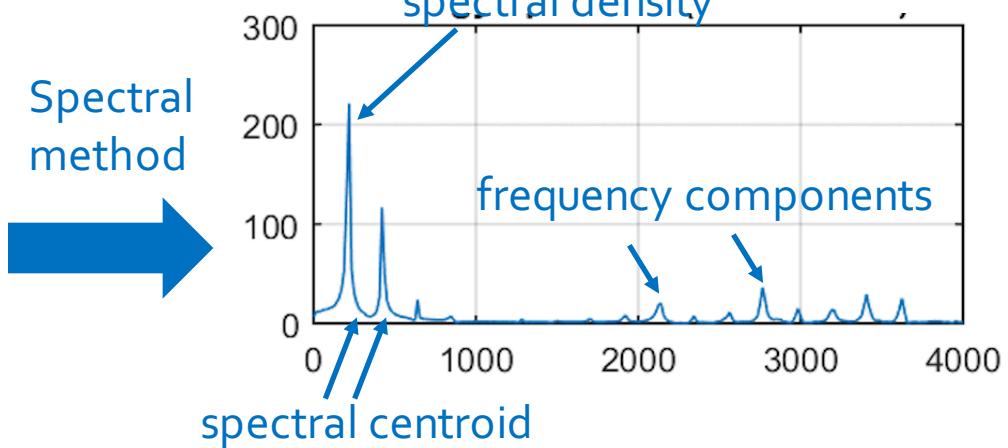
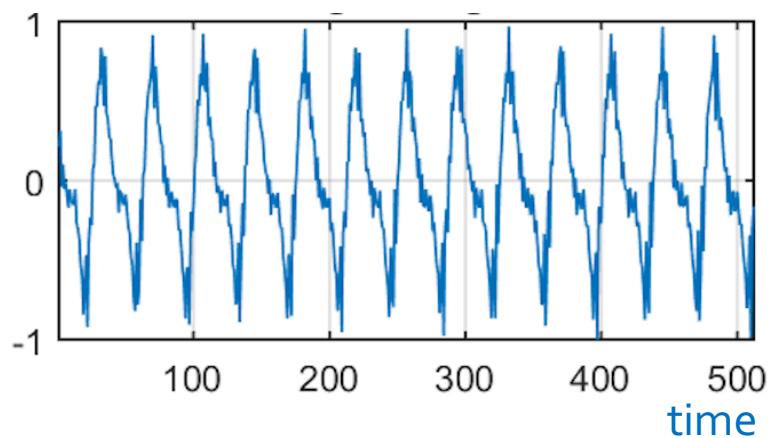
Advertising campaigns

Social spam

Conclusions and future directions

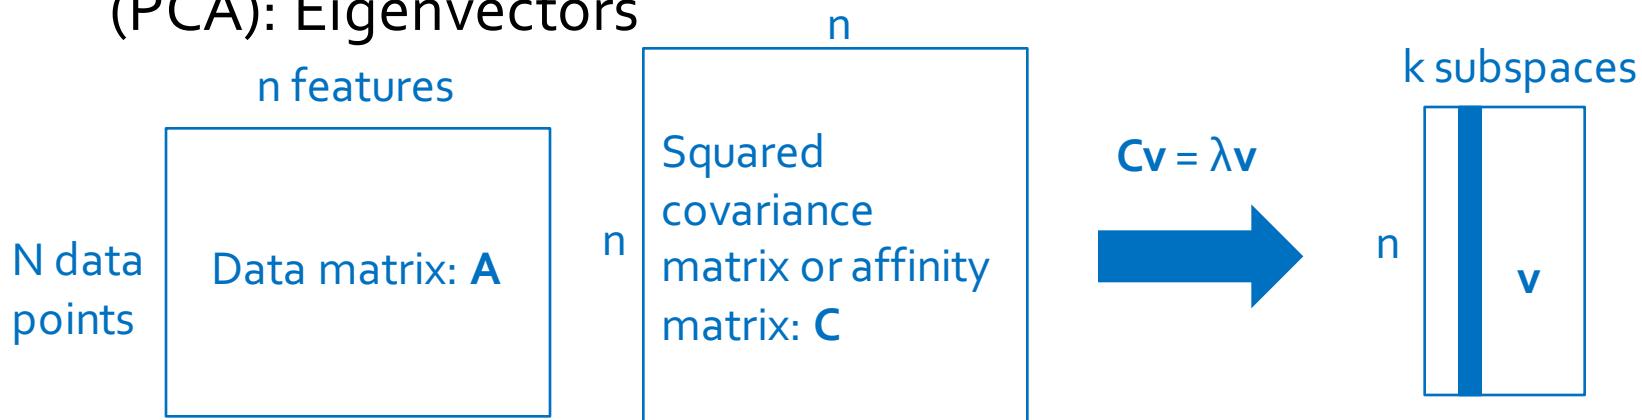
Spectral Features in Speech Signal Processing

- **Temporal features** (time domain features): which are simple to extract and have easy physical interpretation
- **Spectral features (*frequency* based features)**: which are obtained by converting the time based signal into the frequency domain using the Fourier Transform, like “frequency components”, “spectral centroid”, “spectral density”, etc.

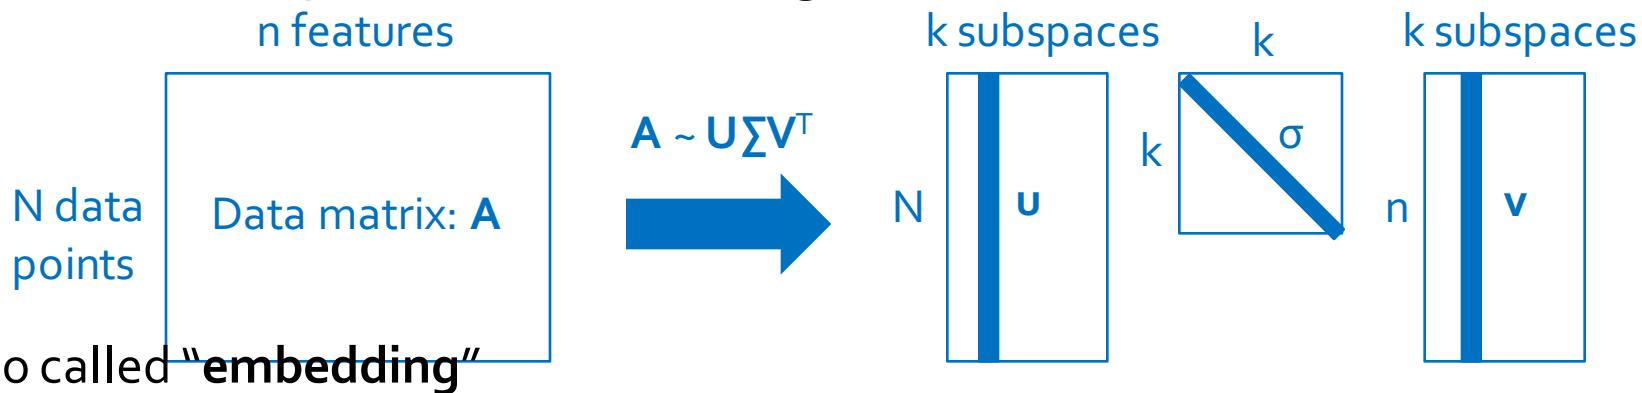


Spectral Subspaces

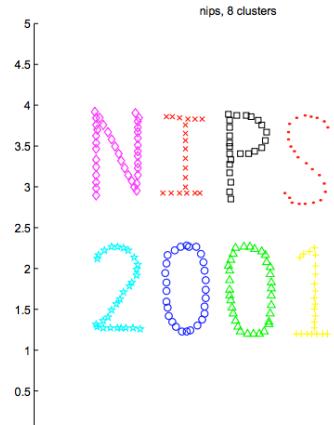
- Frequency components → Principal Component Analysis (PCA): Eigenvectors



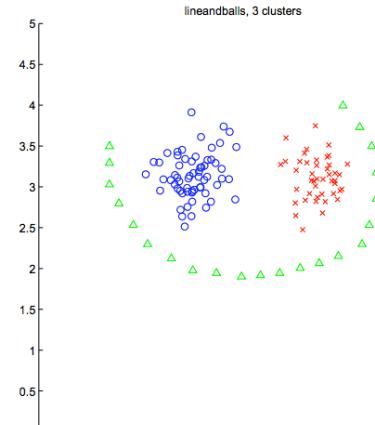
- Other spectral subspace methods. Singular Value Decomposition (SVD): Singular vectors



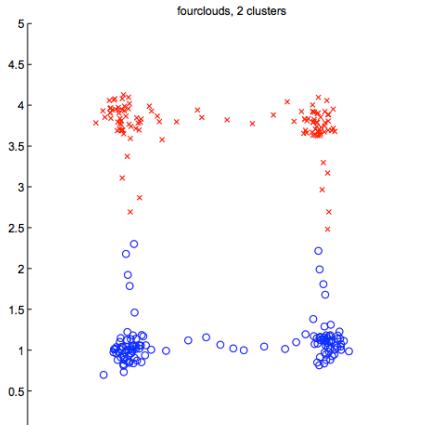
Spectral Clustering



(a)

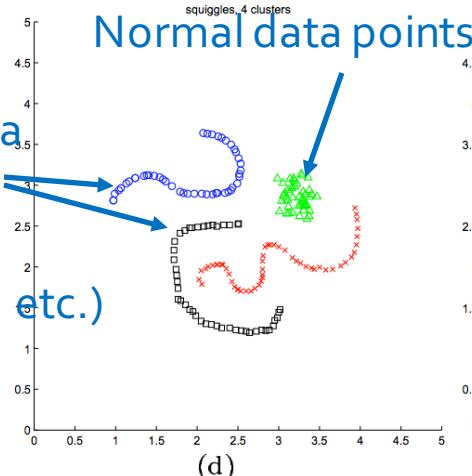


(b)

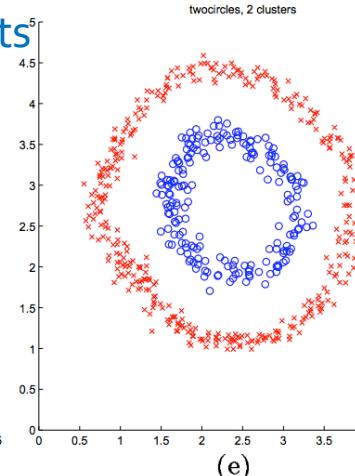


(c)

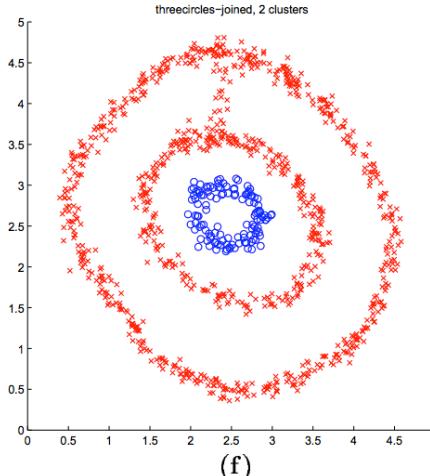
Potential application



(d)



(e)

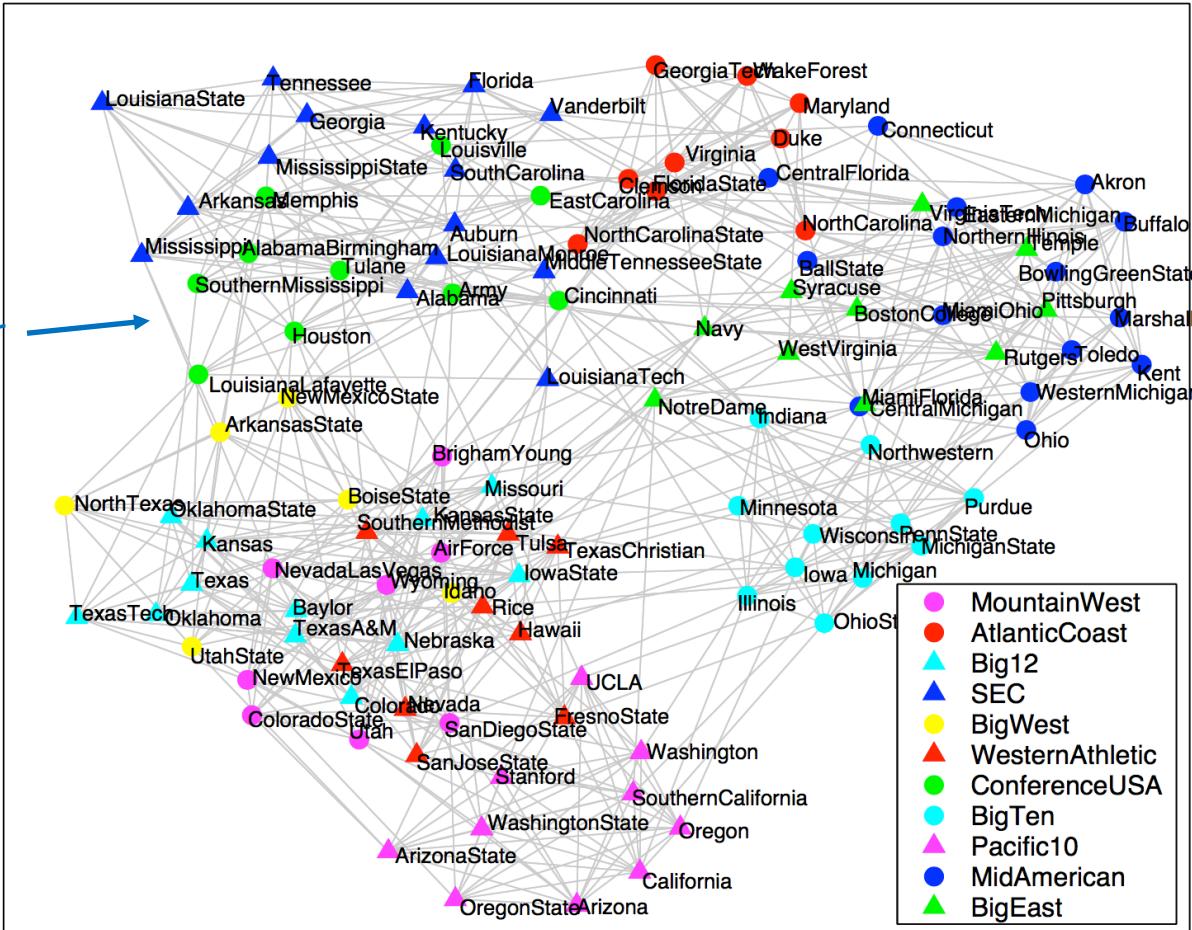


(f)

Andrew Y. Ng, Michael I. Jordan, Yair Weiss. "On Spectral Clustering: Analysis and an algorithm", NIPS 2001.

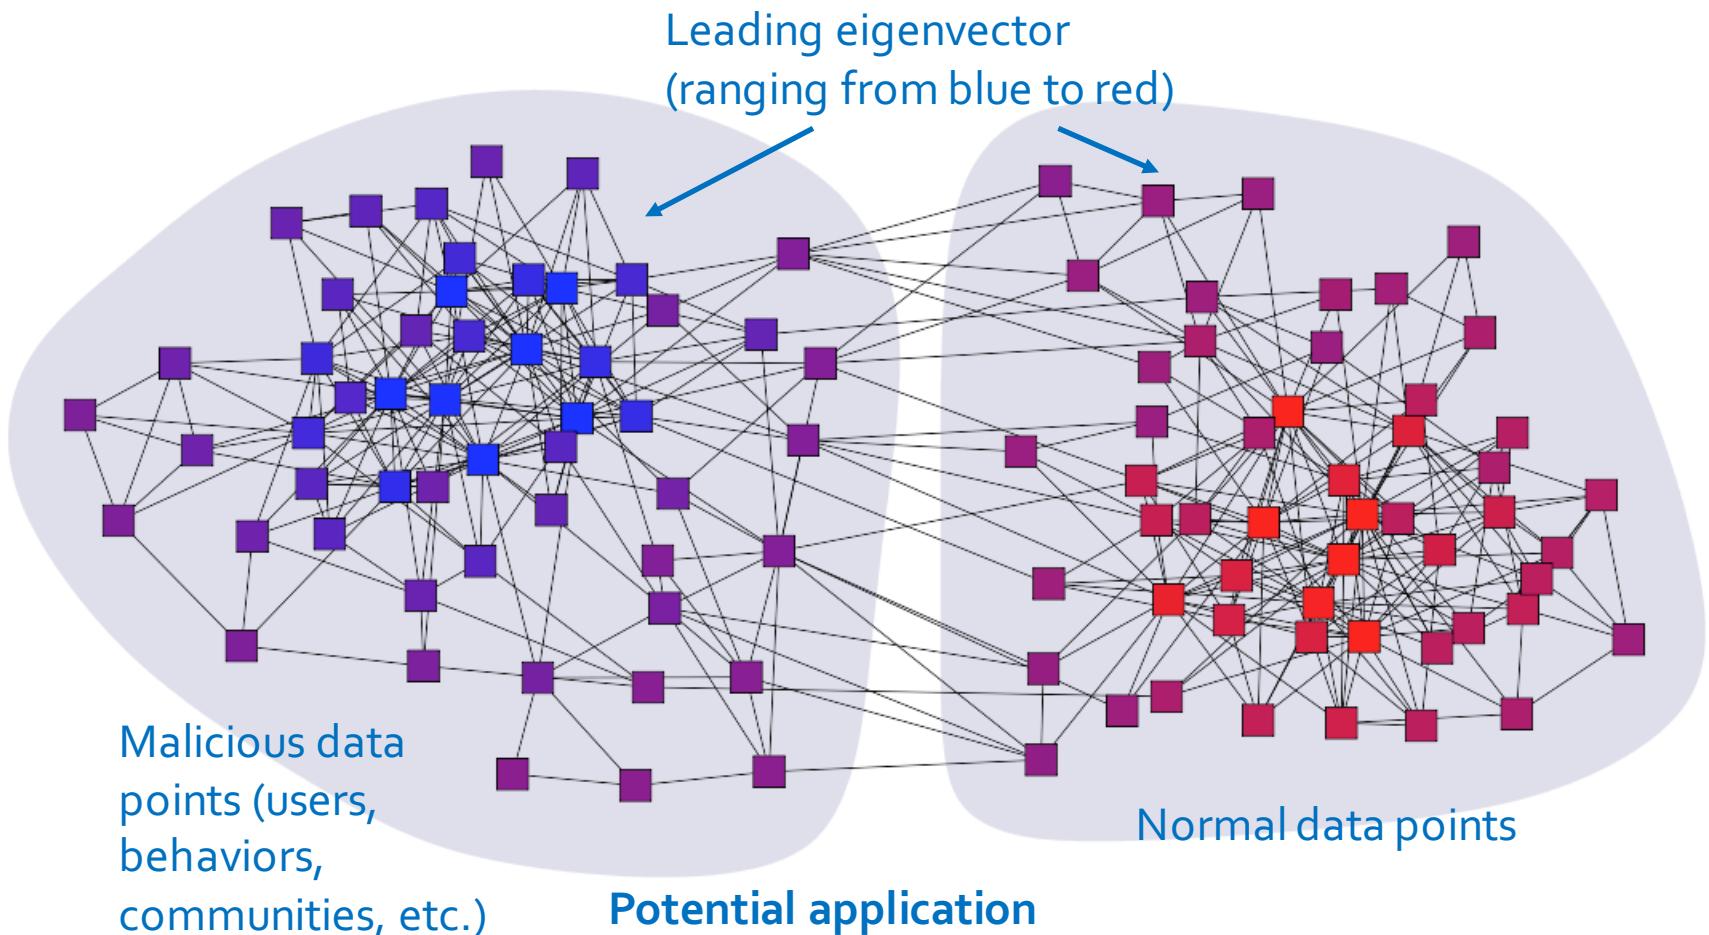
Spectral Methods for Communities

Played together



Scott White and Padhraic Smyth. "A spectral clustering approach to finding communities in graphs", SDM 2005.

Spectral Methods for Communities (cont.)



M. E. J. Newman. "Finding community structure in networks using the eigenvectors of matrices", Physical Review E 2006.

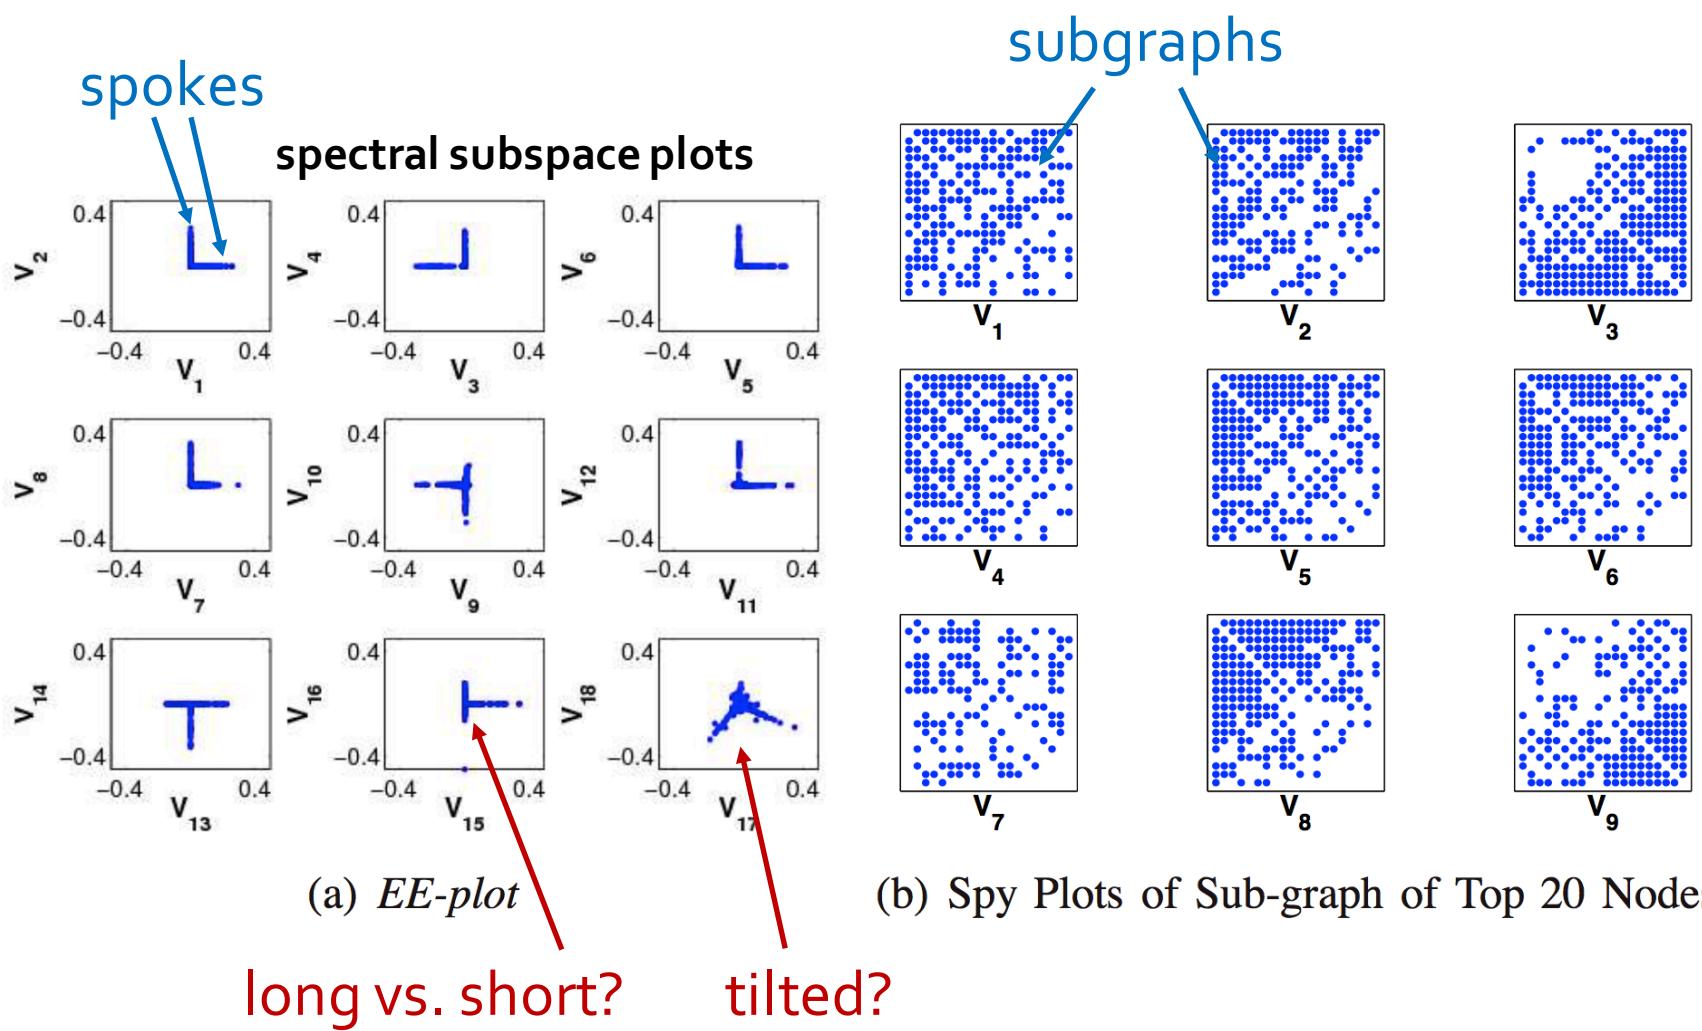
Spectral-based vs. Feature-based

- Advantages
 - Visualization: $k = \text{Number of subspaces}$
 - Feature extraction by data distribution vs. selection
 - “Principal” components: “leading” vectors
 - Data: $N\text{-by-}N$ graphs, $N\text{-by-}N\text{-by-}N$ tensors
 - Applications: Finding communities and anomalies
- Disadvantages
 - Lack of interpretability of the subspaces/features

Finding Surprising Patterns in Large Graphs

- **Problem definition:** Given a social graph based on mobile calls made from/to callers, find caller communities.
- Dataset: Activity over the duration of a month, 186,000 nodes and 464,000 edges.
- Key contribution: Discovery of the “spokes” phenomenon
 - **The singular vectors of the graph, when plotted against each other, often have clear separate lines, typically aligned with axes.**
 - Use EigenEigen (EE) plots to identify communities in the form of **cliques or near-cliques, perfect or near-perfect bipartite-cores**. (*Malicious?)

Spokes and Dense Cliques



Inferring Lockstep Behavior from Connectivity Patterns

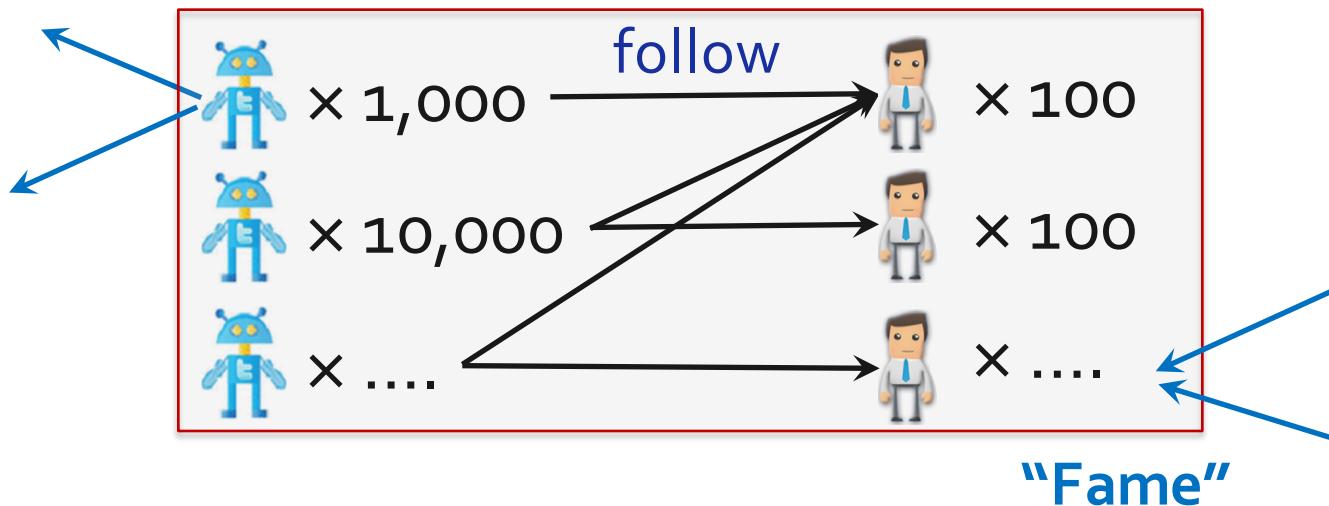
- **Problem definition:** Given a large graph, from spectral subspace plots, can we infer **lockstep behavior patterns**?



Inferring Lockstep Behavior from Connectivity Patterns

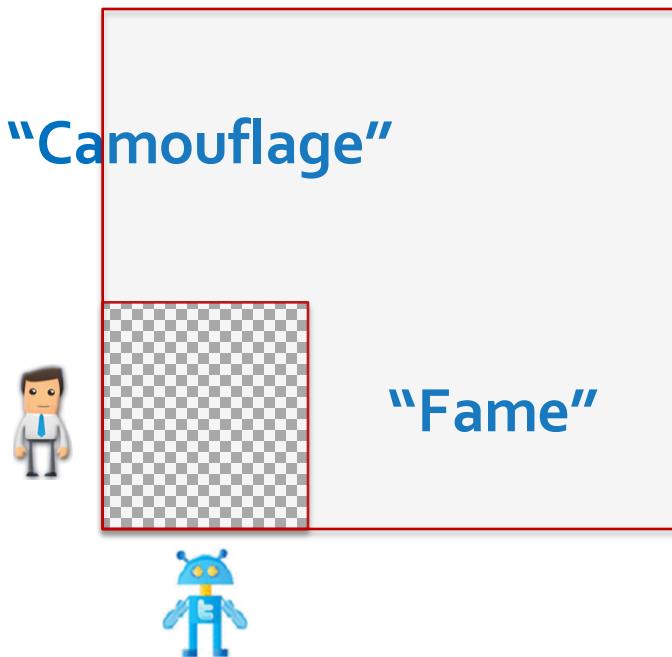
- Problem definition: Given a large graph, from **spectral subspace plots**, can we infer **lockstep behavior** patterns?

“Camouflage”



Inferring Lockstep Behavior from Connectivity Patterns

- **Problem definition:** Given a large graph, from **spectral subspace plots**, can we infer **lockstep behavior** patterns?

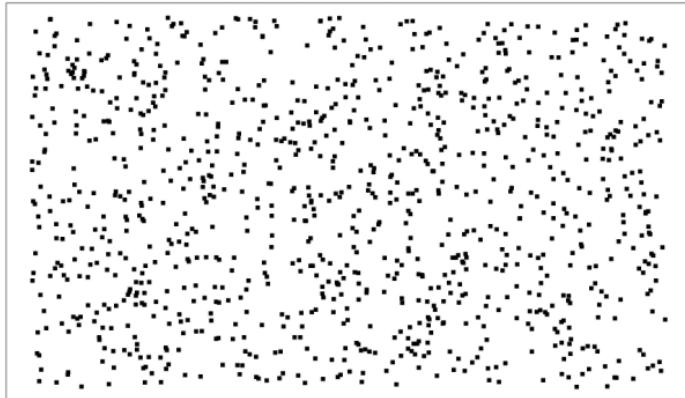


Case #0

- No lockstep behavior: Scatter

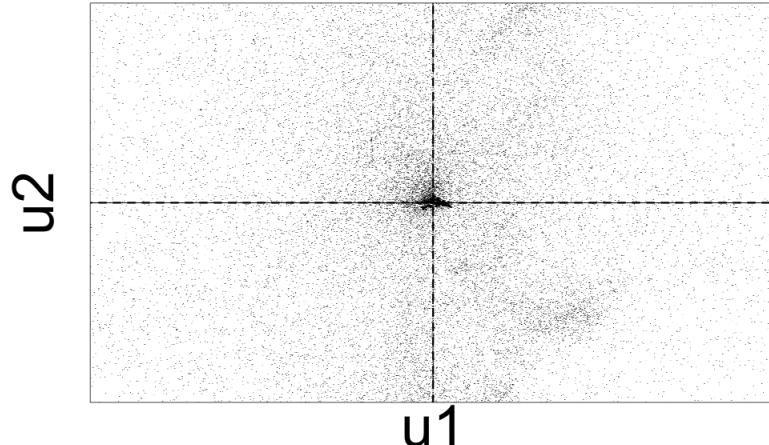
synthetic follower

Adjacency Matrix



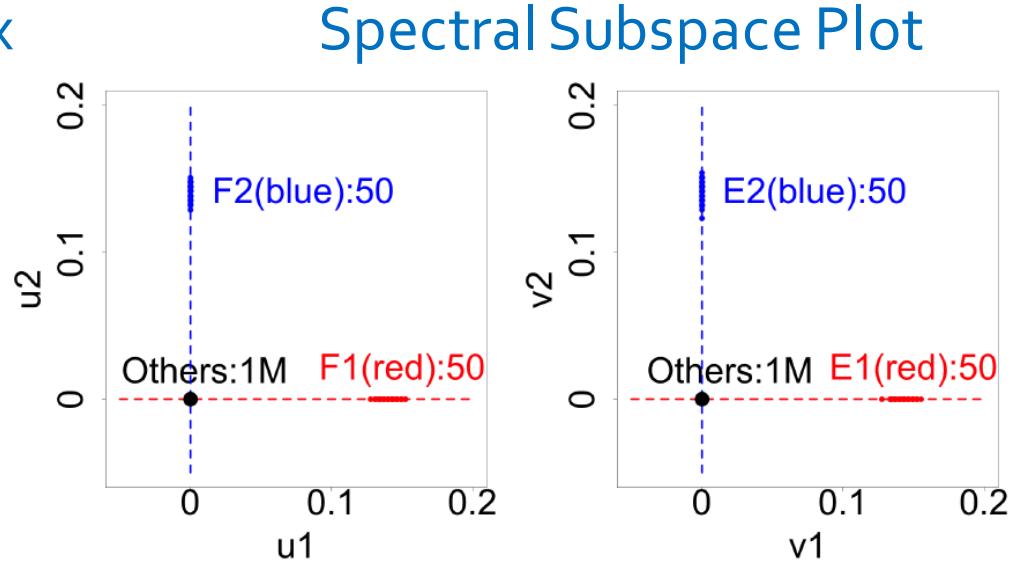
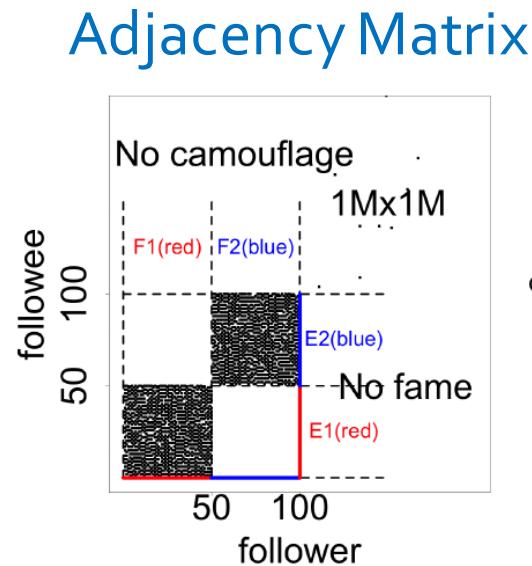
synthetic follower

Spectral Subspace Plot



Case #1

- Non-overlapping lockstep: “Rays”

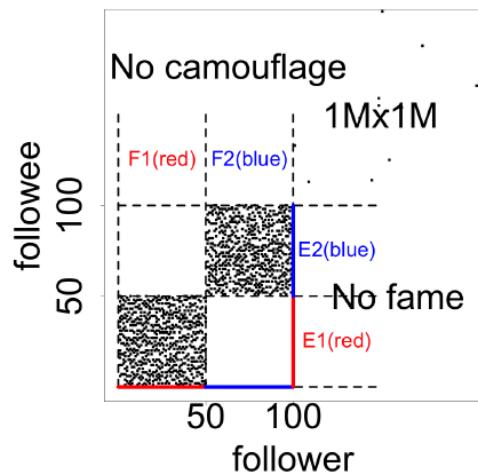


Rule 1 (short “rays”): two blocks, high density (90%), no “camouflage”, no “fame”

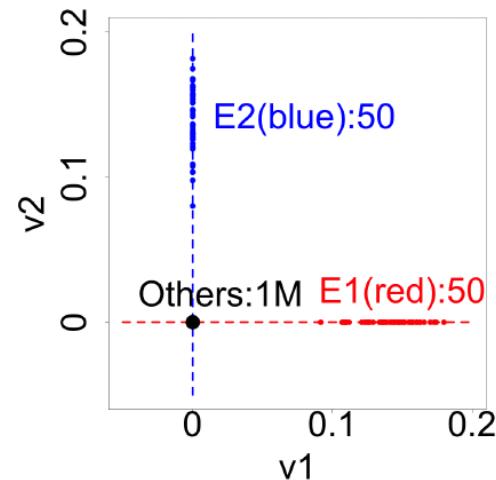
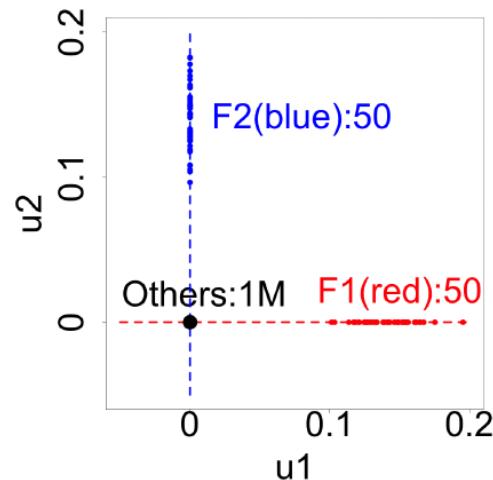
Case #2

- Non-overlapping: Low density, Elongation

Adjacency Matrix



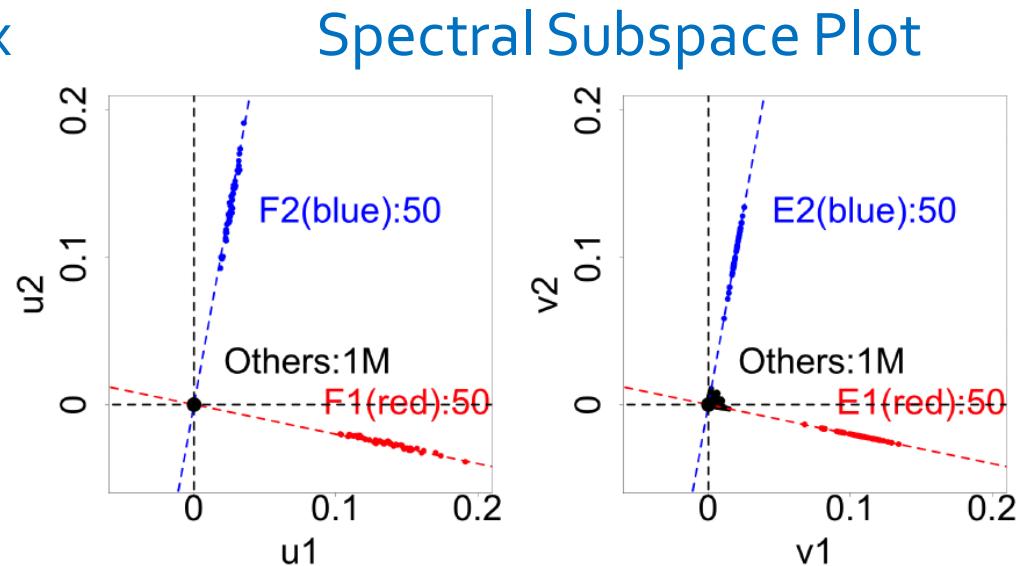
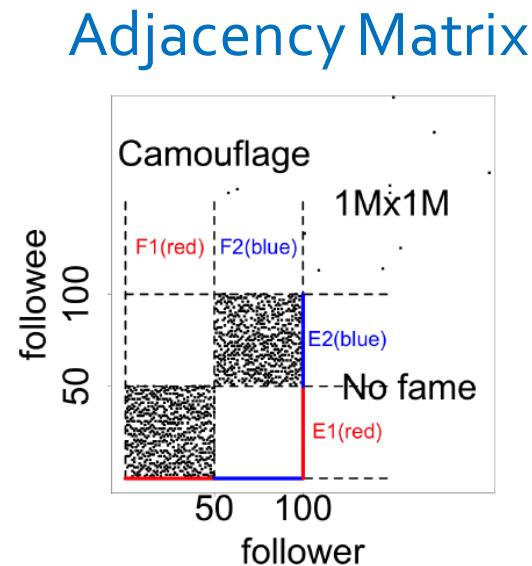
Spectral Subspace Plot



Rule 2 (long “rays”): two blocks, low density (50%), no “camouflage”, no “fame”

Case #3

- Non-overlapping: “Camouflage”/ “Fame”, Tilting

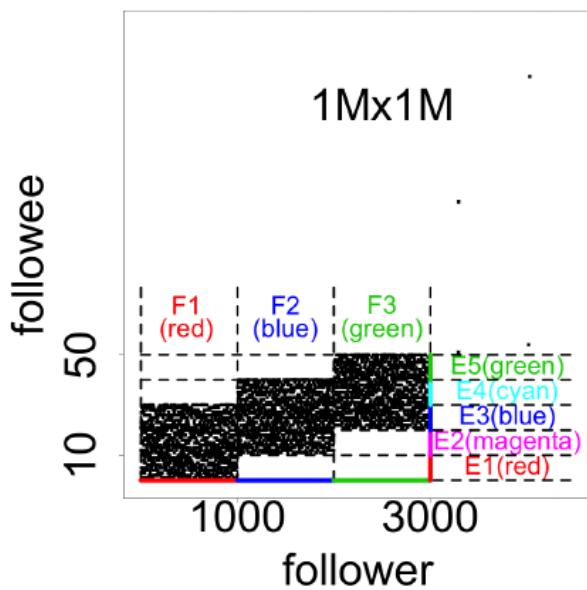


Rule 3 (tilting “rays”): two blocks, with “camouflage”, no “fame”

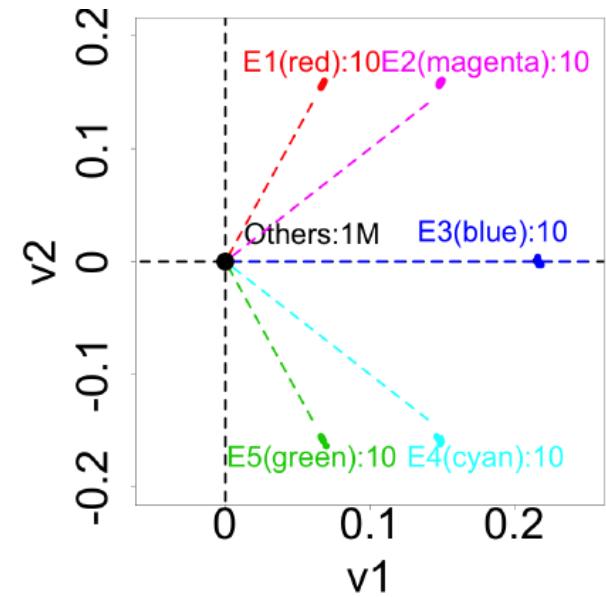
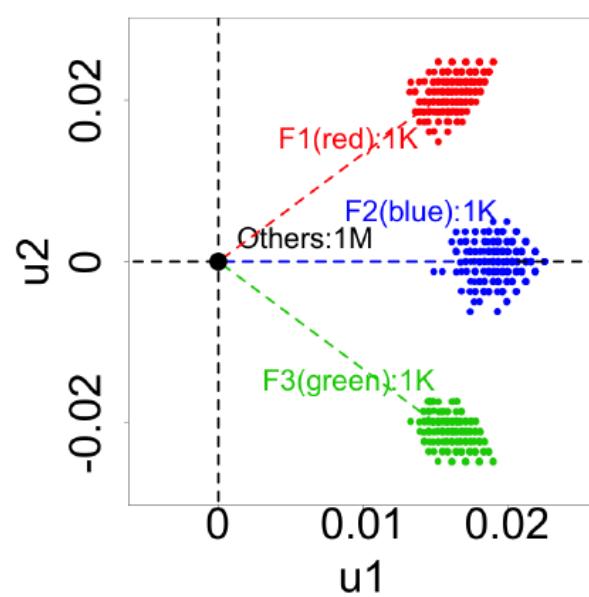
Case #4

- Overlapping: “Staircase”, “Pearls”

Adjacency Matrix

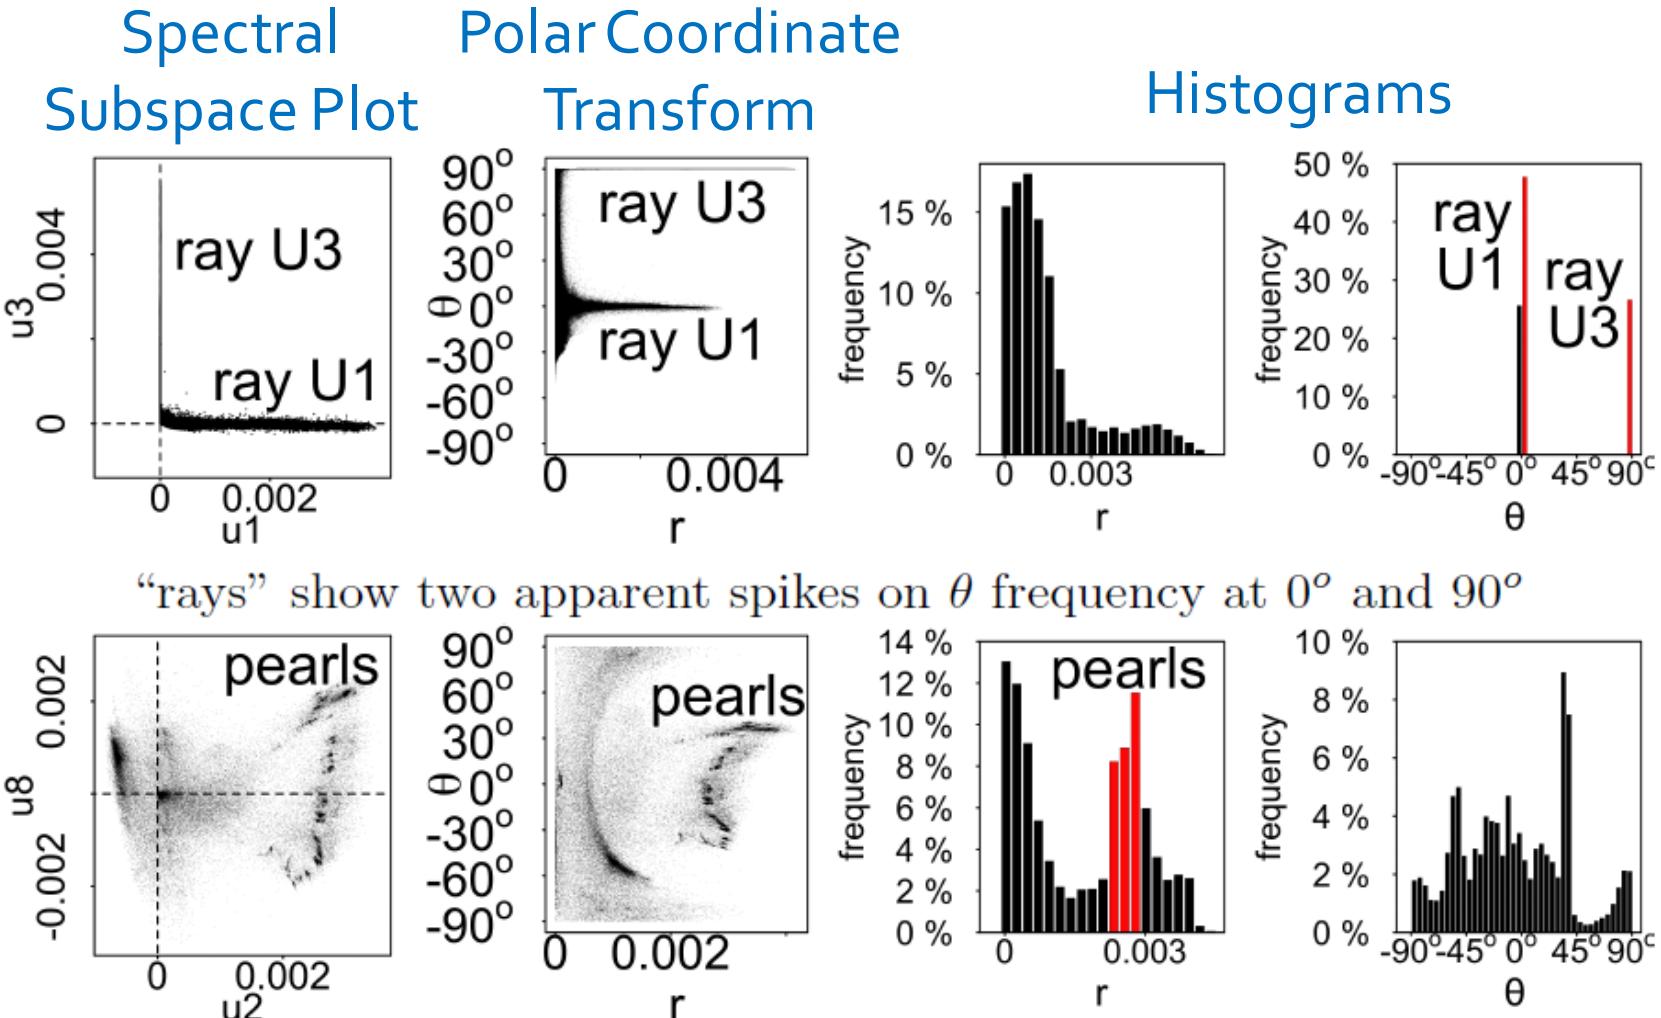


Spectral Subspace Plot



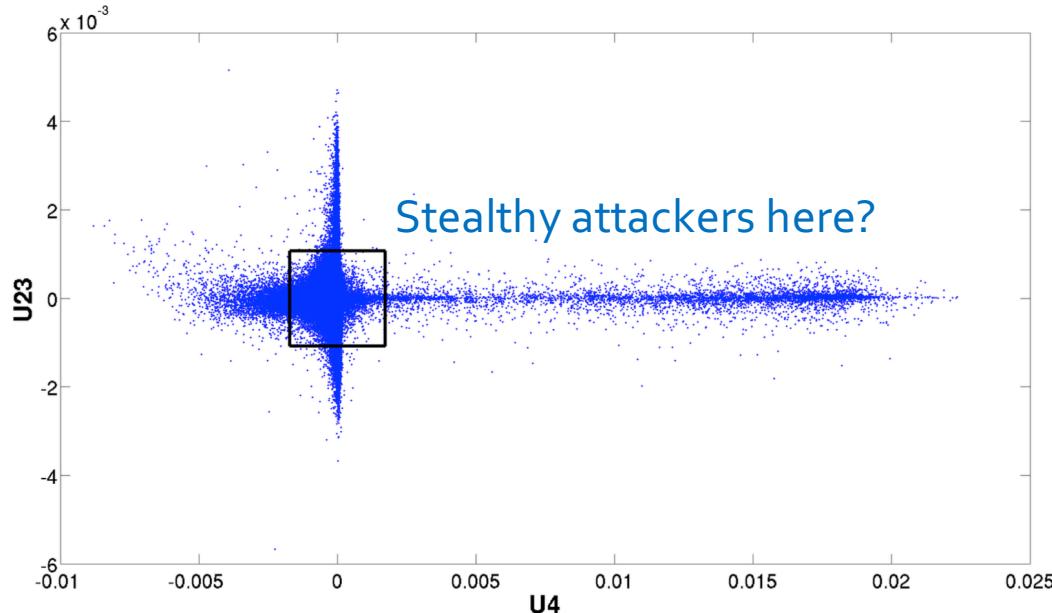
Rule 4 (“pearls”): a “staircase” of three partially overlapping blocks.

LockInfer Algorithm: Reading Spectral Subspace Plots



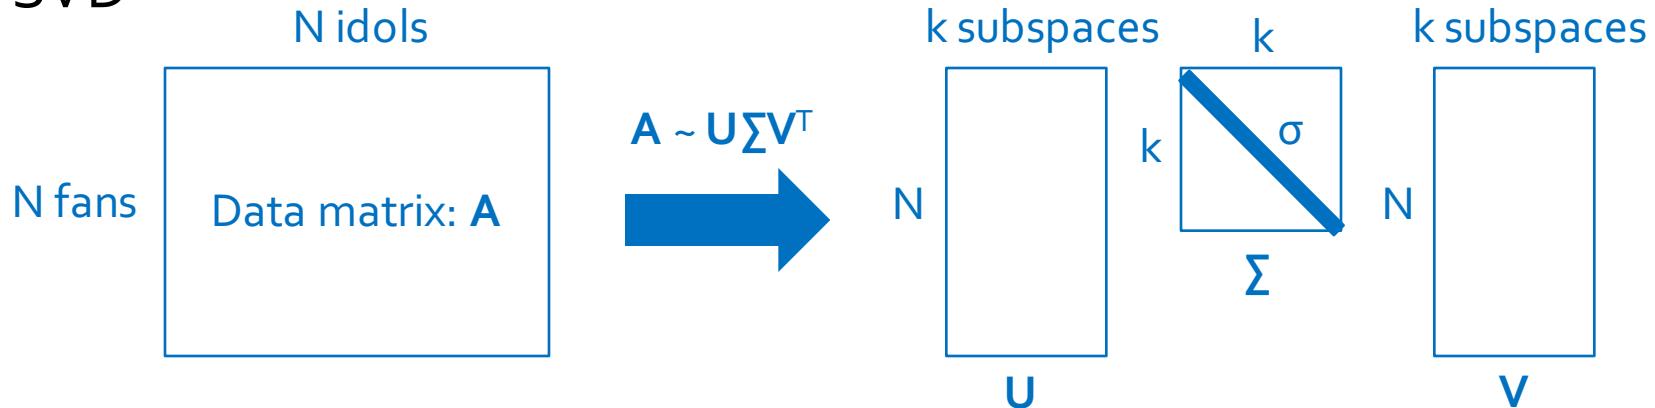
Spotting Small-Scale, Stealthy Attacks

- **Problem definition:** Can we catch stealthy attacks which are missed by traditional spectral methods?
- Dataset: Twitter “who-follows-whom” social graph, 41.7 million nodes, 1.5 billion edges

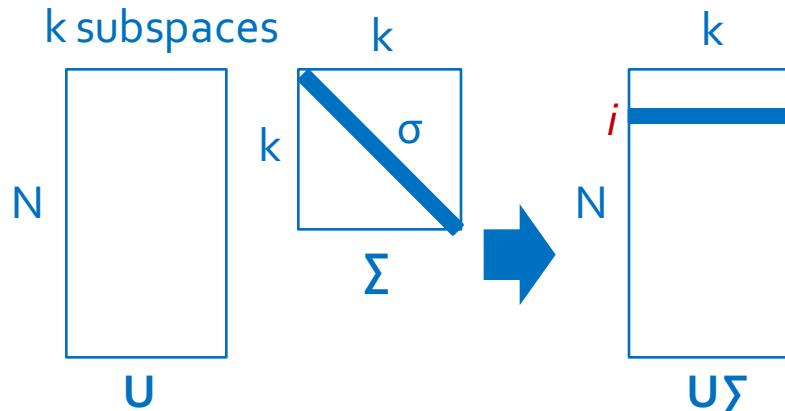


fBox: Reconstructed Degrees

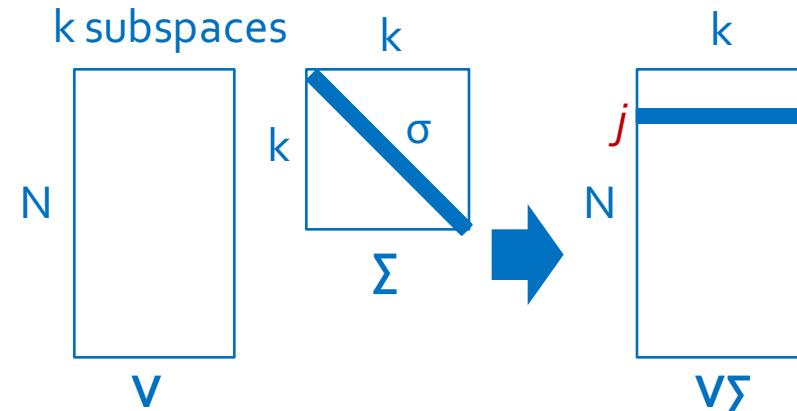
SVD



Reconstructed out-degree(i) = $\|(\mathbf{U}\Sigma)_i\|_2$



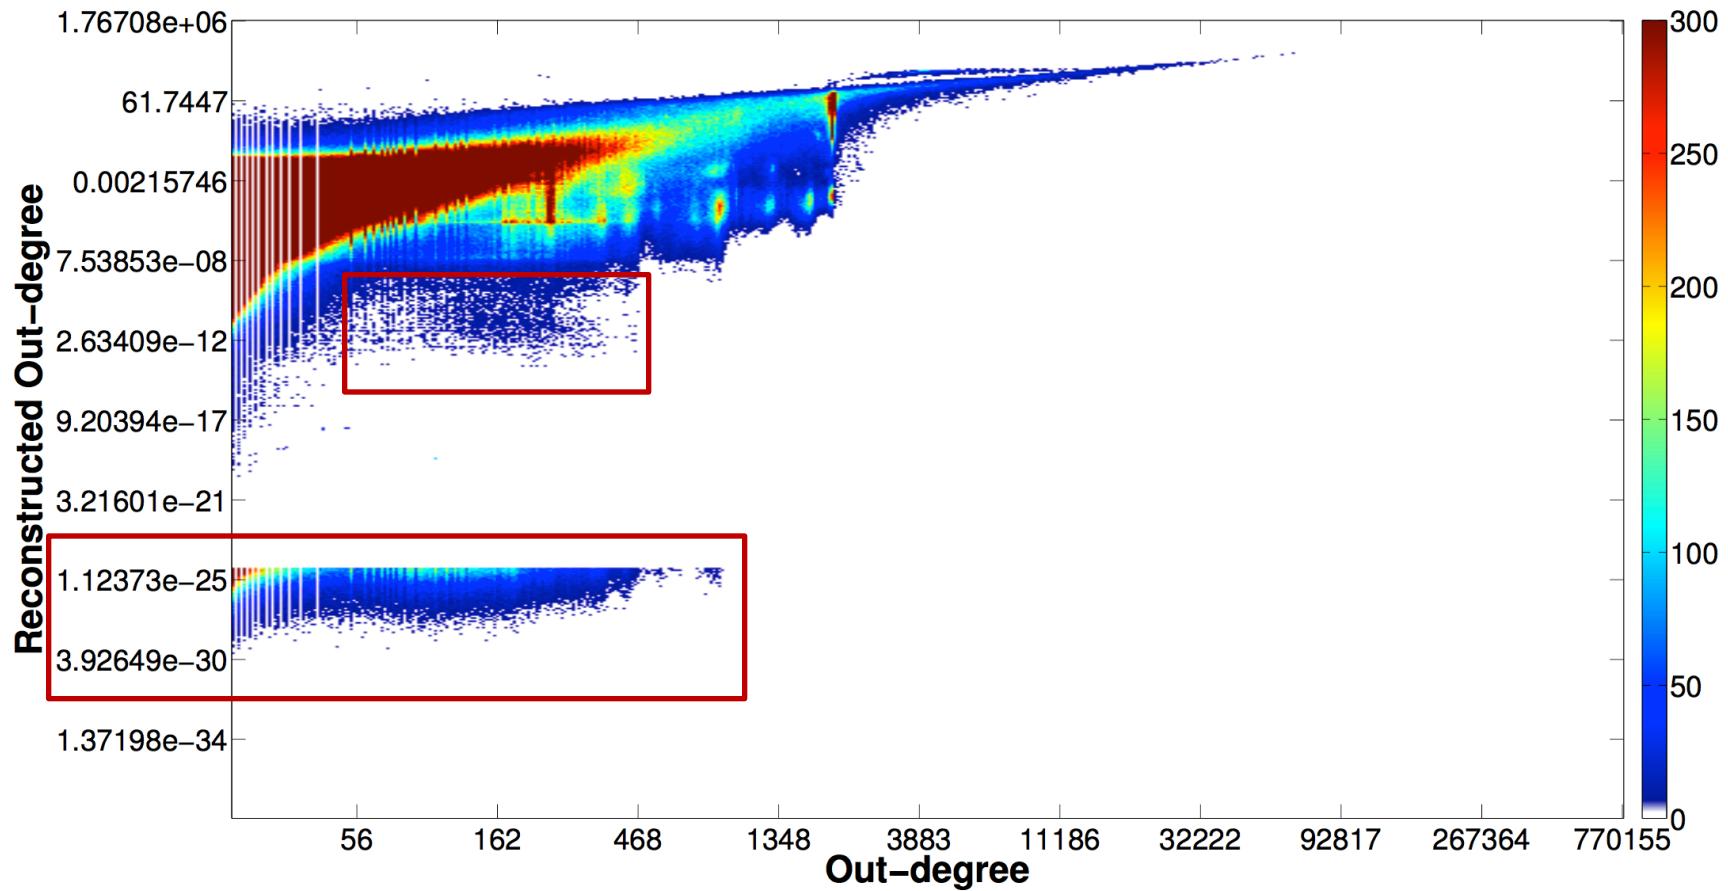
Reconstructed in-degree(j) = $\|(\mathbf{V}\Sigma)_j\|_2$



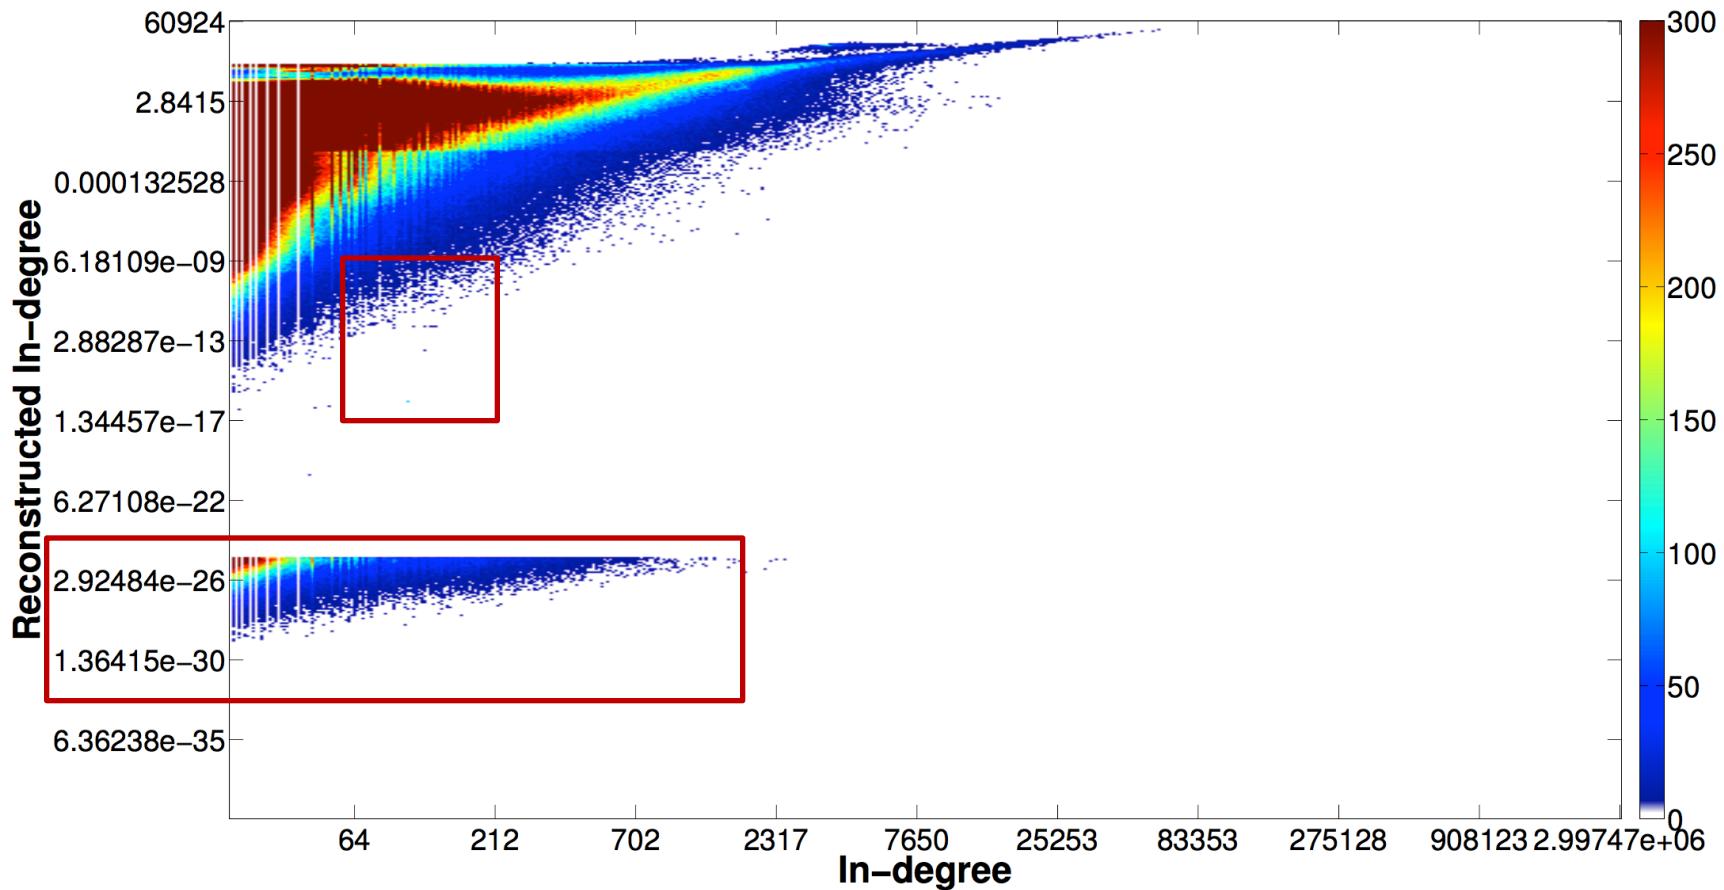
Norm-Preserving Property of SVD

- The row vectors of a full rank decomposition and associated projection will retain the same l_2 norm or vector length as in the original space:
 - For $k = \text{rank}(A)$, $\|A_i\|_2 = \|(\mathbf{U}\Sigma)_i\|_2$ and $\|A^T j\|_2 = \|(\mathbf{V}\Sigma)_j\|_2$
- Plots
 - **Reconstructed out-degree vs. Real out-degree**
 - **Reconstructed in-degree vs. Real in-degree**
 - Can we find communities (groups of a number of users) whose reconstructed degrees are very small but in the same range?

Reconstructed out-degree vs. Real out-degree



Reconstructed in-degree vs. Real in-degree



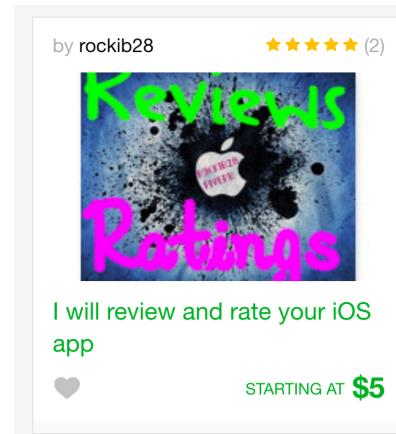
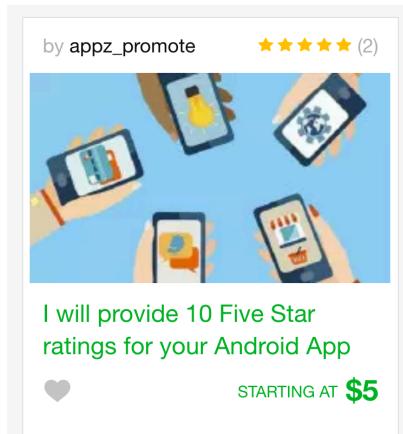
Reconstructed in-degree vs. Real in-degree (cont.)



Bounding Graph Fraud in Camouflage

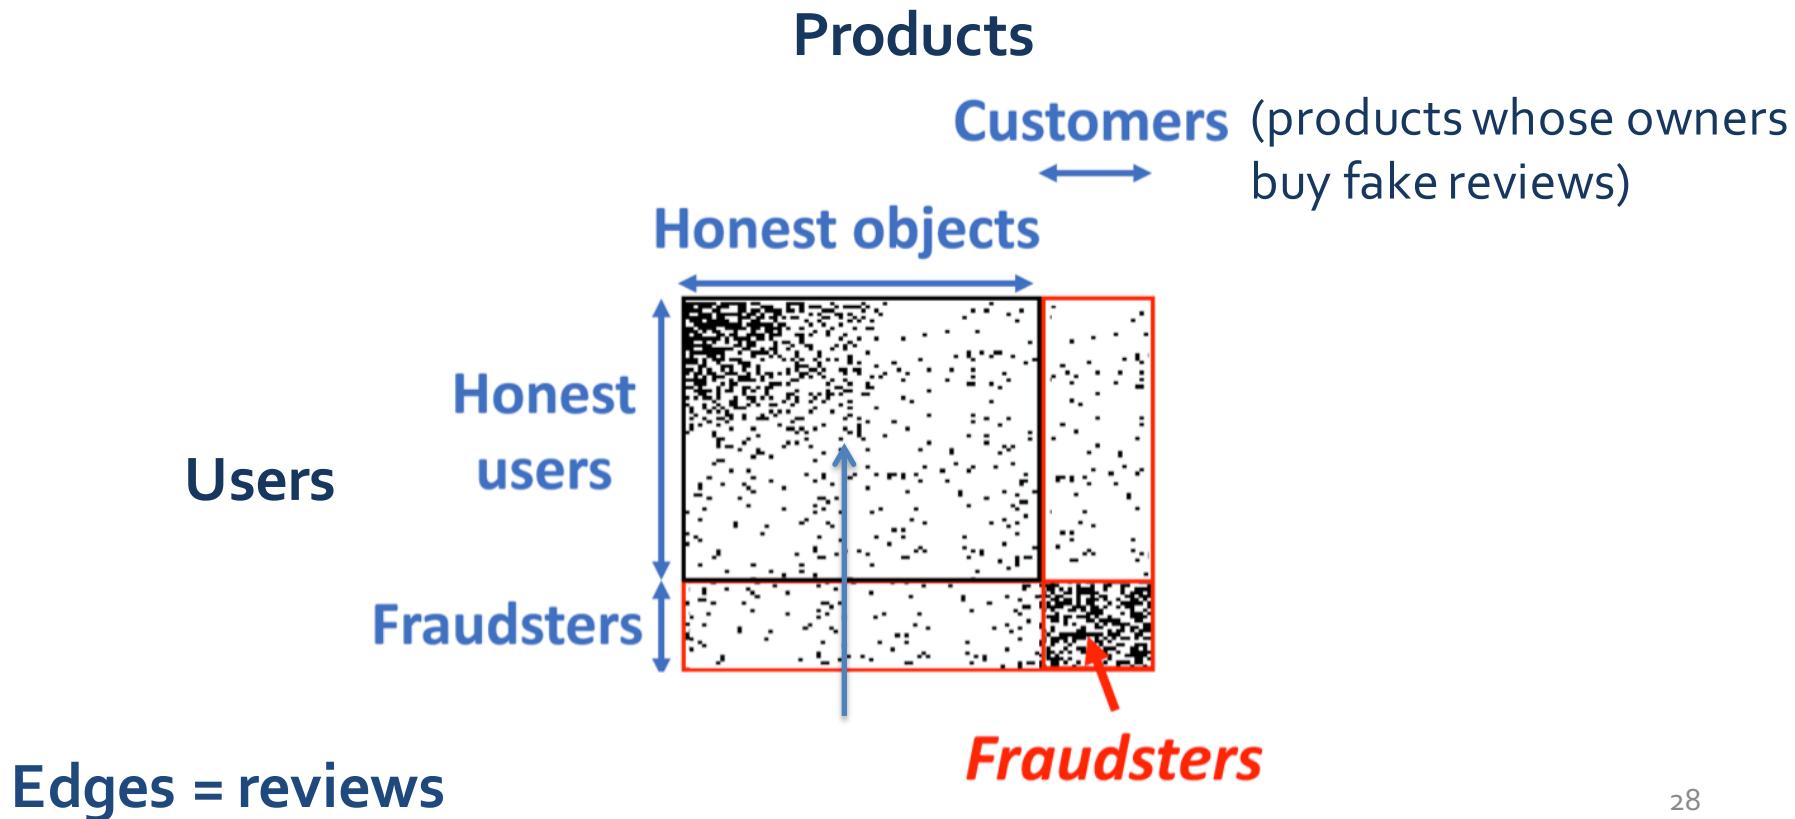
- One application: Fake reviews

I will do 5 five star reviews, all from real profiles



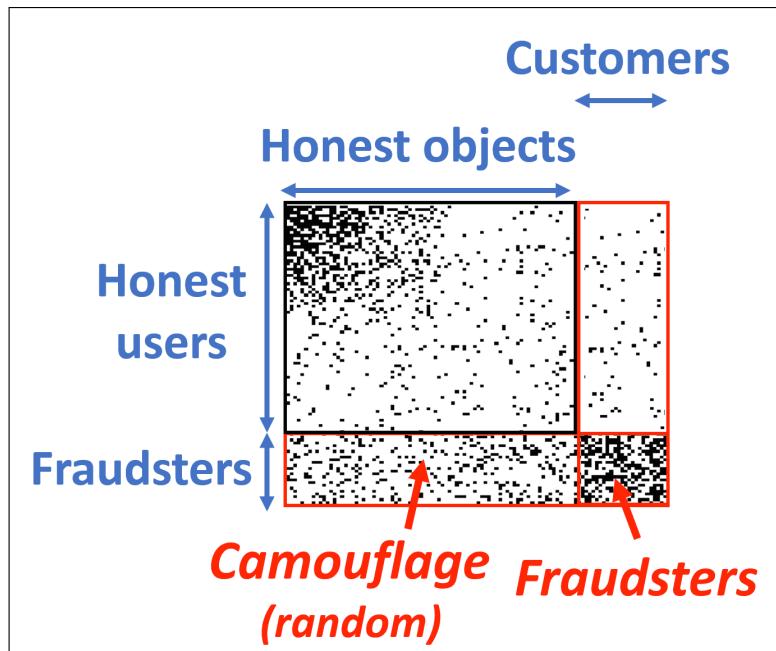
“User-Product” Review Graph

- Problem definition: Given a “user-product” review graph, can we spot fraudsters and customers?



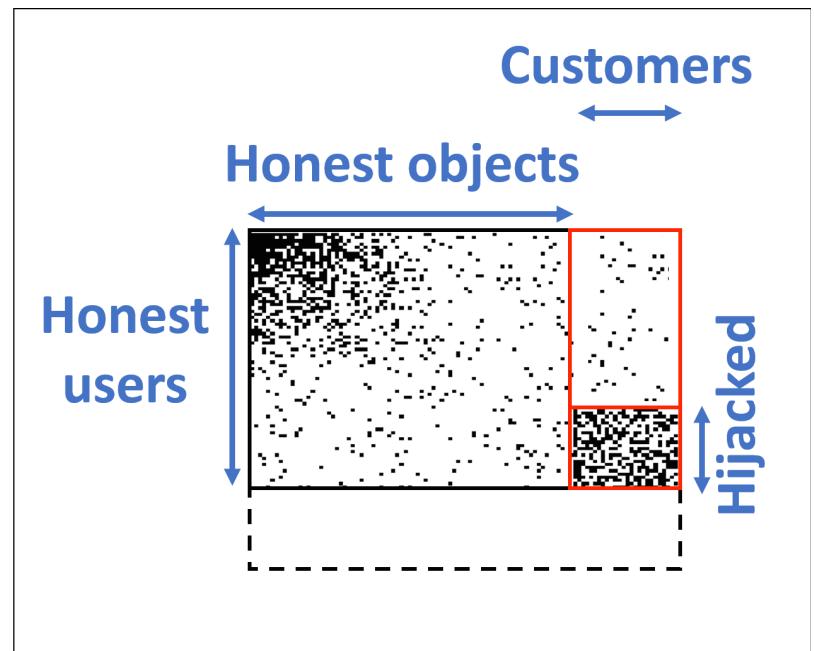
Camouflage: Evading Detection

Random camouflage



= “camouflage” in LockInfer

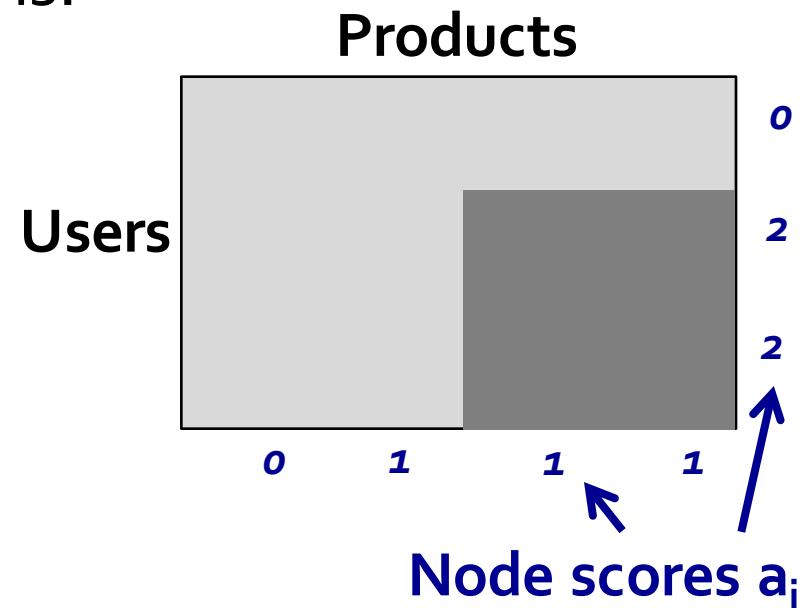
Hijacked user accounts



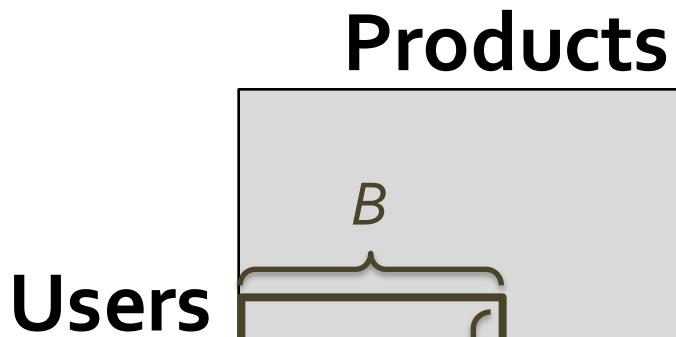
= “fame” in LockInfer

Formal Problem Definition

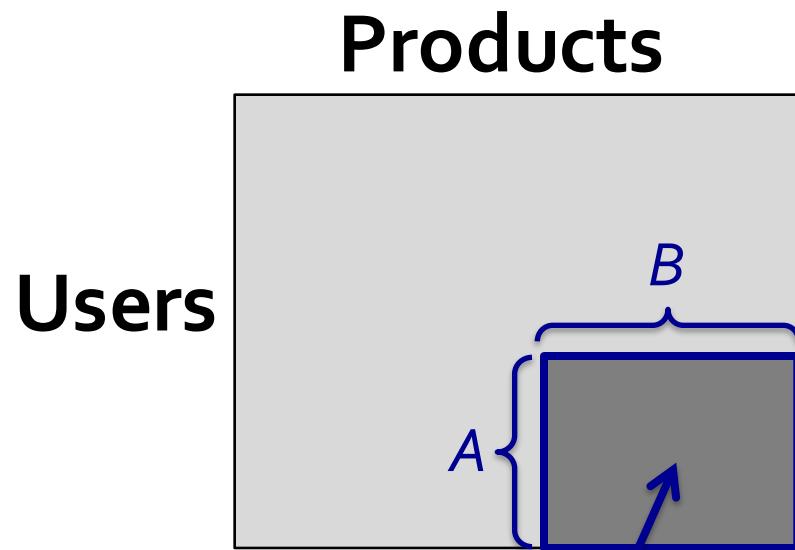
- Given:
 - Bipartite graph between users and products
 - (optional: prior node suspiciousness scores)
- Detect attacks in a way that is:
 - Camouflage-resistant
 - Near-linear time
 - Offers provable bounds
 - Works well in practice



Suspiciousness Metric



$$g(A, B) = 5$$



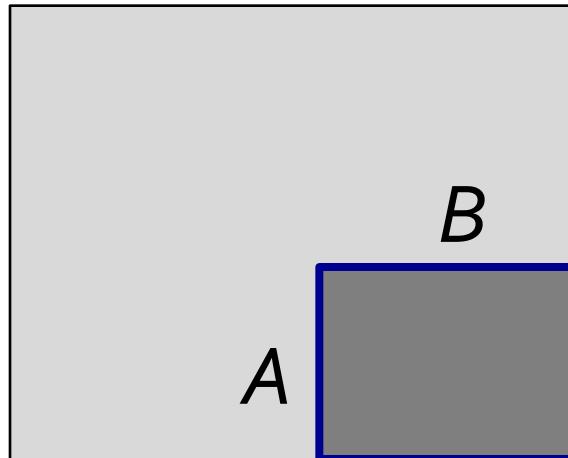
$$g(A, B) = 20$$

Camouflage-Resistance

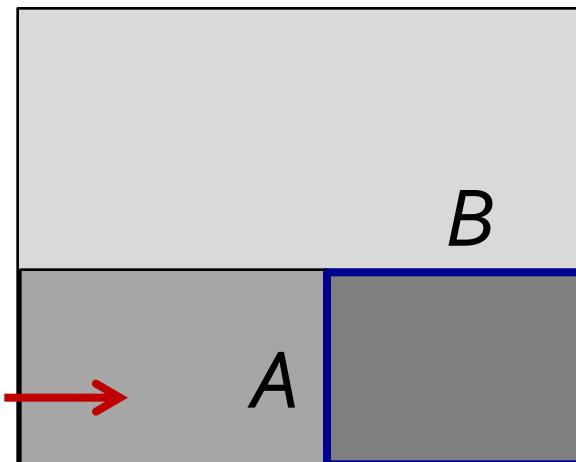
- g is *camouflage-resistant* if $g(A, B)$ does not decrease when camouflage is added to A .

Products

Users



Camouflage



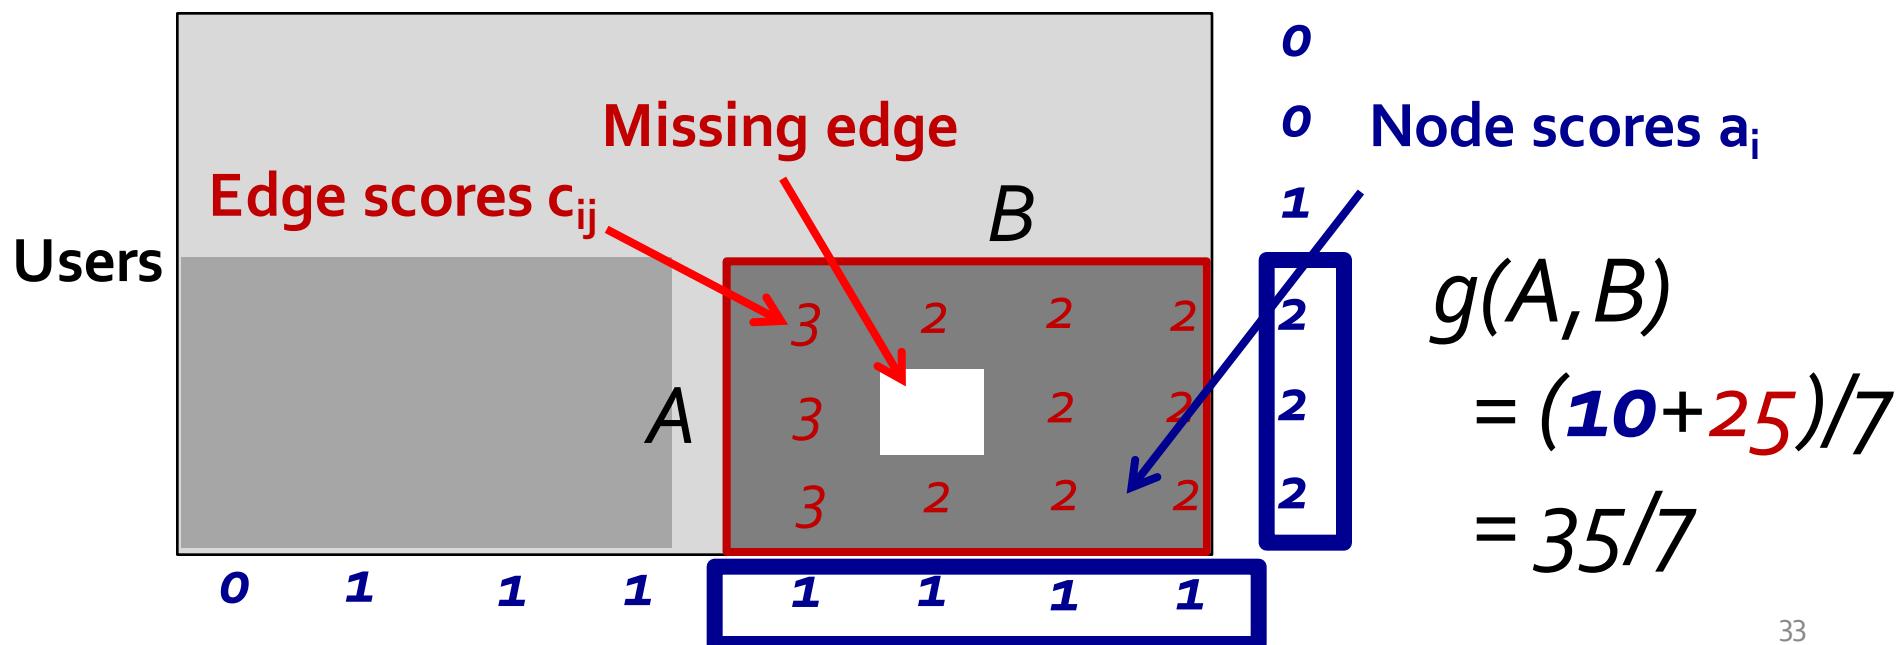
$$g(A, B) = 50 \longrightarrow g(A, B) = 50$$

Proposed Suspiciousness Metric

"Average suspiciousness" $g(A, B)$

$$= \frac{(\text{sum of node susp.}) + (\text{sum of edge susp.})}{|A| + |B|}$$

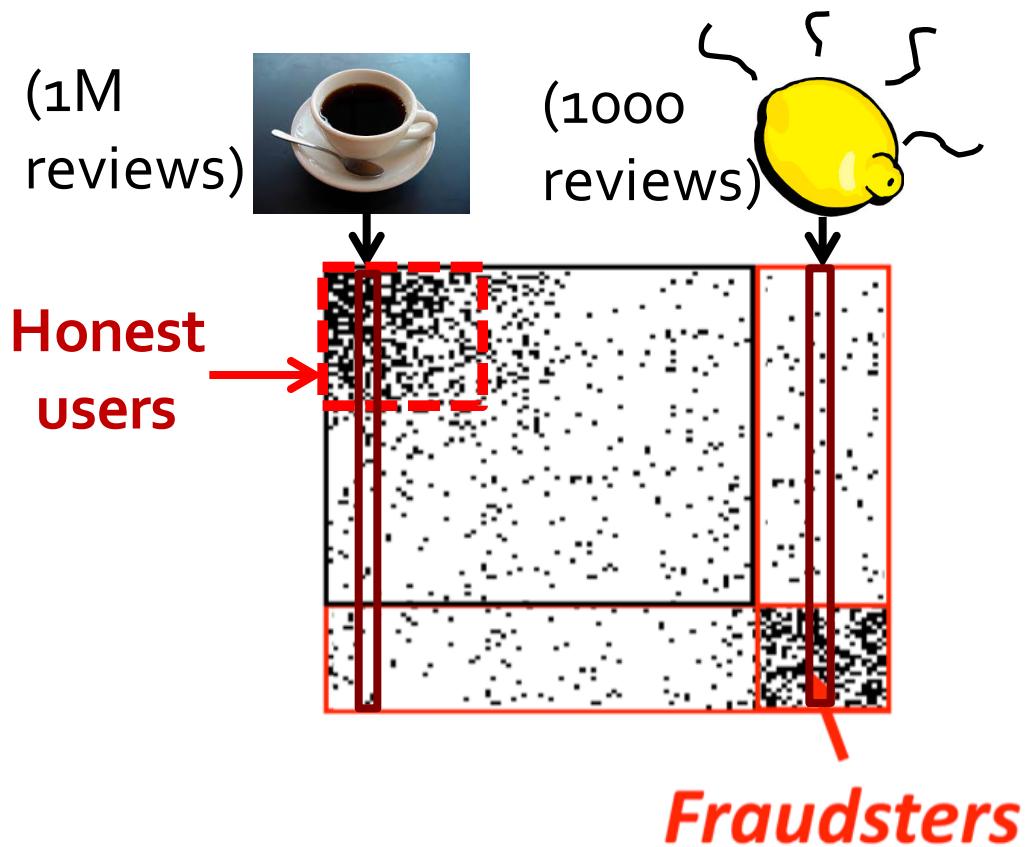
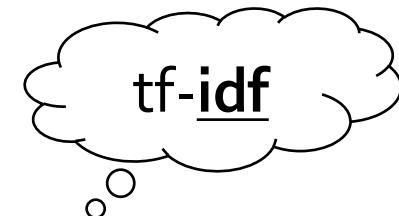
Products



Edge Scores c_{ij}

Proposed weighting scheme:

$$c_{ij} = 1 / \log(\text{unweighted sum of } j\text{-th column})$$



Why?

- Popular products are not necessarily suspicious
- Fraudulent products have a high fraction of edges from fraud

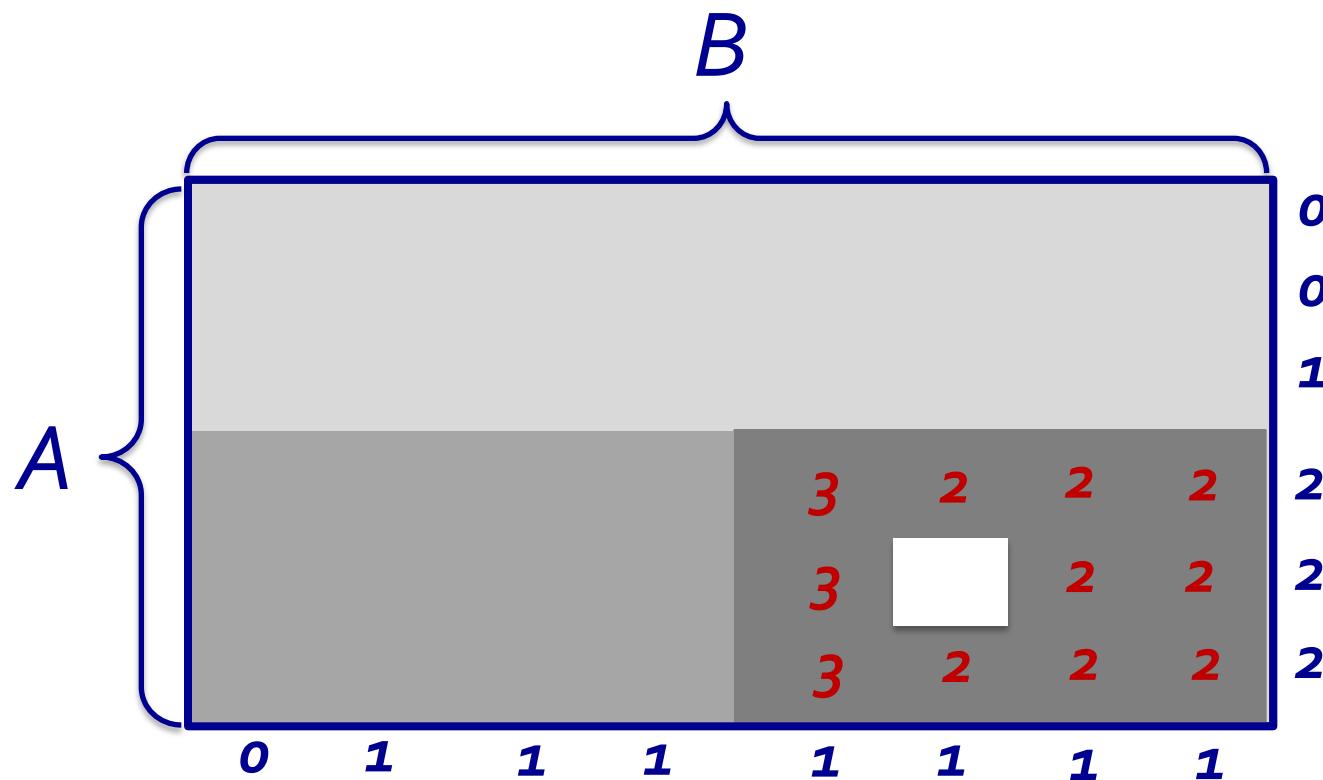
Metric Properties

Average suspiciousness $g(A, B)$:

- Can be optimized in near-linear time**
- Provable bounds
- Camouflage-resistant
- Works in practice

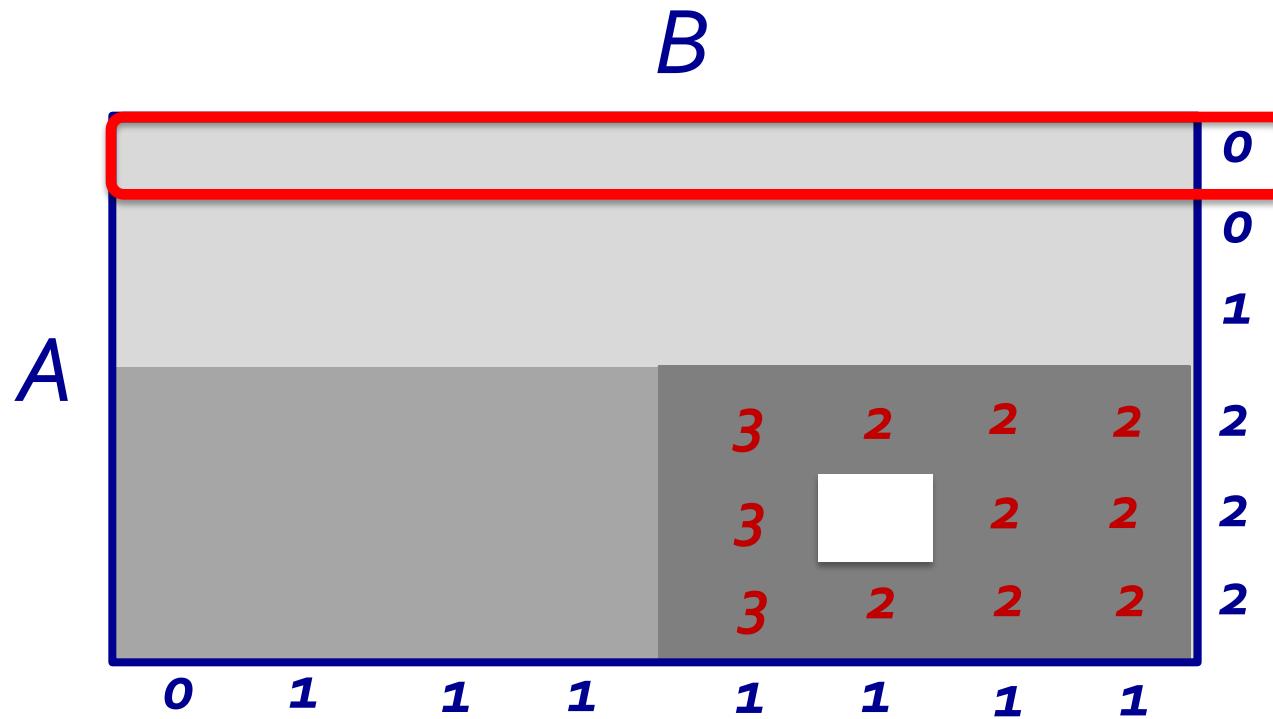
FRAUDAR: Greedy Algorithm

- Start with sets A, B as all users / products



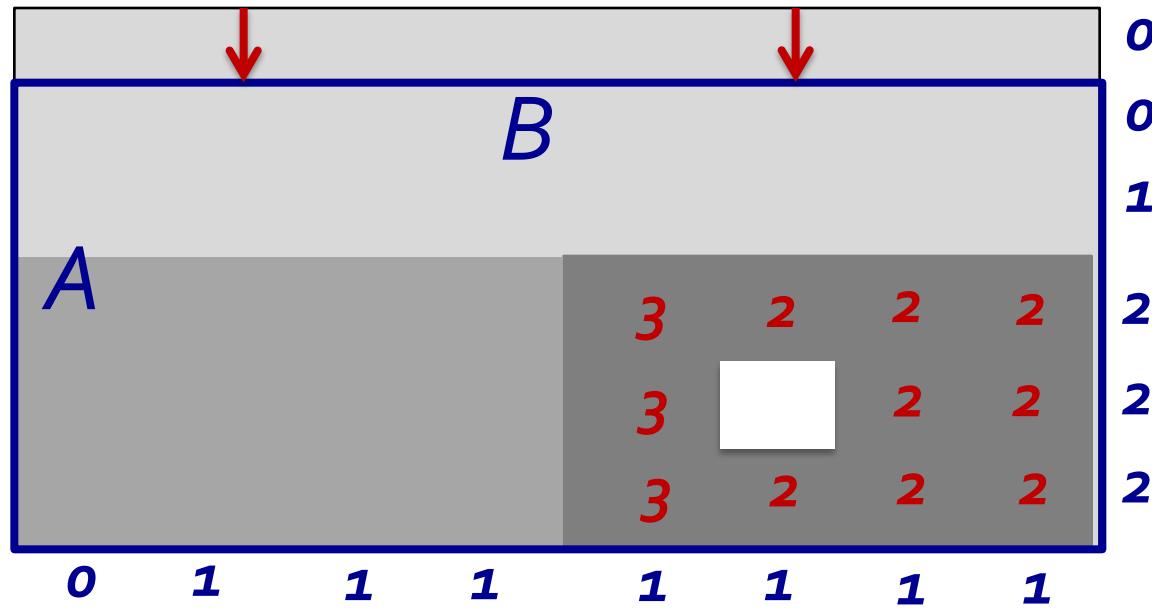
FRAUDAR: Greedy Algorithm (cont.)

- Delete rows / columns greedily to maximize g (average suspiciousness)



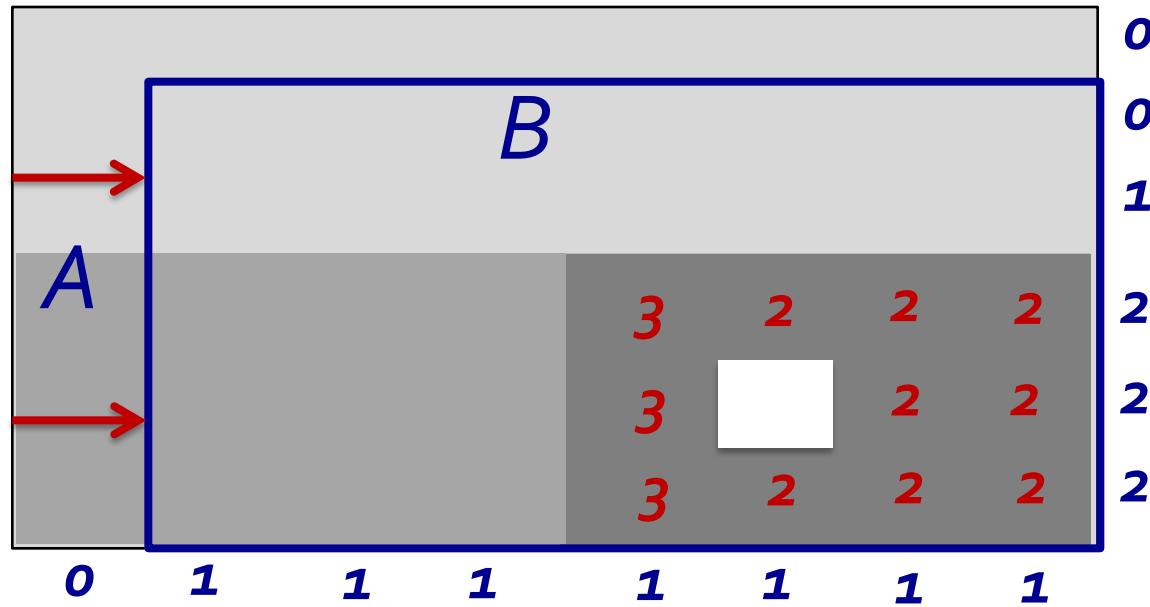
FRAUDAR: Greedy Algorithm (cont.)

- Delete rows / columns greedily to maximize g (average suspiciousness)



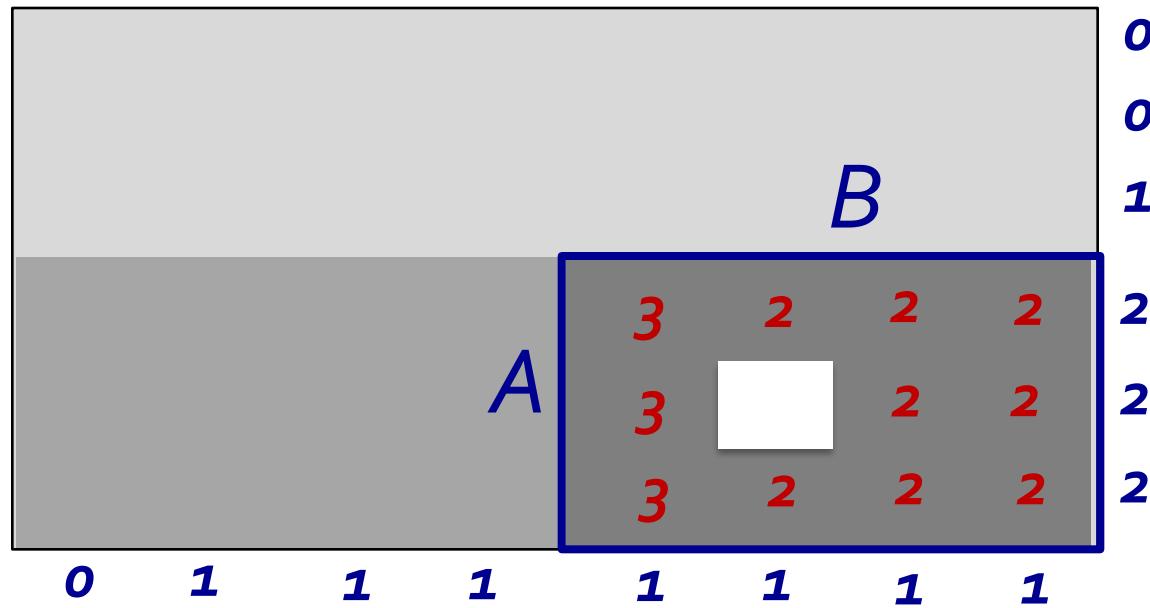
FRAUDAR: Greedy Algorithm (cont.)

- Delete rows / columns greedily to maximize g (average suspiciousness)



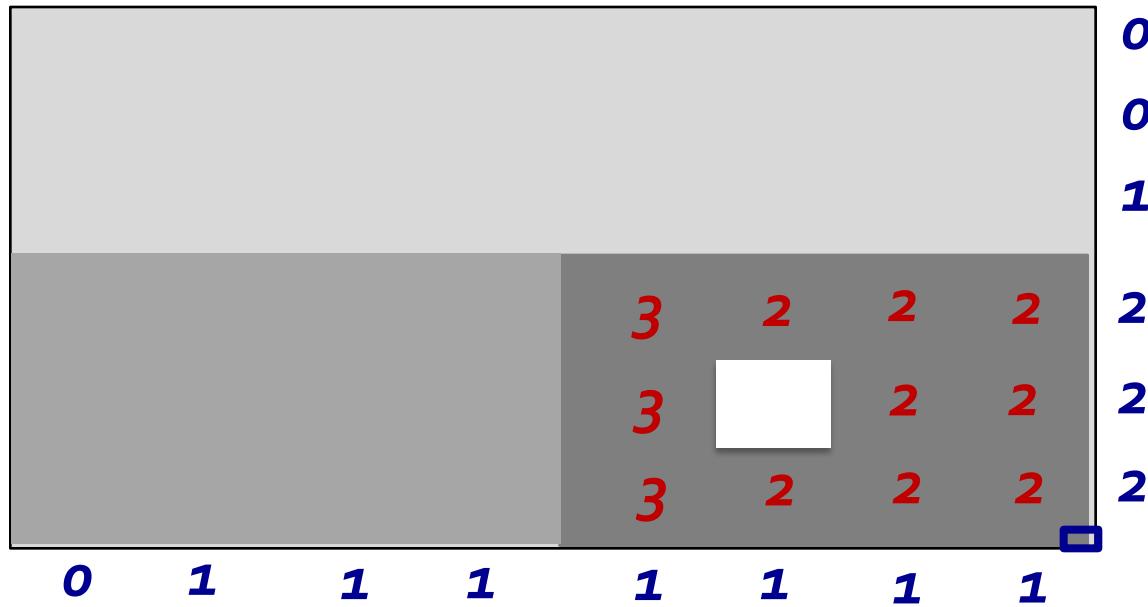
FRAUDAR: Greedy Algorithm (cont.)

- Delete rows / columns greedily to maximize g (average suspiciousness)



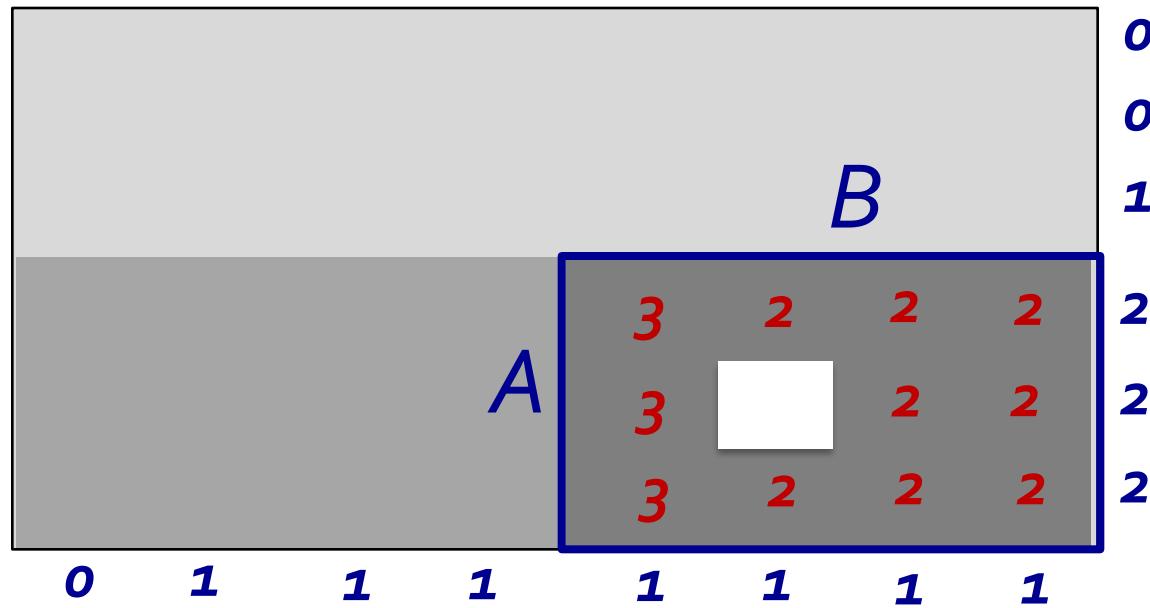
FRAUDAR: Greedy Algorithm (cont.)

- Continue until A and B are empty



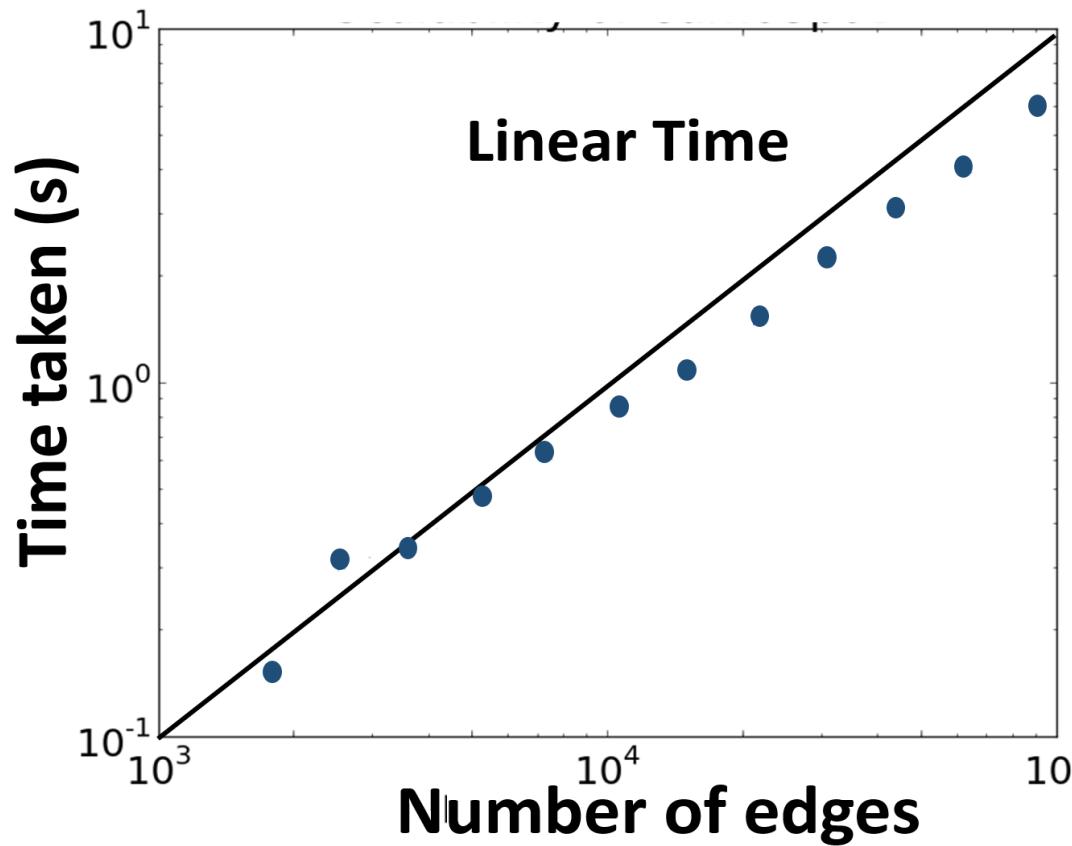
FRAUDAR: Greedy Algorithm (cont.)

- Return the best subsets A and B seen so far (based on g)



Computation Time

- $O(|E| \log(|V|))$: using appropriate data structures



Metric Properties

Average suspiciousness $g(A, B)$:

- Can be optimized in near-linear time
- Provable bounds**
- Camouflage-resistant
- Works in practice

Theoretical Guarantee

- **Theorem 1:** The subgraph (A, B) returned by FRAUDAR satisfies

$$g(\mathcal{A} \cup \mathcal{B}) \geq \frac{1}{2}g_{OPT}$$

FRAUDAR subgraph

Optimum value of g

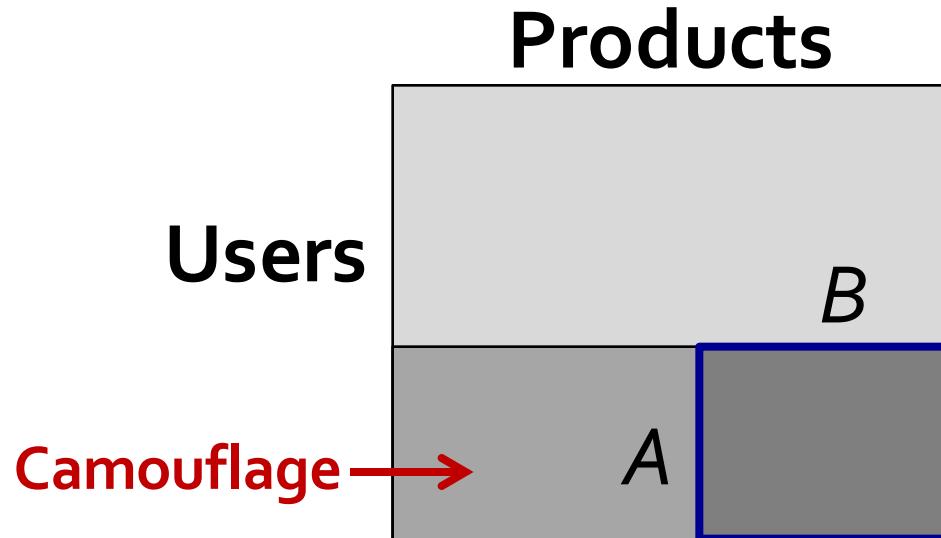
Metric Properties

Average suspiciousness $g(A, B)$:

- Can be optimized in near-linear time
- Provable bounds
- Camouflage-resistant**
- Works in practice

Camouflage Resistance

- **Theorem 2:** If c_{ij} is a *column weighting* (i.e. c_{ij} is any function of the j -th column), then \mathbf{g} is **camouflage-resistant**.



$(c_{ij} = 1 / \log(\text{sum of } j\text{-th column})$ satisfies this)

Metric Properties

Average suspiciousness $g(A, B)$:

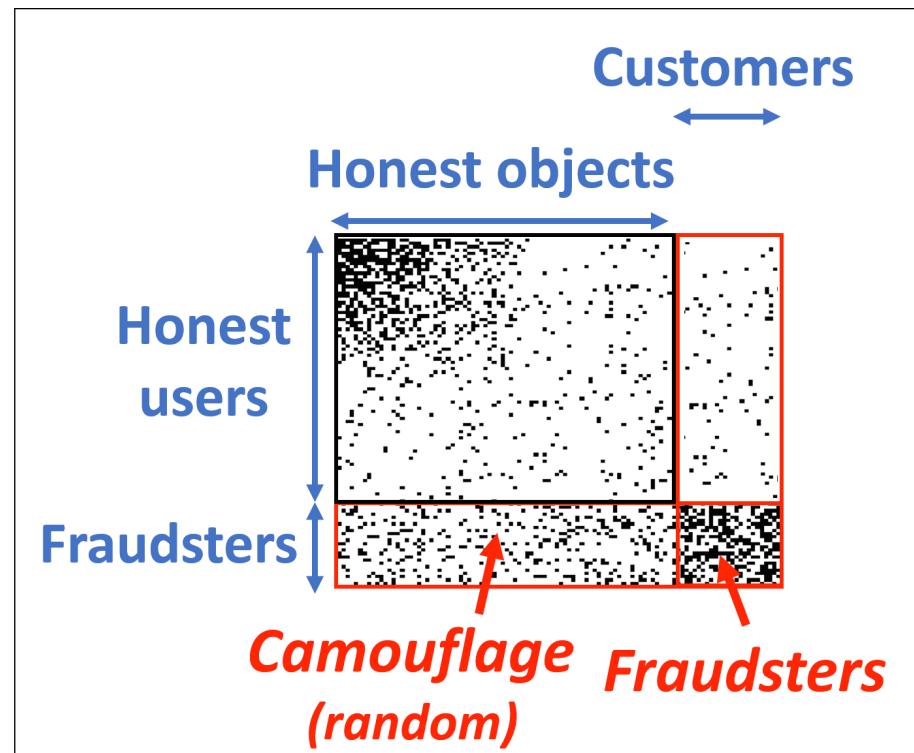
- Can be optimized in near-linear time
- Provable bounds
- Camouflage-resistant
- Works in practice**

Experiments: Detecting Injection of Various Types of “Camouflage”

- Amazon Review Graph:
24K users, 4K products
- Injected 200 x 200 blocks
with various types of
camouflage
 - None
 - Random camouflage
 - Biased camouflage
 - Hijacked accounts

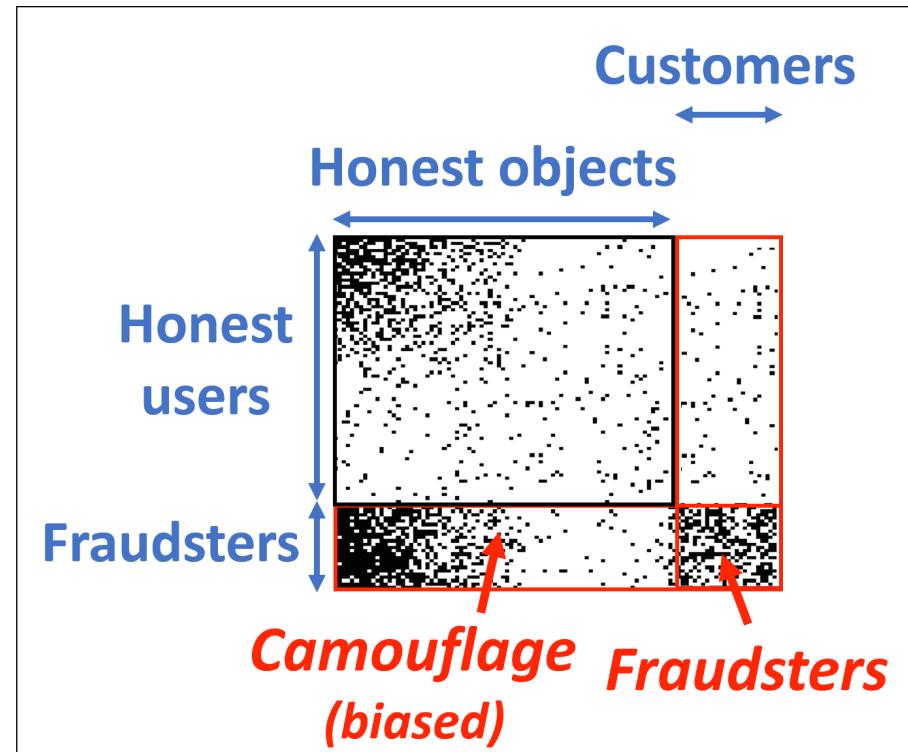
Experiments: Detecting Injection of Various Types of “Camouflage”

- Amazon Review Graph:
24K users, 4K products
- Injected 200 x 200 blocks
with various types of camouflage
 - None
 - **Random camouflage**
 - Biased camouflage
 - Hijacked accounts



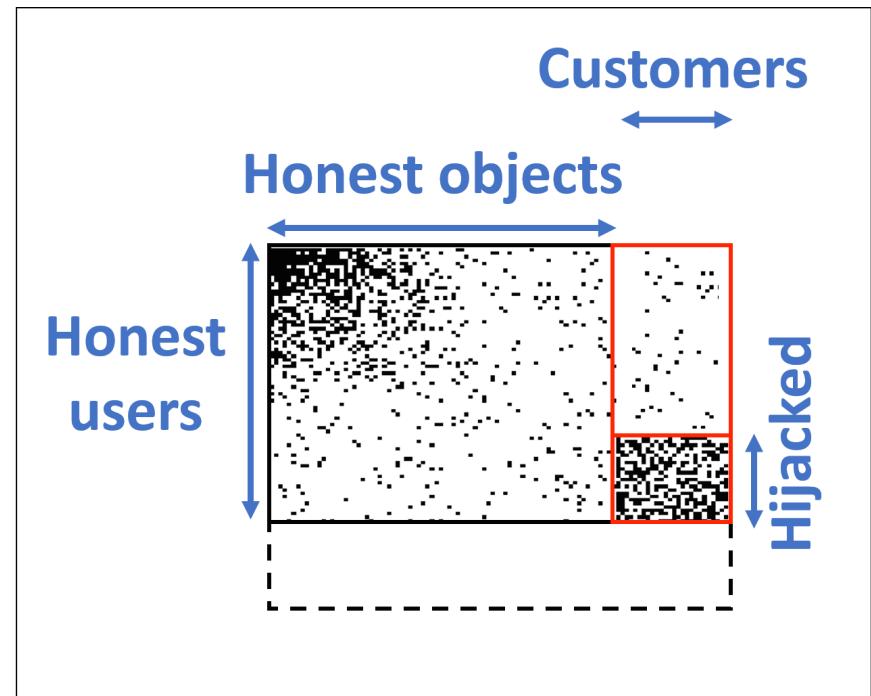
Experiments: Detecting Injection of Various Types of “Camouflage”

- Amazon Review Graph:
24K users, 4K products
- Injected 200 x 200 blocks
with various types of camouflage
 - None
 - Random camouflage
 - **Biased camouflage**
 - Hijacked accounts

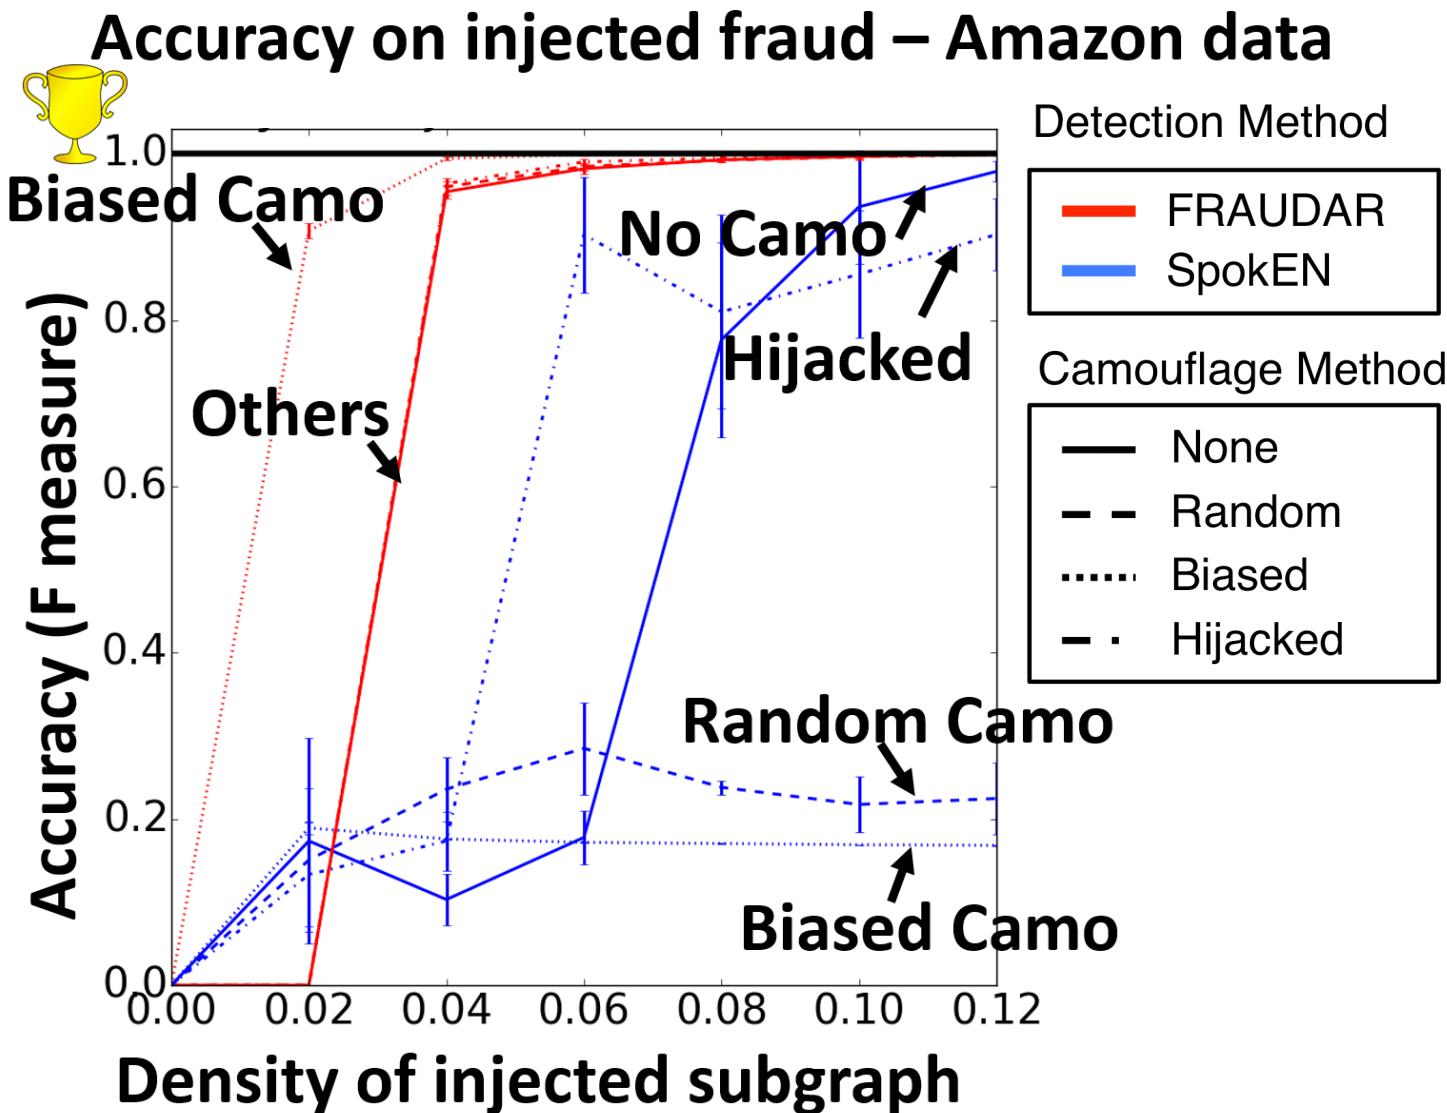


Experiments: Detecting Injection of Various Types of “Camouflage”

- Amazon Review Graph:
24K users, 4K products
- Injected 200 x 200 blocks
with various types of camouflage
 - None
 - Random camouflage
 - Biased camouflage
 - **Hijacked accounts**



Accuracy on Detecting Injected Fraud



Summary

- Spectral methods
 - Spectral clustering and community detection
- Spectral subspaces and spectral subspace plots
- **EigenSpokes** (singular vectors and “spokes”)
- **LockInfer** (“**camouflage**”, “fame”, “pearls”, “staircase”, etc.)
- **fBox** (small-scale, stealthy attacks; reconstructed degrees)
- **FRAUDAR** (theoretical guarantees for bounding graph fraud in the face of **camouflage**)
- Applications: Mobile calls, Twitter social network, “user-product” reviews

References

- Andrew Y. Ng, Michael I. Jordan, Yair Weiss. "On Spectral Clustering: Analysis and an algorithm", NIPS 2001.
- Scott White and Padhraic Smyth. "A spectral clustering approach to finding communities in graphs", SDM 2005.
- M. E. J. Newman. "Finding community structure in networks using the eigenvectors of matrices", Physical Review E 2006.
- Aditya Prakash, Mukund Seshadri, Ashwin Sridharan, Sridhar Machiraju, Christos Faloutsos. "EigenSpokes: Surprising Patterns and Scalable Community Chipping in Large Graphs", PAKDD 2010.
- Meng Jiang, Peng Cui, Alex Beutel, Christos Faloutsos, Shiqiang Yang. "Inferring lockstep behavior from connectivity pattern in large graphs", KAIS 2016.
- Neil Shah, Alex Beutel, Brian Gallagher, Christos Faloutsos. "Spotting suspicious link behavior with fBox: An adversarial perspective", ICDM 2014.
- Bryan Hooi, Hyun Ah Song, Alex Beutel, Neil Shah, Kijung Shin, Christos Faloutsos. "FRAUDAR: Bounding Graph Fraud in the Face of Camouflage", KDD 2016 Best Research Paper Award.



KDD 2017

Halifax, Nova Scotia - Canada

August 13 - 17, 2017

Tutorial: Data-Driven Approaches towards Malicious Behavior Modeling



Meng Jiang
University of Notre Dame



Srijan Kumar
Stanford University



Christos Faloutsos
Carnegie Mellon
University



V.S. Subrahmanian
University of Maryland,
College Park