

Tutorial: Data-Driven Approaches towards Malicious Behavior Modeling



Meng Jiang
University of Notre Dame



Srijan Kumar
Stanford University



Christos Faloutsos
Carnegie Mellon University



V.S. Subrahmanian
University of Maryland, College Park

Tutorial link: <http://bit.ly/kdd2017>

Outline

Introduction

Feature-based
algorithms

Spectral-based
algorithms

Density-based
algorithms

Sockpuppets

Vandals

Hoaxes

Visualization: “spokes”, “blocks”, “staircases”

Camouflage

Theoretical guarantee

Ill-gotten Likes

Synchronized Behaviors

Advertising campaigns

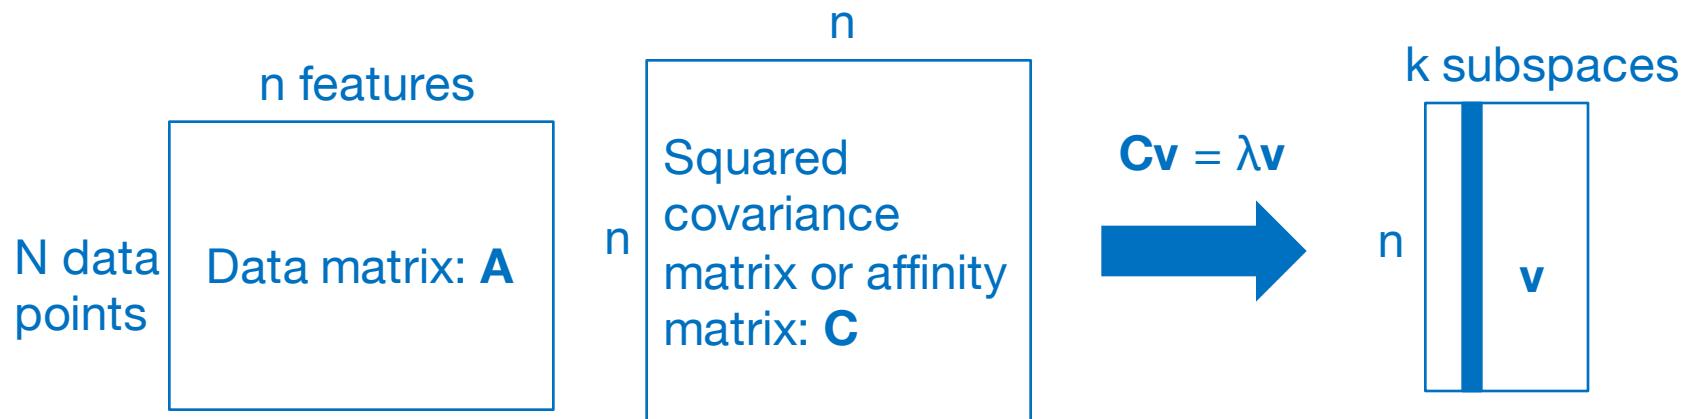
Social spam

Conclusions and future directions

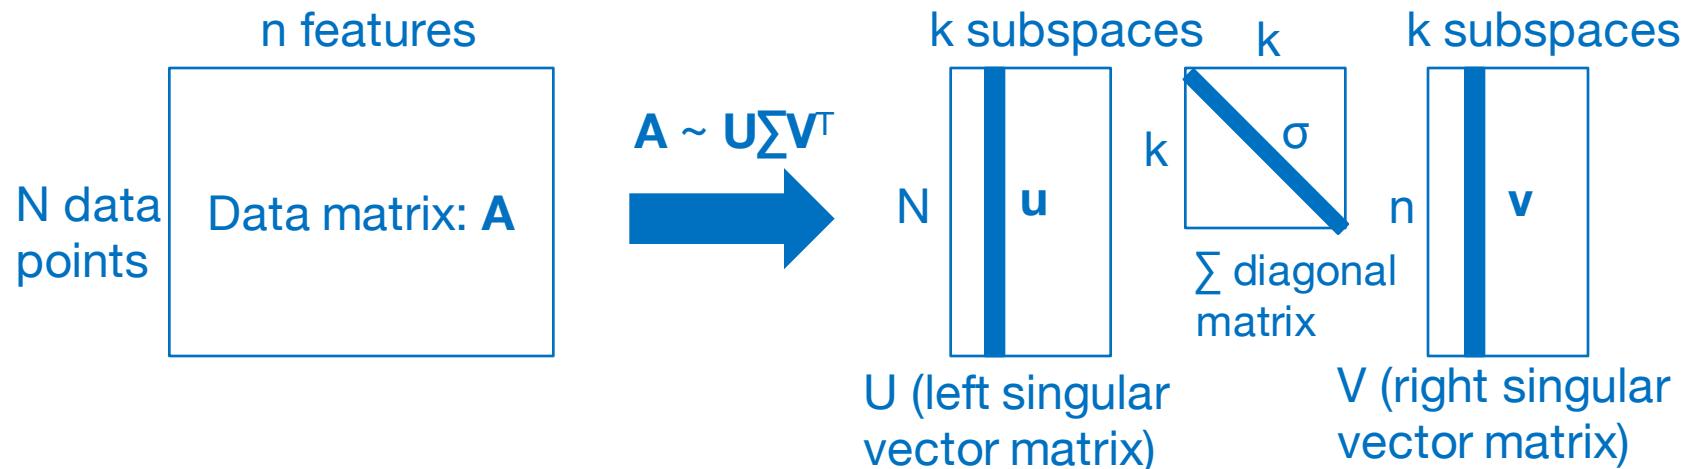
Tutorial link: <http://bit.ly/kdd2017>

How to get spectral subspaces?

- Frequency components → Principal Component Analysis (PCA): Eigenvectors



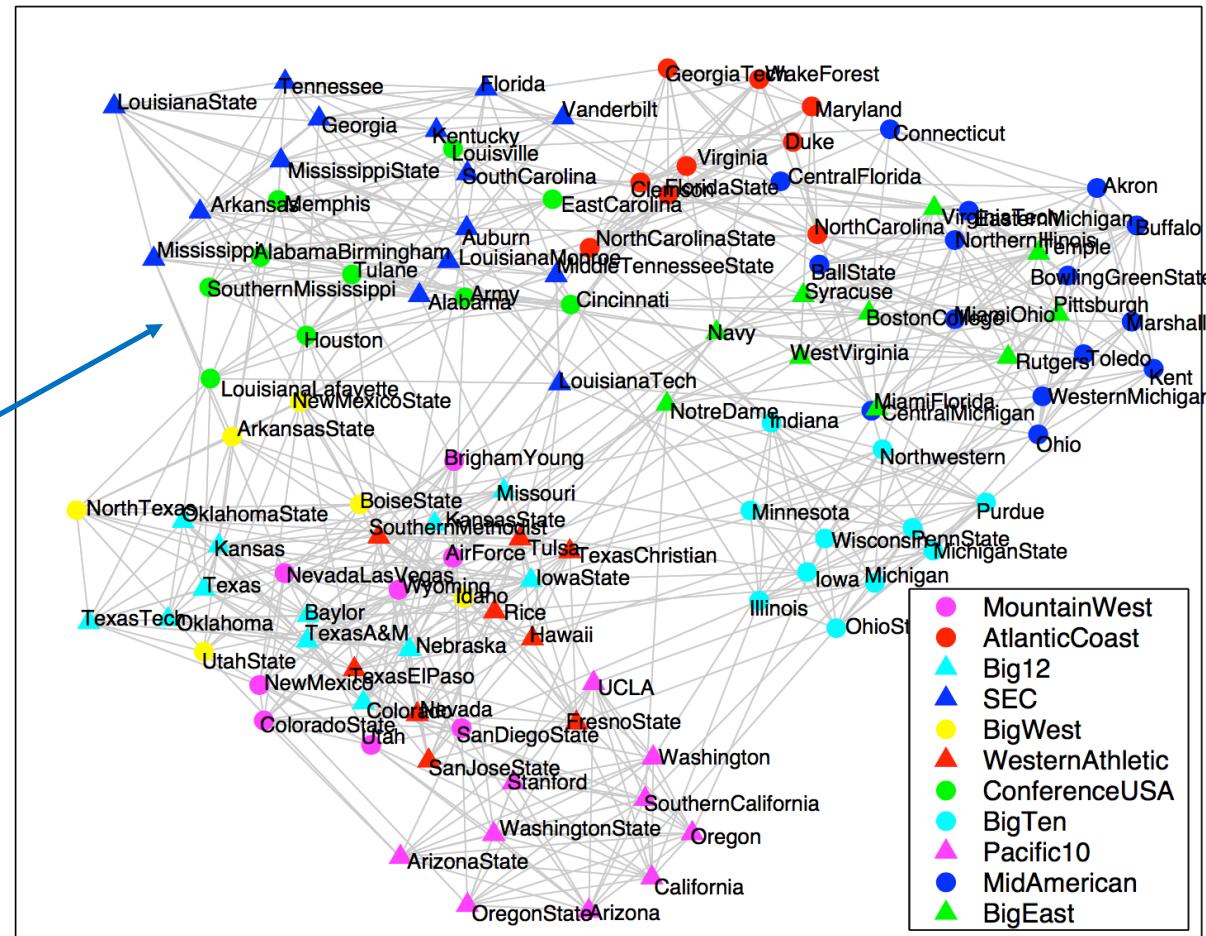
- Other spectral decomposition methods. Singular Value Decomposition (SVD): Singular vectors



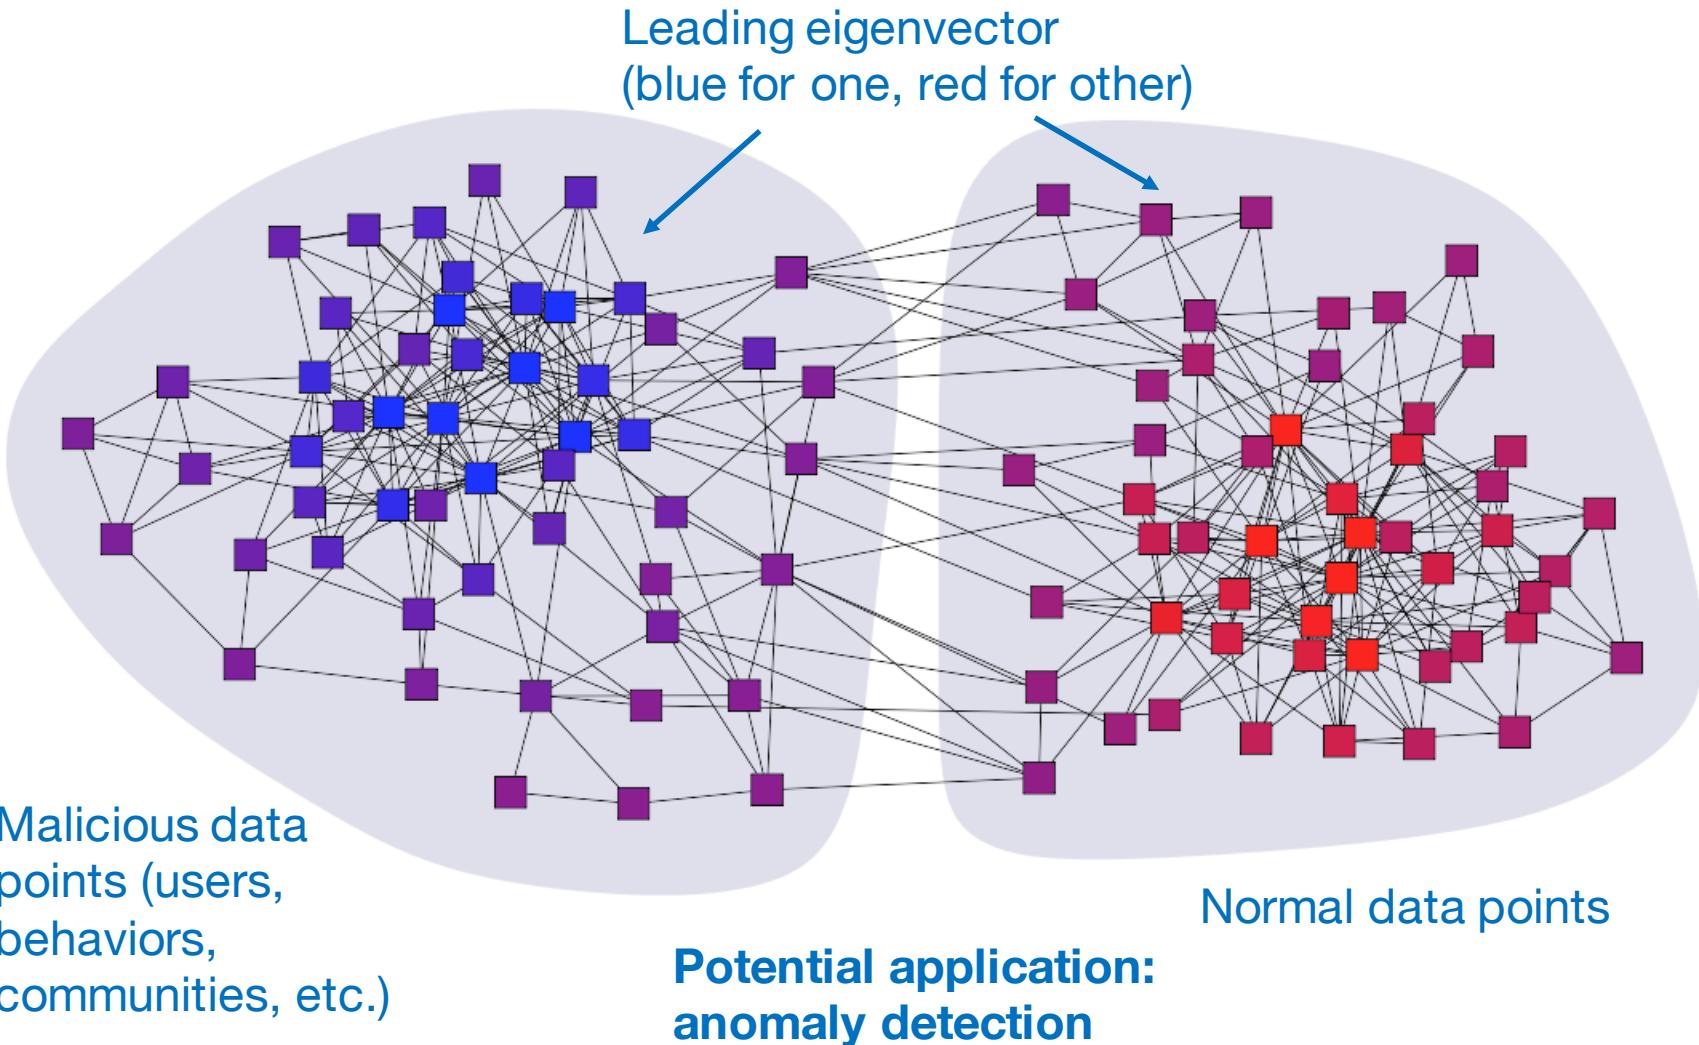
Spectral methods: Community Identification

Nodes are USA college football teams and edges represent which team played with which other team.

Communities represent groups of frequently co-playing teams.



Spectral methods: Anomaly detection



Spectral-based methods

- Advantages
 - Visualization: tunable value of k = number of subspaces
 - Feature extraction by data distribution rather than manual or automatic selection
 - “Principal” components represent “leading” vectors
 - Data: Can easily work with N -by- N graphs, N -by- N -by- N tensors
 - Applications: Finding communities and anomalies
- Disadvantages
 - Lack of interpretability of the subspaces/features

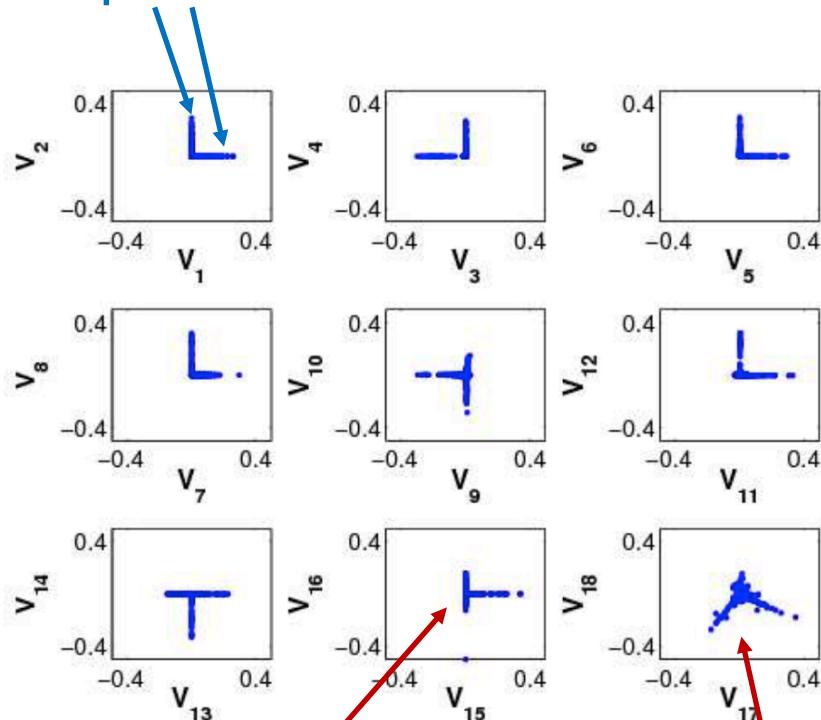
Finding Surprising Spectral Patterns in Large Graphs

- **Problem definition:** Given a social graph based on mobile calls made from/to callers, find caller communities.
- Dataset: Activity over the duration of a month, 186,000 nodes and 464,000 edges.
- Key contribution: Discovery of the “spokes” phenomenon
 - **The singular vectors of the graph, when plotted against each other, often have clear separate lines, typically aligned with axes.**
 - Use EigenEigen (EE) plots to identify communities in the form of **cliques or near-cliques, perfect or near-perfect bipartite-cores**.

Spokes and Dense Cliques

V-V plots: Right spectral subspaces

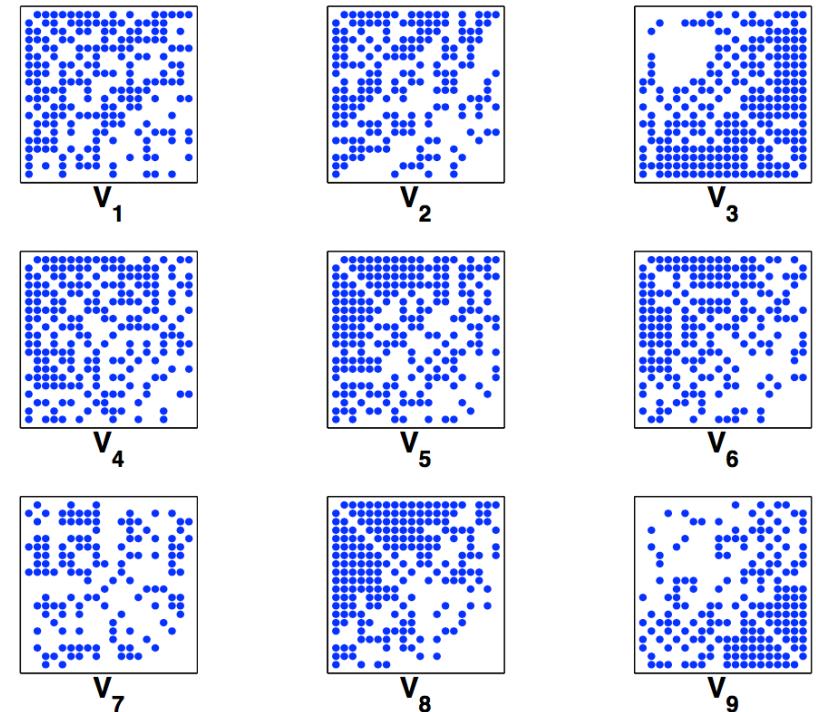
Spokes



(a) *EE-plot*

Long vs. Short?

Subgraphs of nodes identified as important by different V vectors: all are tightly knit



(b) Spy Plots of Sub-graph of Top 20 Nodes

Tilted?

Spectral subspace

- What is the meaning of spokes, elongated spokes, tilted spokes?
- Are there other patterns?
- Can these patterns be used to identify malicious behavior?

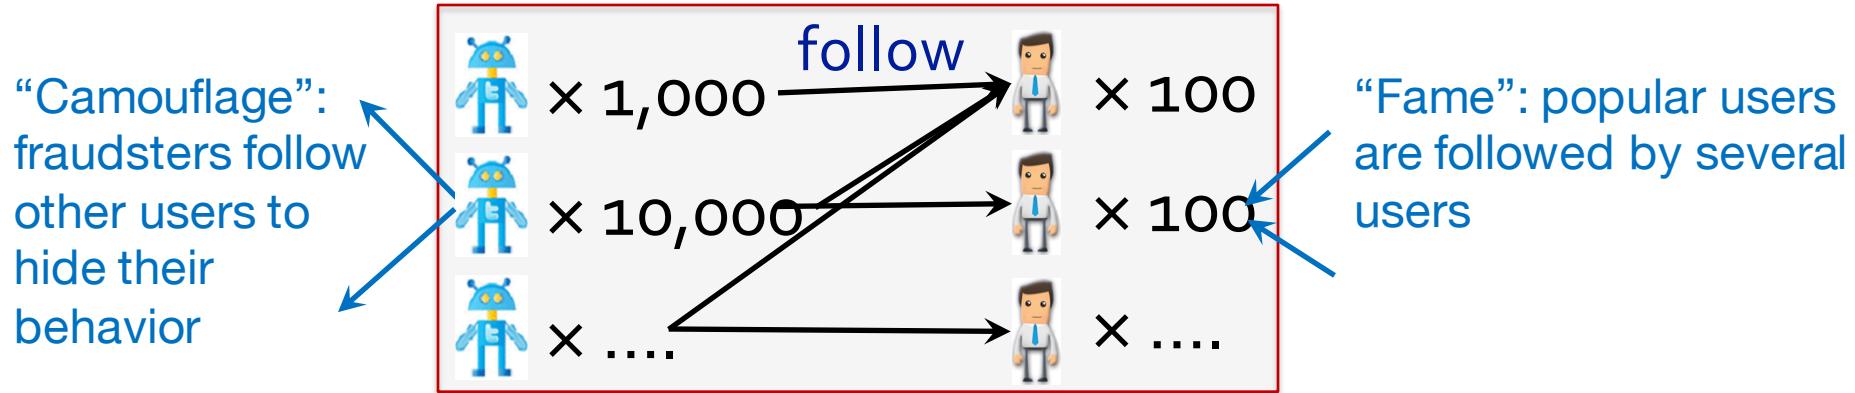
Inferring Lockstep Behavior from Connectivity Patterns

- Problem definition: Given a large graph, from **spectral subspace plots**, can we infer **lockstep behavior** patterns?



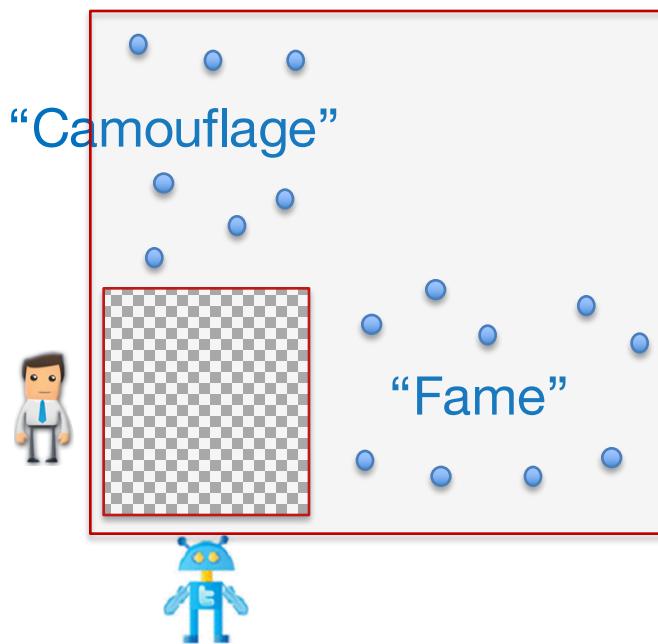
Inferring Lockstep Behavior from Connectivity Patterns

- Problem definition: Given a large graph, from **spectral subspace plots**, can we infer **lockstep behavior patterns**?



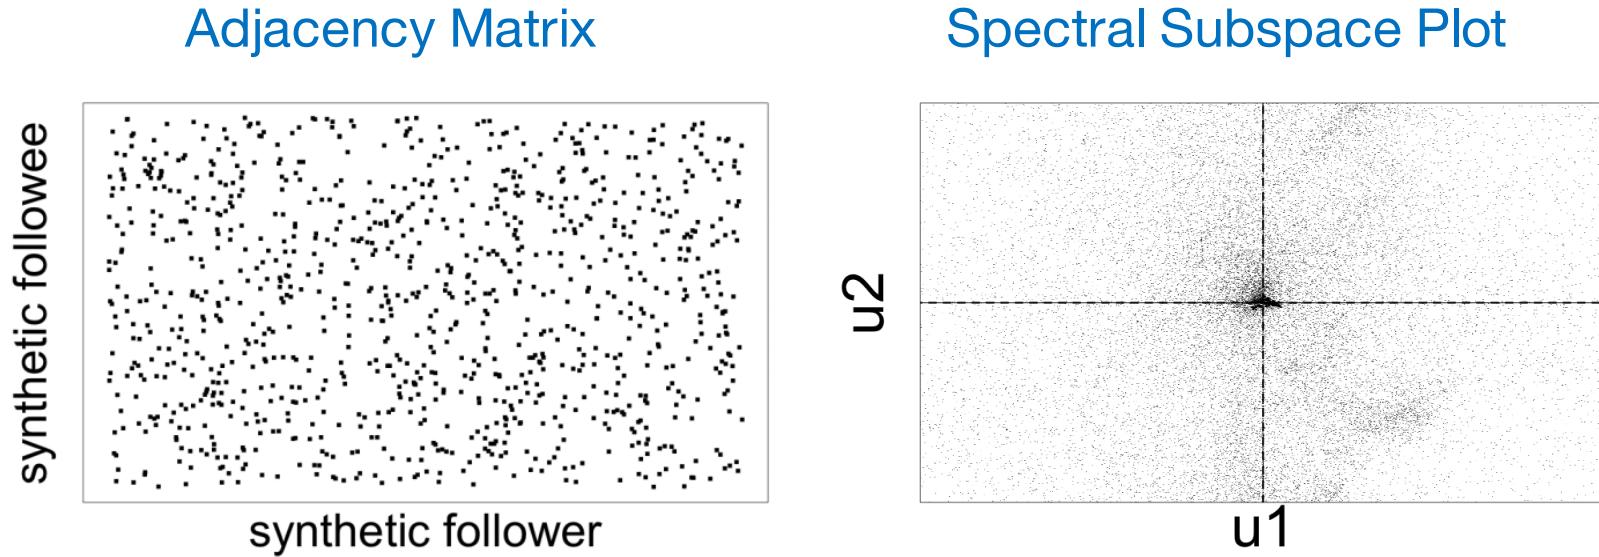
Inferring Lockstep Behavior from Connectivity Patterns

- Problem definition: Given a large graph, from **spectral subspace plots**, can we infer **lockstep behavior** patterns?



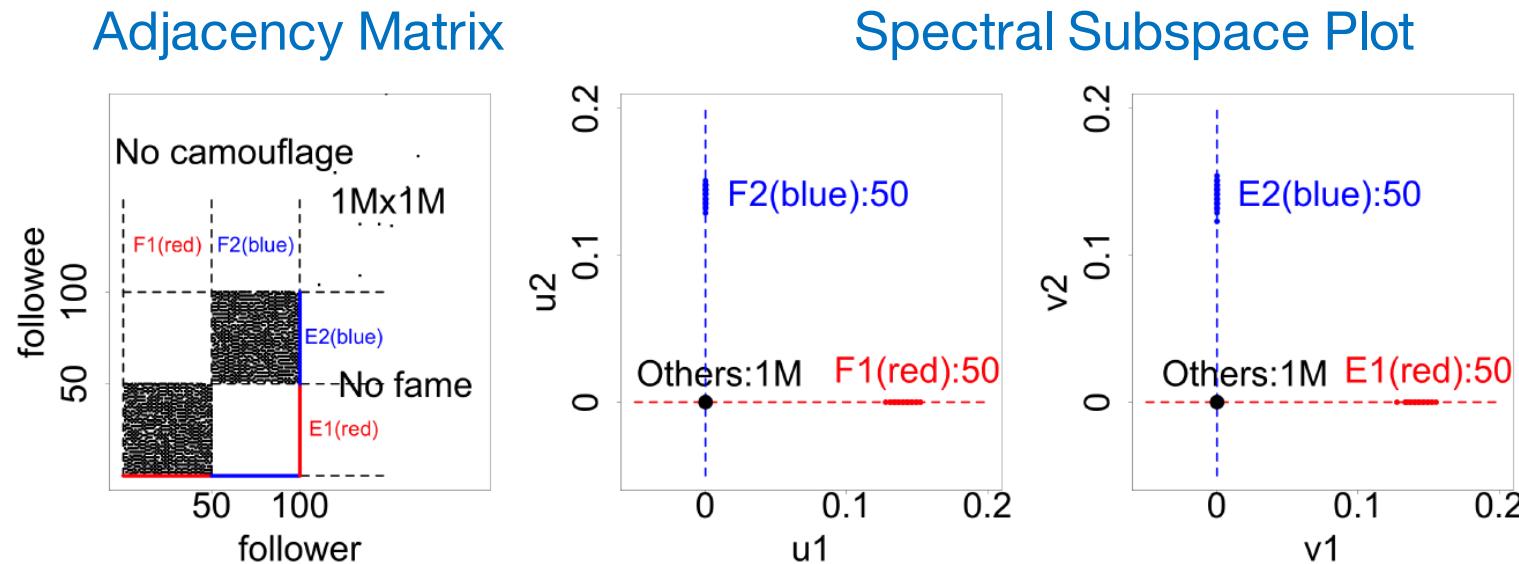
Case 0: No lockstep behavior

- No blocks in adjacency matrix lead to scattering and no patterns in spectral subspace



Case 1: Non-overlapping dense lockstep

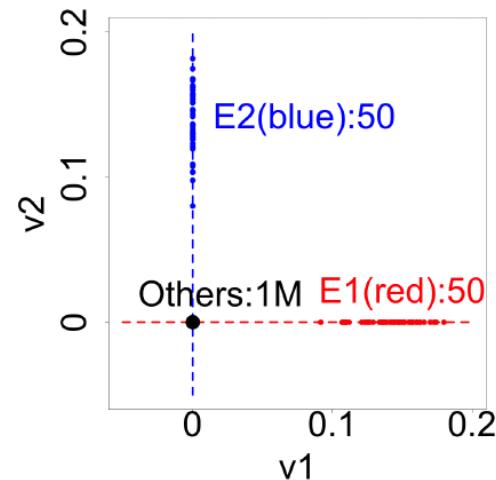
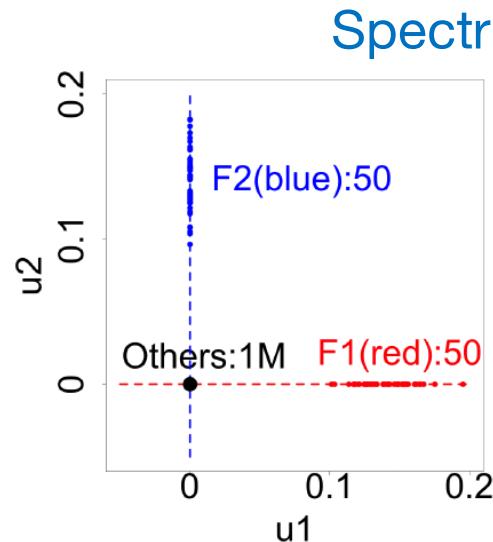
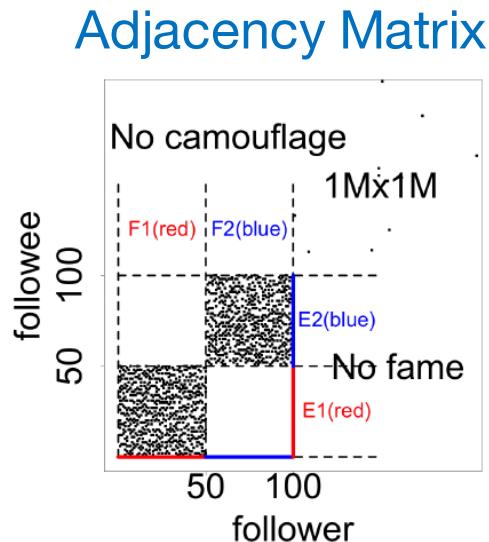
- Dense blocks in adjacency matrix generates “rays” in spectral subspace



Rule 1 (short “rays”): two blocks, high density (90%), no “camouflage”, no “fame”

Case 2: Non-overlapping sparse lockstep

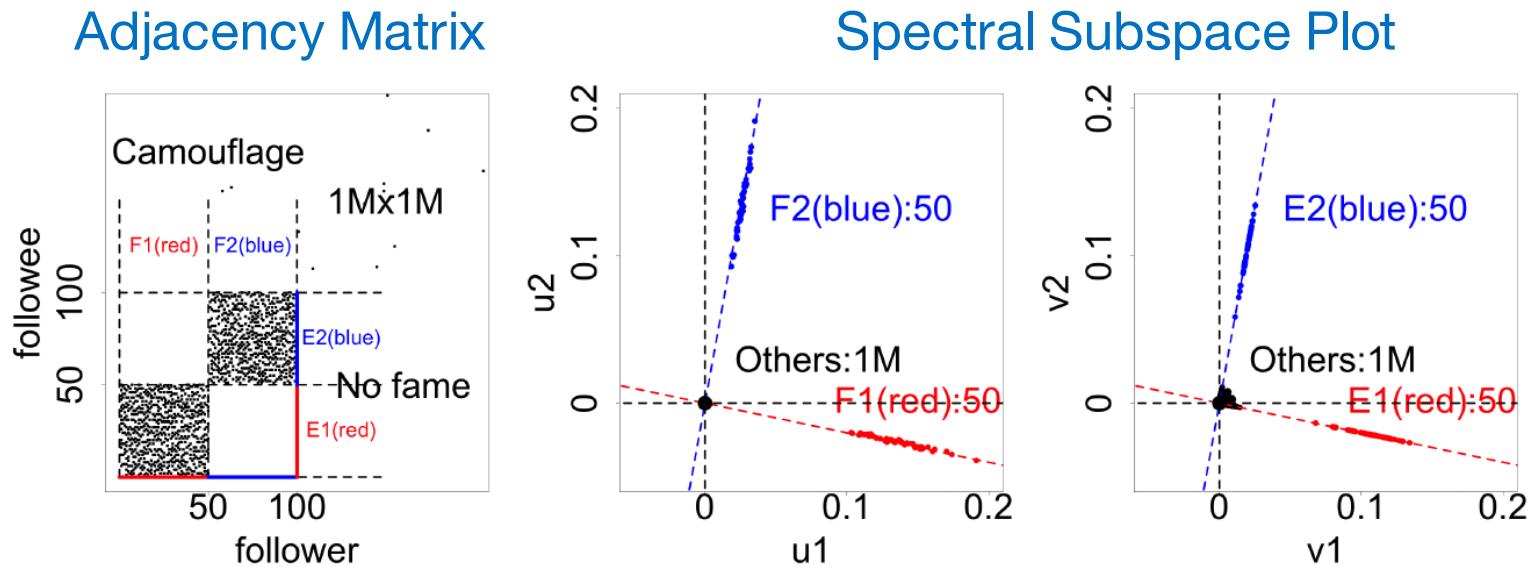
- Low density blocks in adjacency matrix leads to elongation of rays, indicating more varied behavior



Rule 2 (long “rays”): two blocks, low density (50%), no “camouflage”, no “fame”

Case 3: Non-overlapping lockstep with outside edges

- Edges to or from blocks in adjacency matrix leads to tilting of rays in spectral subspace
- Edges going out of block: “camouflage” by fraudsters
- Edges into the block: “fame” edges to popular users

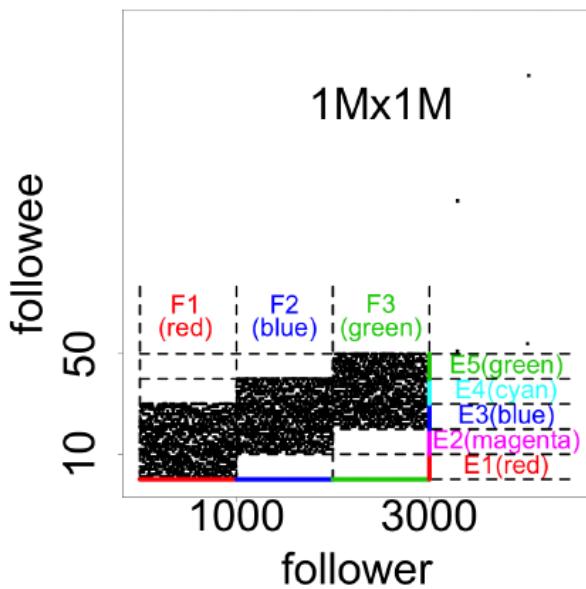


Rule 3 (tilting “rays”): two blocks, with “camouflage”, no “fame”

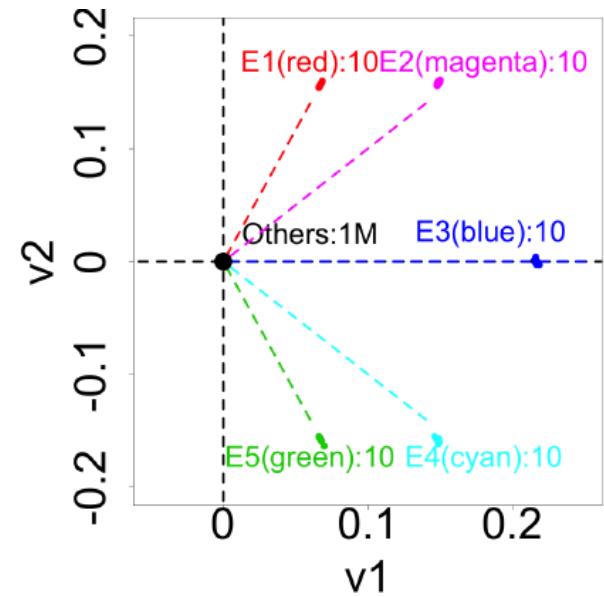
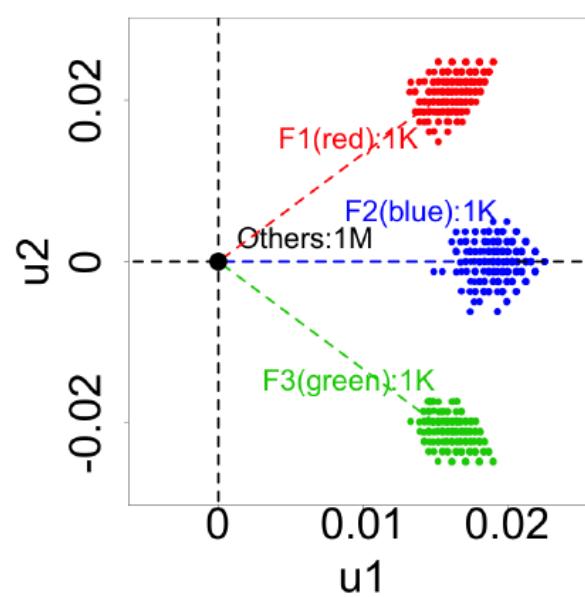
Case 4: Overlapping lockstep

- “Staircase”, i.e. sequentially overlapping blocks, in adjacency matrix generates “Pearls” in spectral subspace

Adjacency Matrix
with staircase



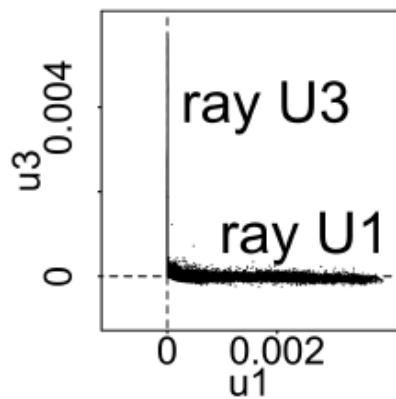
Spectral subspace plot
showing pearls



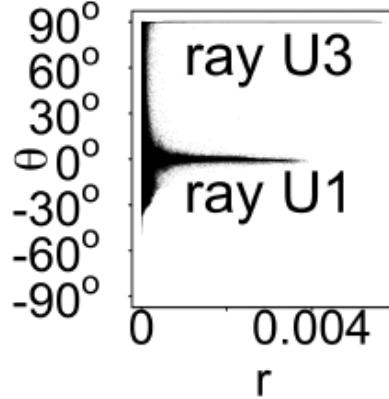
Rule 4 (“pearls”): a “staircase” of three partially overlapping blocks.

LockInfer Algorithm: Reading Spectral Subspace Plots

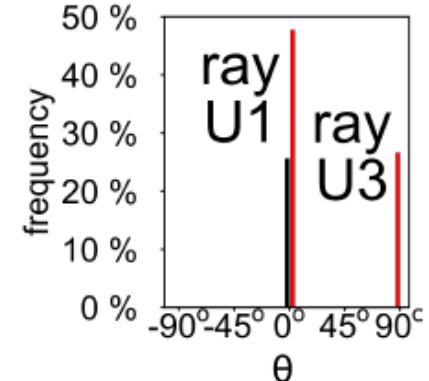
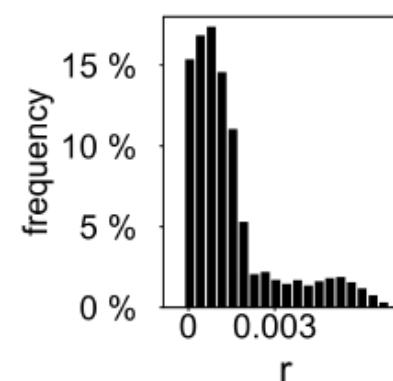
Spectral
Subspace Plot



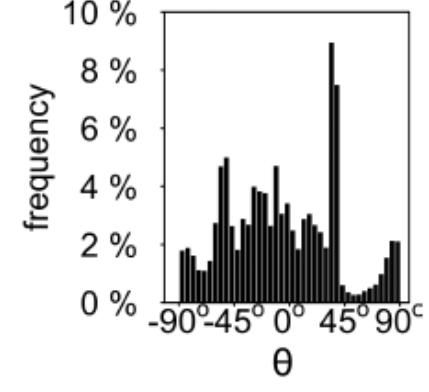
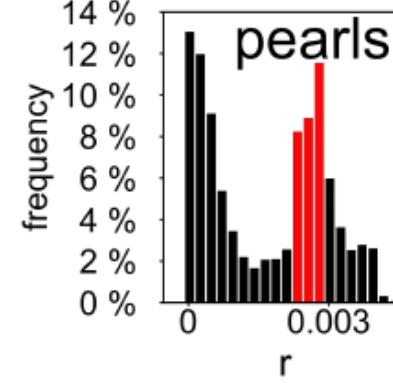
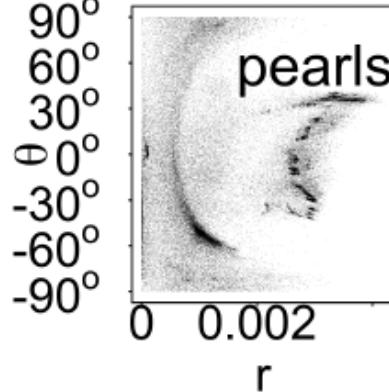
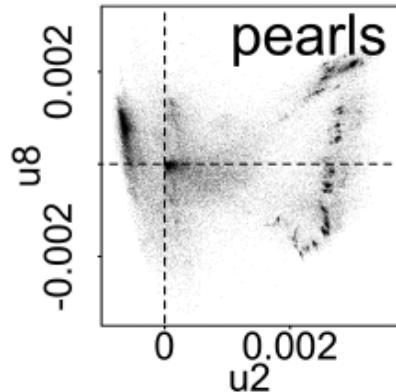
Polar Coordinate
Transform



Histograms



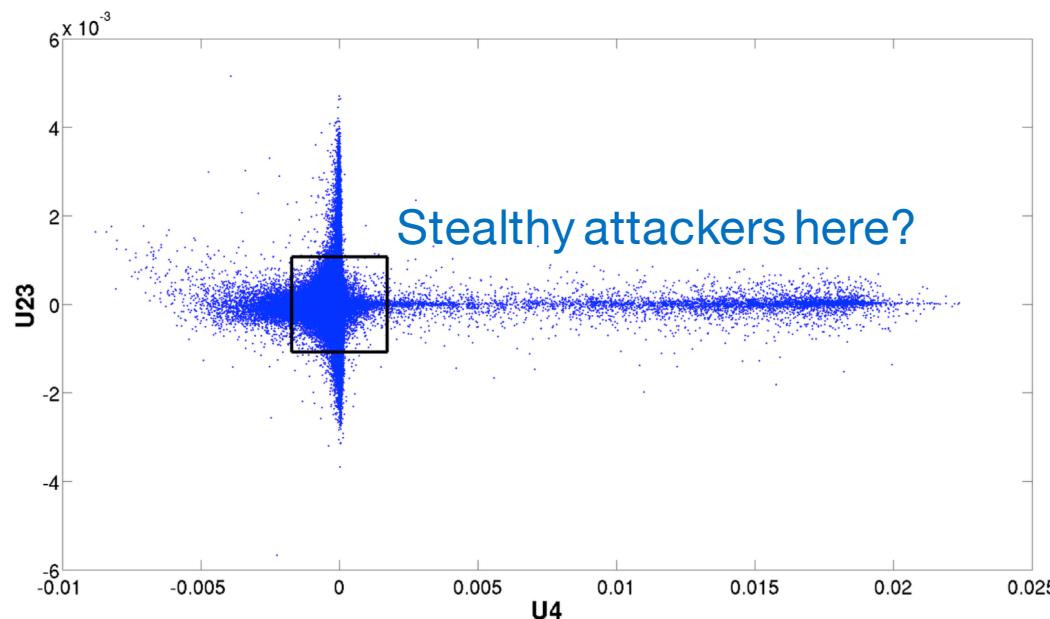
“rays” show two apparent spikes on θ frequency at 0° and 90°



“pearls” show a spike on r frequency at a much-greater-than-zero value

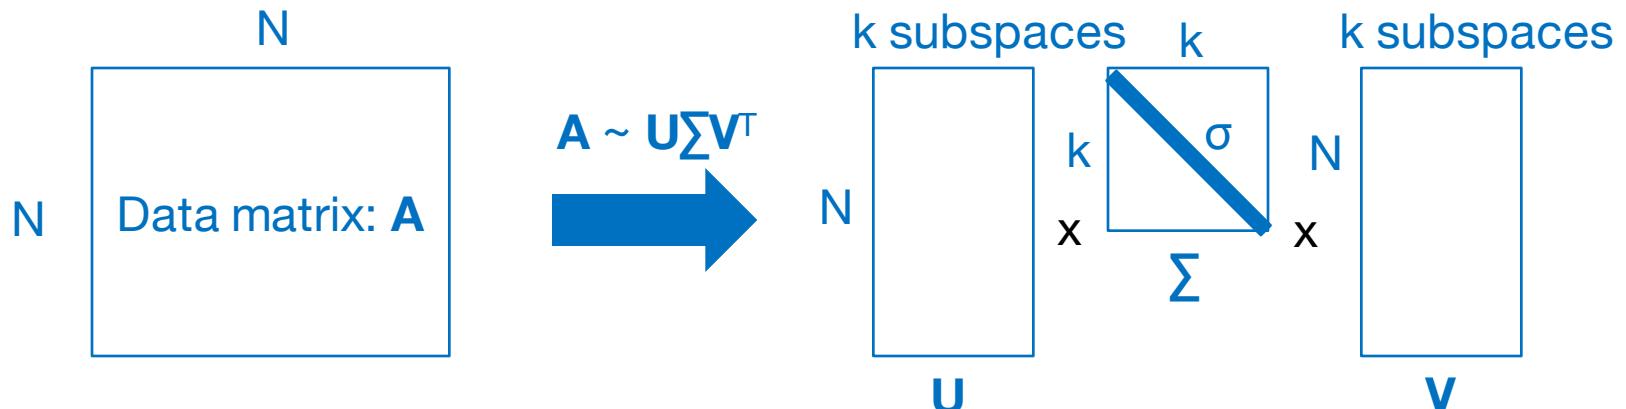
Spotting Small-Scale, Stealthy Attacks

- **Problem definition:** Can we catch stealthy attacks that are missed by traditional spectral methods?
- Dataset: Twitter “who-follows-whom” social graph, 41.7 million nodes, 1.5 billion edges



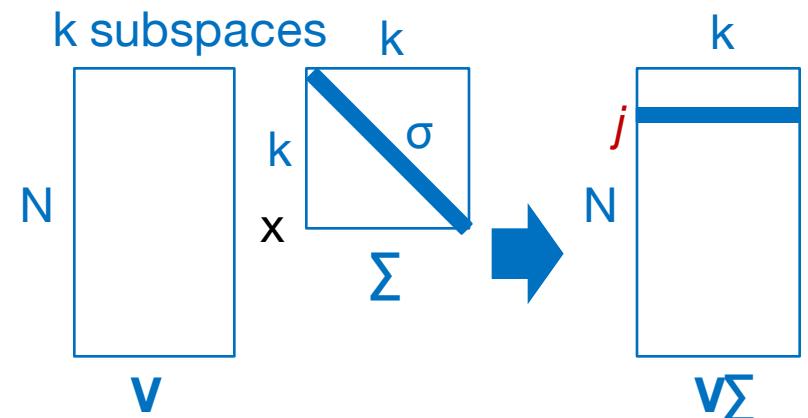
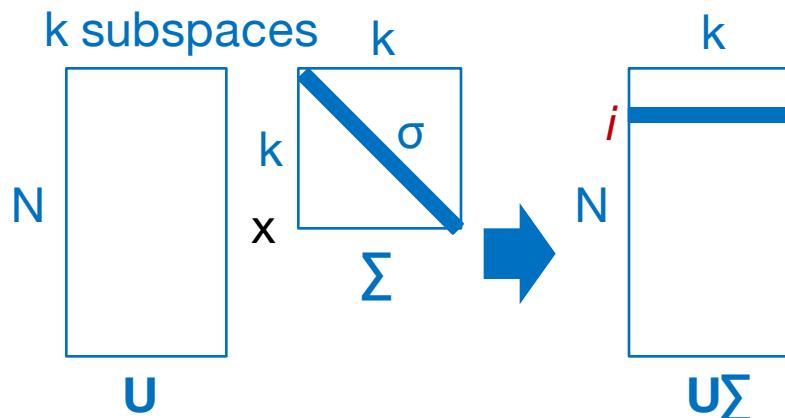
fBox: Reconstructed Degrees

SVD



$$\text{Reconstructed out-degree}(i) = \|(\mathbf{U}\Sigma)_i\|_2$$

$$\text{Reconstructed in-degree}(j) = \|(\mathbf{V}\Sigma)_j\|_2$$



Norm-Preserving Property of SVD

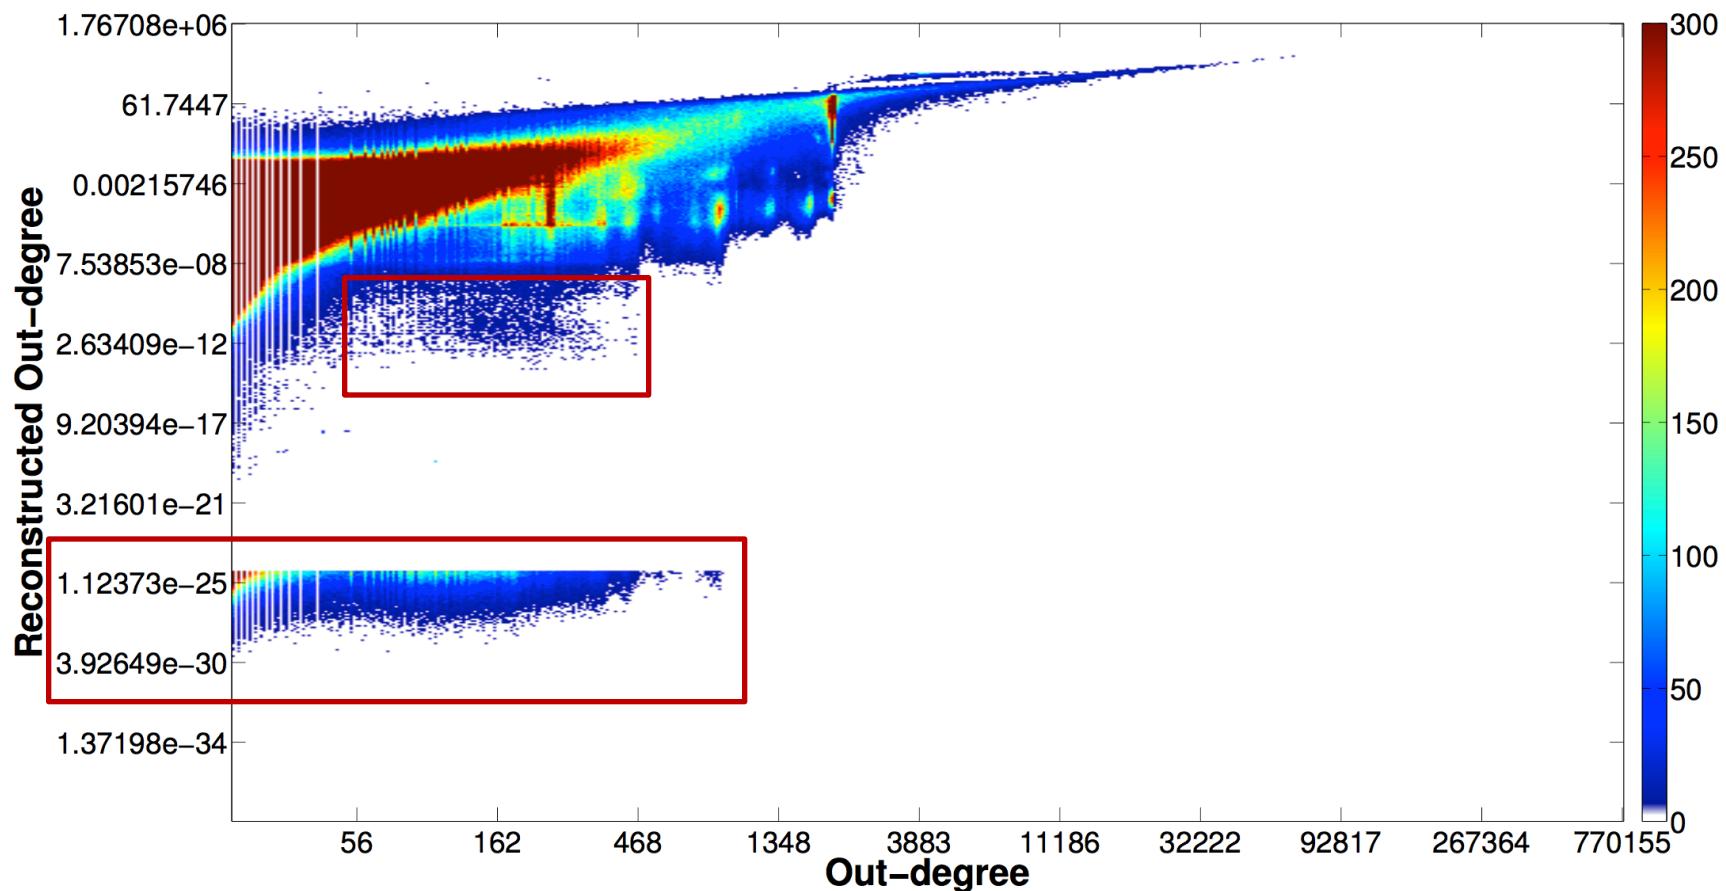
- The row vectors of a full rank decomposition and associated projection will retain the same L_2 norm or vector length as in the original space:
 - For $k = \text{rank}(\mathbf{A})$, $\|\mathbf{A}_i\|_2 = \|(\mathbf{U}\Sigma)_i\|_2$ and $\|\mathbf{A}^T j\|_2 = \|(\mathbf{V}\Sigma)_j\|_2$
- So, compare:
 - Reconstructed out-degree vs. real out-degree
 - Reconstructed in-degree vs. real in-degree

Why does fBox work?

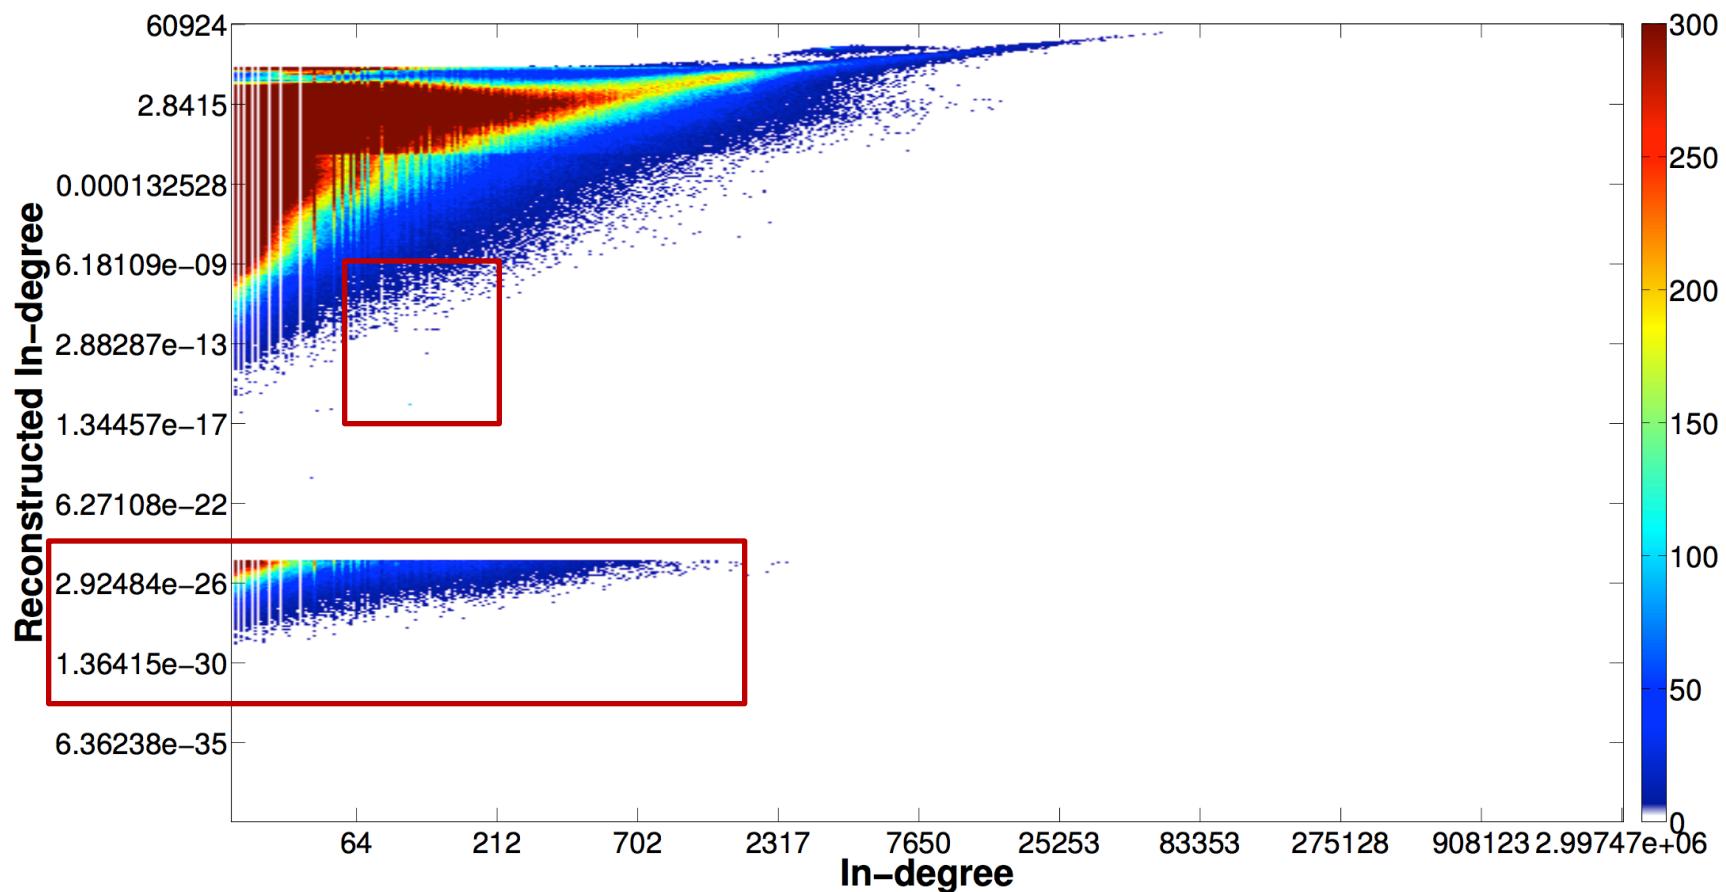
For $k < \text{rank}(A)$, dishonest users' reconstruction is poor compared to that of honest users.

- Dishonest users who either form isolated components or link to dishonest objects will project poorly and have characteristically low reconstruction degrees
- Honest users who are well-connected to real products and brands should project more strongly and have characteristically higher reconstruction degrees

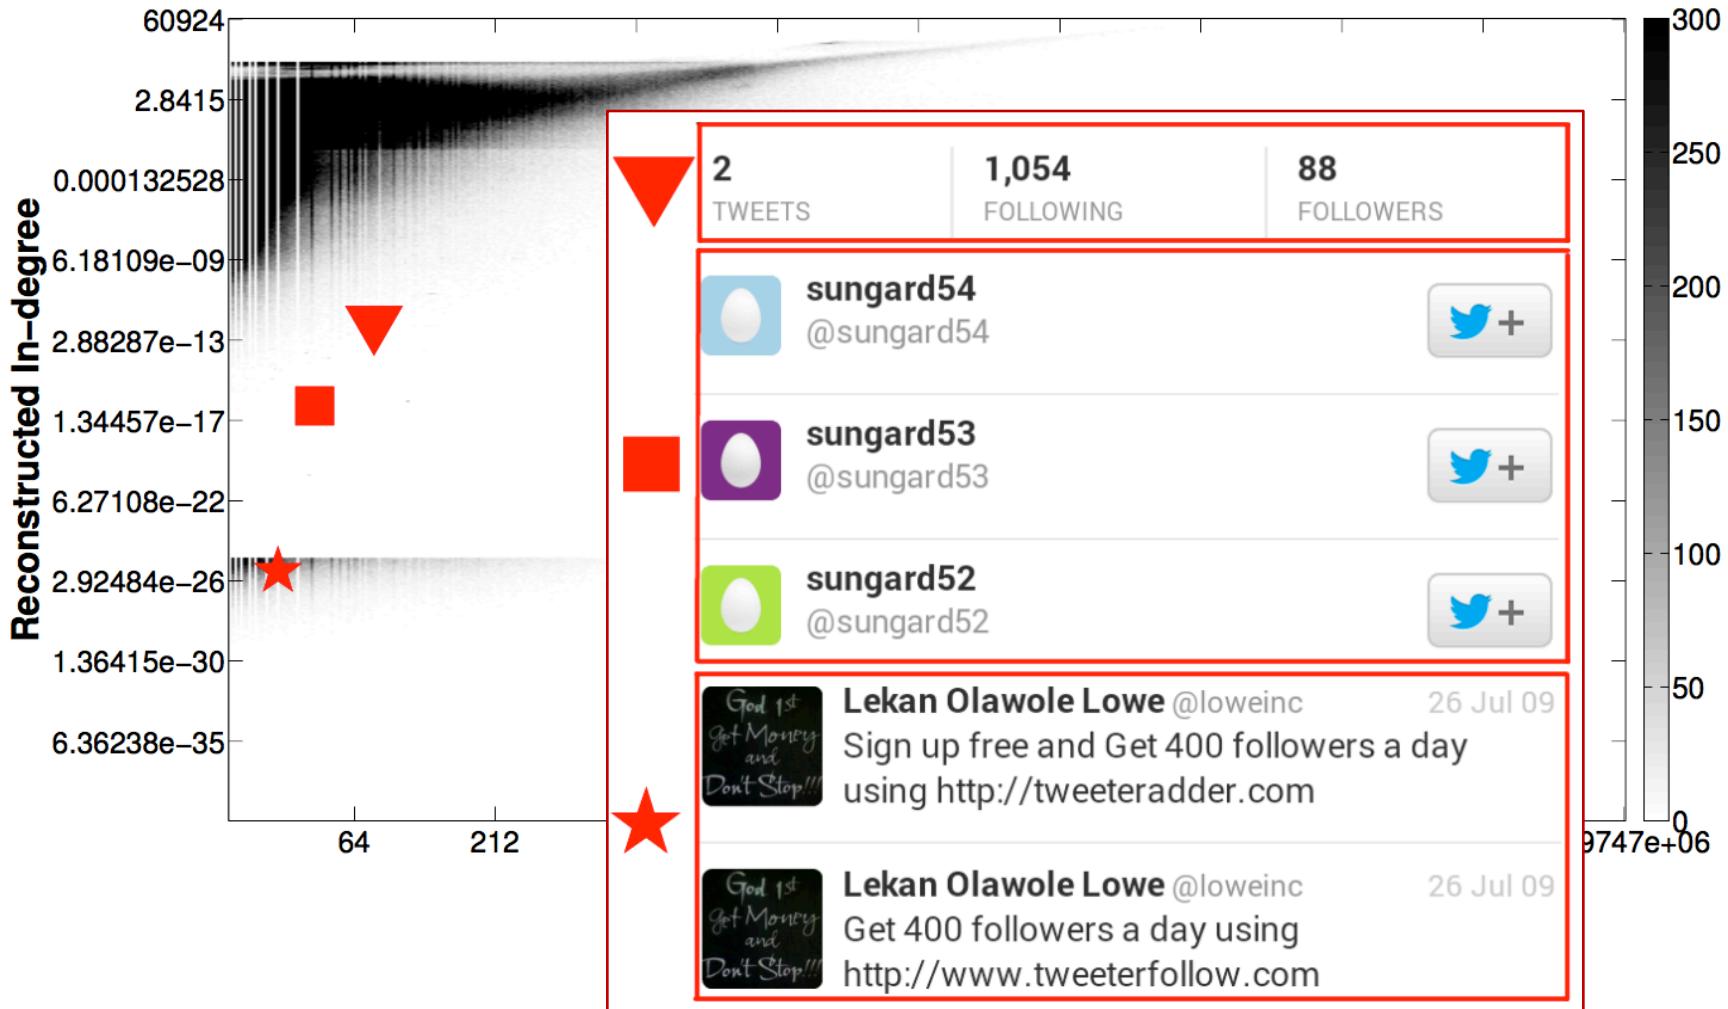
Reconstructed out-degree vs. Real out-degree



Reconstructed in-degree vs. Real in-degree



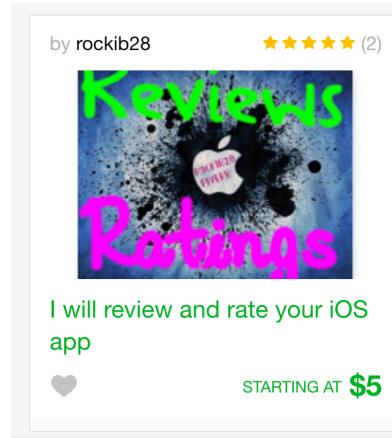
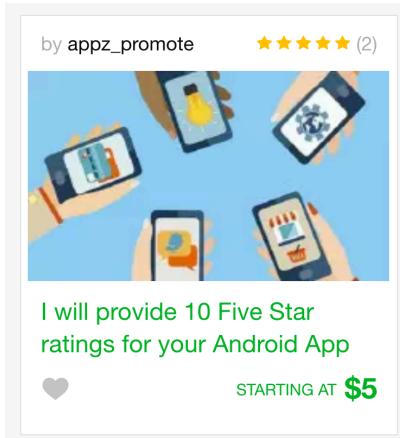
Reconstructed in-degree vs. Real in-degree (cont.)



Bounding Graph Fraud in Camouflage

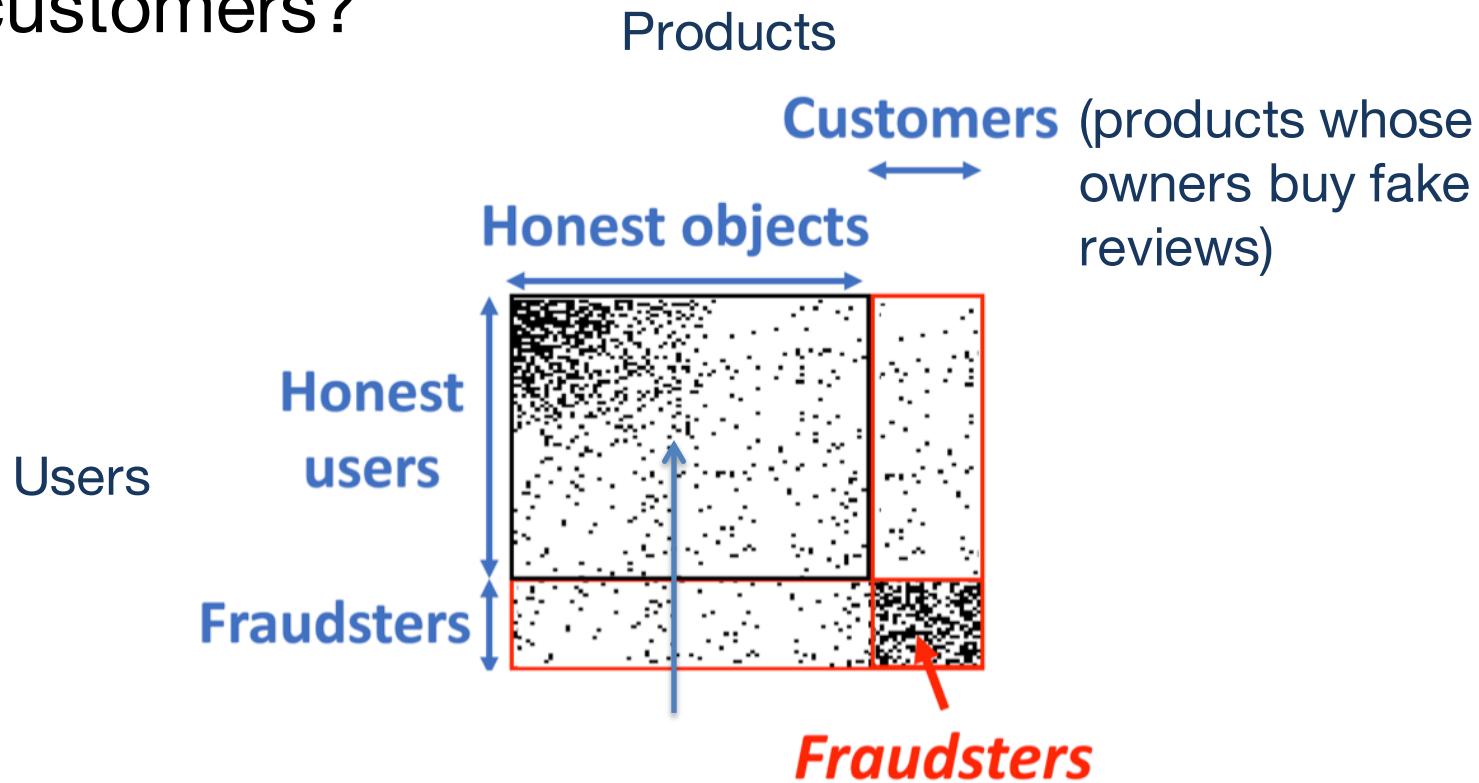
- An application: Fake reviews

I will do 5 five star reviews, all from real profiles



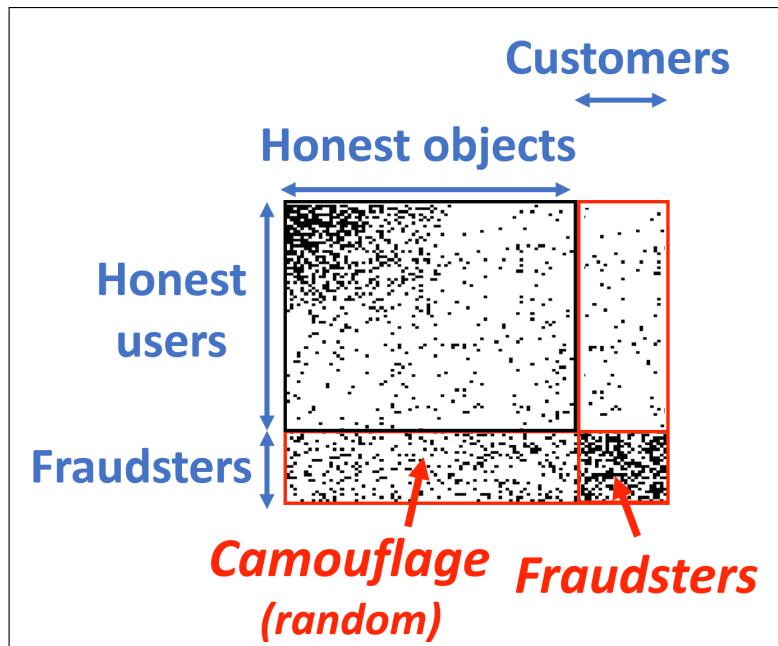
“User-Product” Review Graph

- Problem definition: Given a “user-product” review graph, can we spot fraudsters and customers?

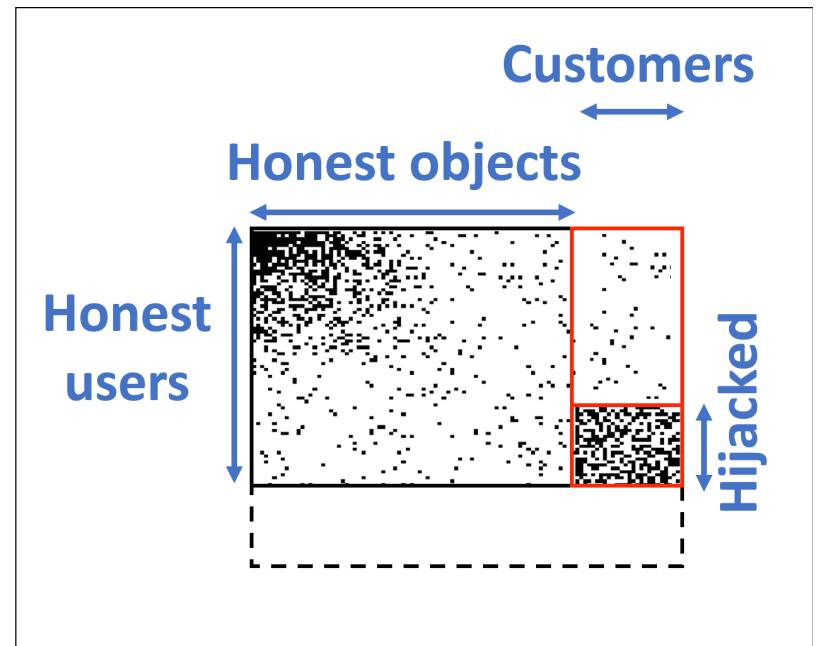


Camouflage: Evading Detection

Random camouflage



Hijacked user accounts

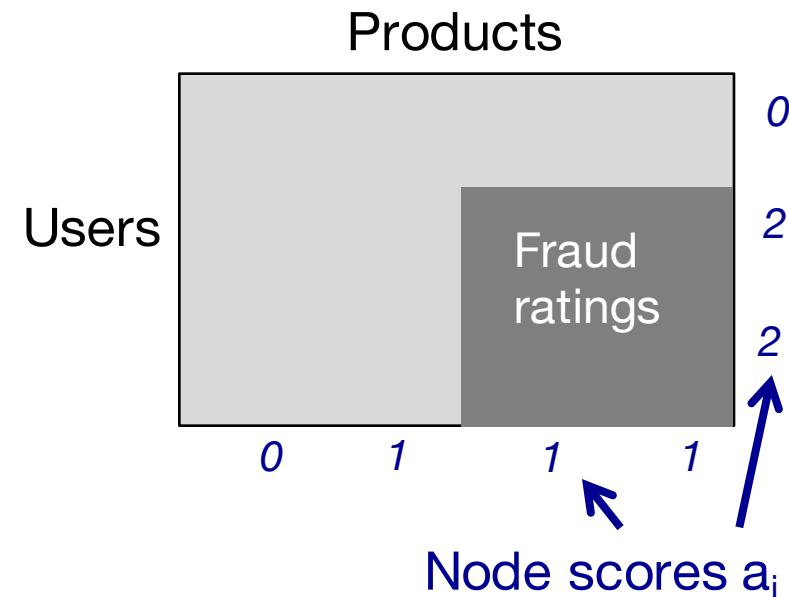


“camouflage” in LockInfer

“Fame” in LockInfer

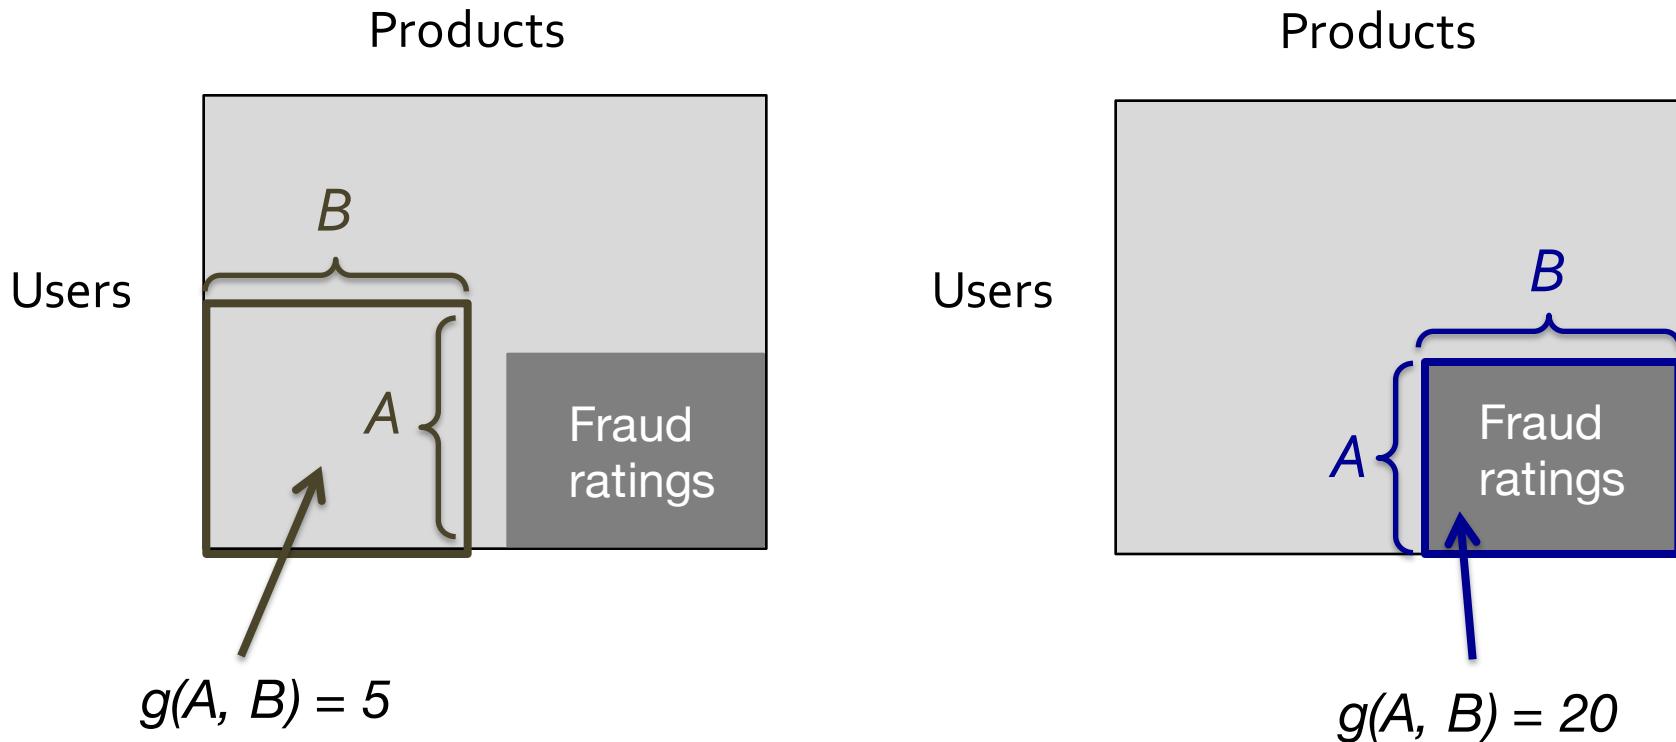
Formal Problem Definition

- Given:
 - Bipartite graph between users and products
 - May have prior node suspiciousness scores
- Develop detection metric that is:
 - **Camouflage-resistant**
 - **Near-linear time**
 - **Offers provable bounds**
 - **Works well in practice**



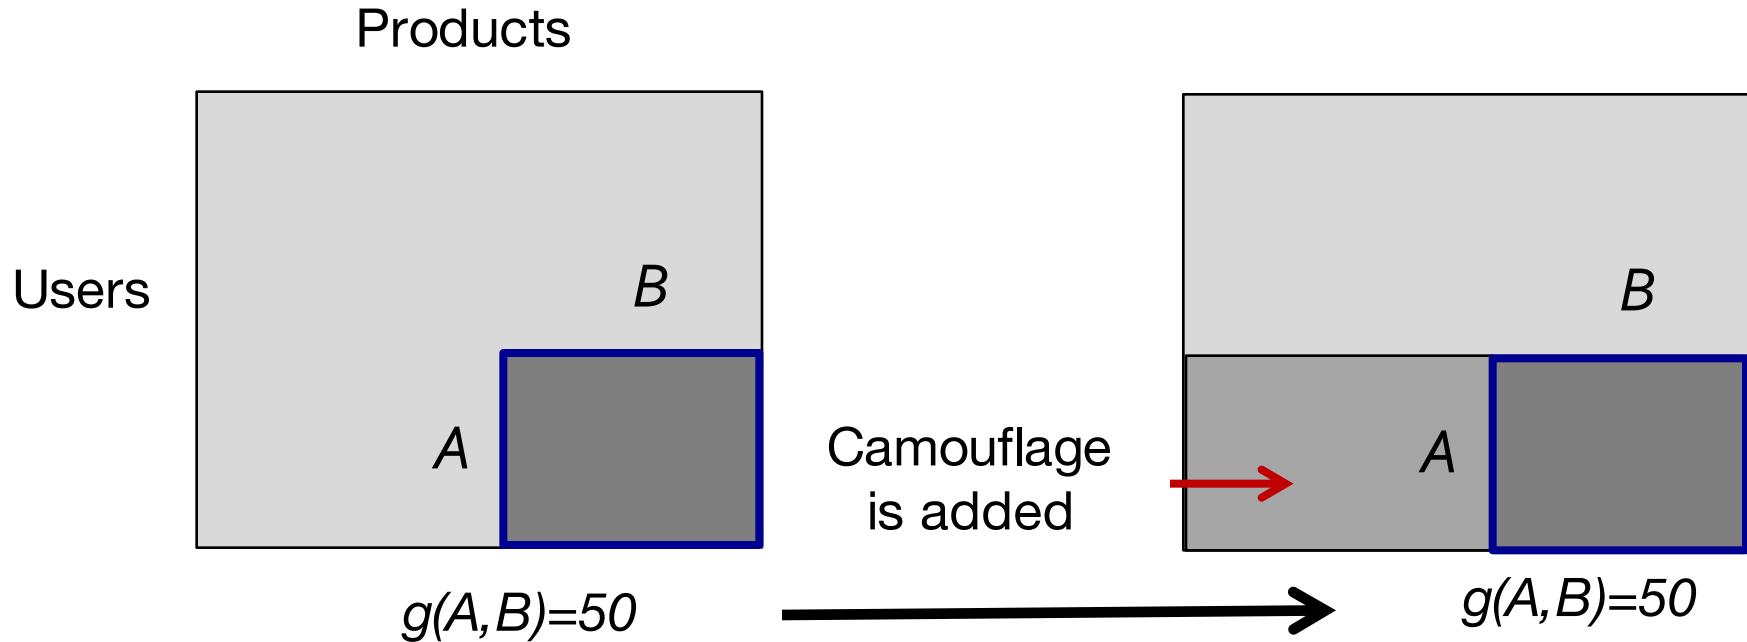
Suspiciousness Metric

$g(A, B)$ is a density metric for edges from set of users A to set of products B.



Camouflage-Resistance

Metric g is **camouflage-resistant** if $g(A, B)$ does not decrease when camouflage is added to A .

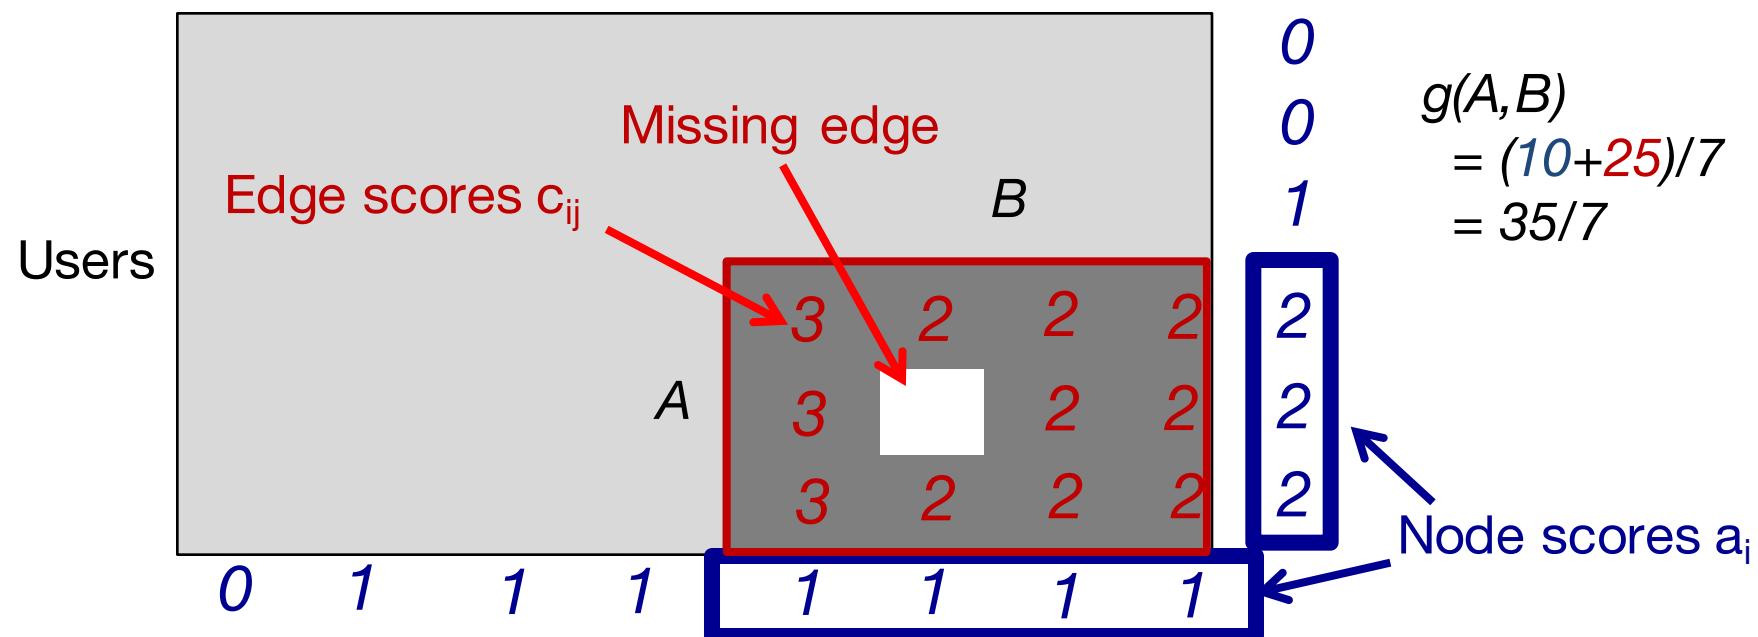


Proposed Suspiciousness Metric

“Average suspiciousness” $g(A, B)$

$$= \frac{(\text{sum of node susp.}) + (\text{sum of edge susp.})}{|A| + |B|}$$

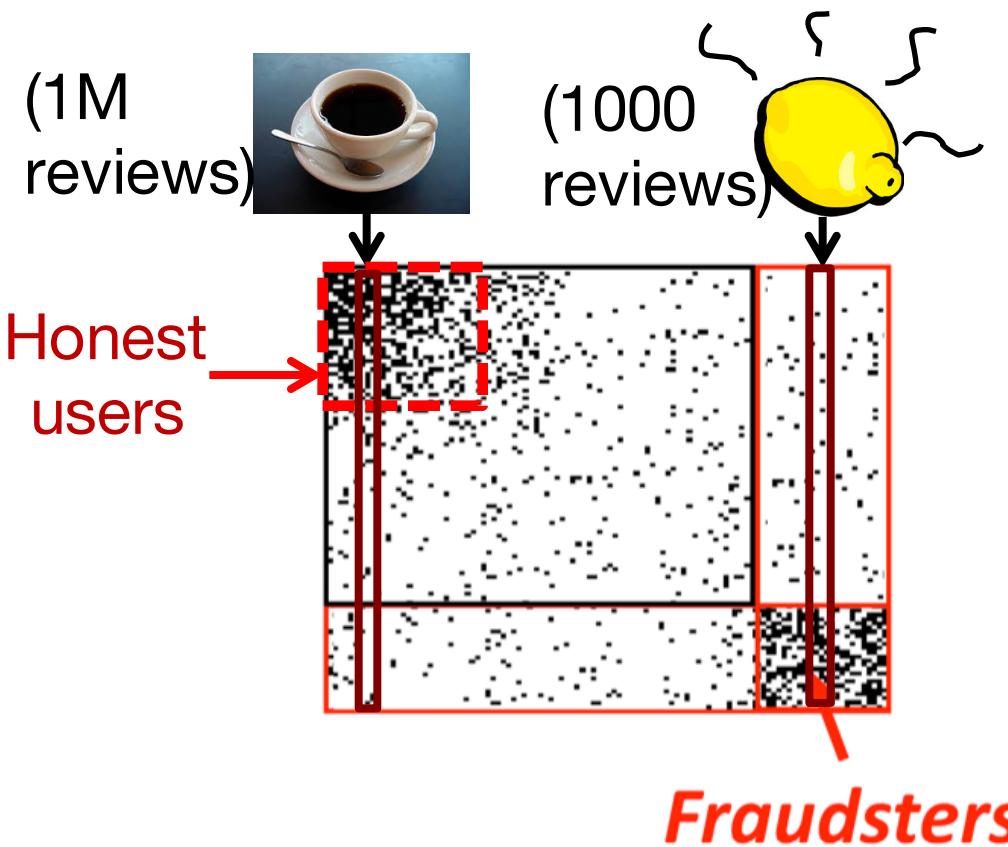
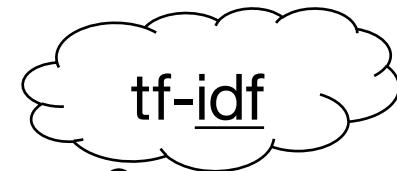
Products



Edge Scores c_{ij}

Proposed weighting scheme:

$$c_{ij} = 1 / \log(\text{unweighted sum of } j\text{-th column})$$



Why?

- Popular products are not necessarily suspicious
- Fraudulent products have a high fraction of edges from fraudsters

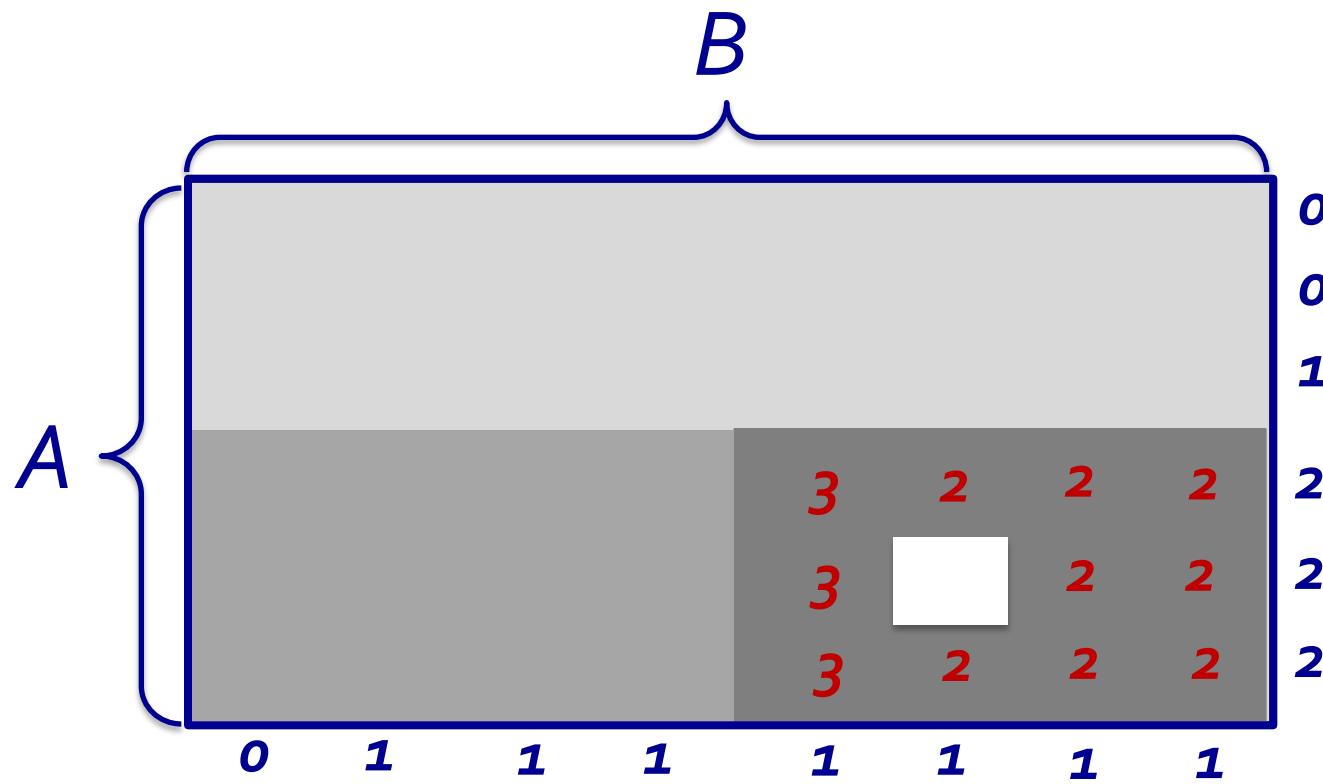
Metric Properties

Average suspiciousness $g(A, B)$:

- Can be optimized in near-linear time**
- Provable bounds
- Camouflage-resistant
- Works in practice

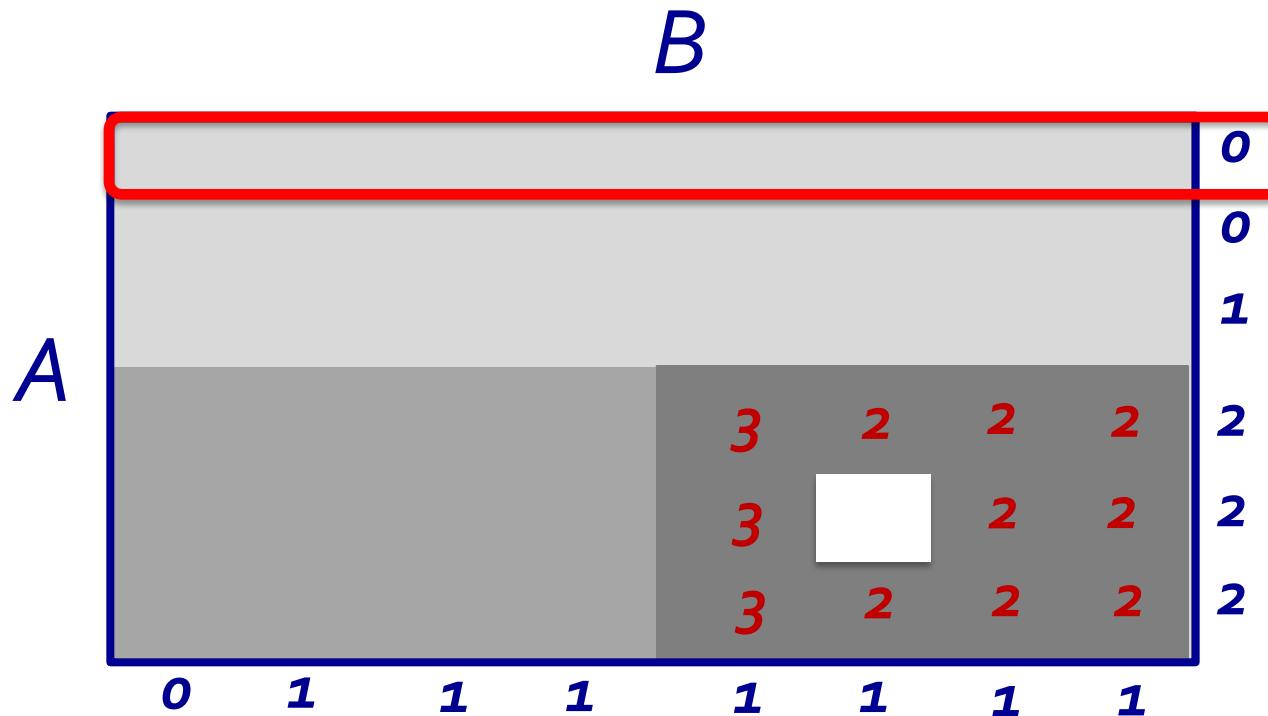
FRAUDAR: Greedy Algorithm

- Start with sets A, B as all users / products



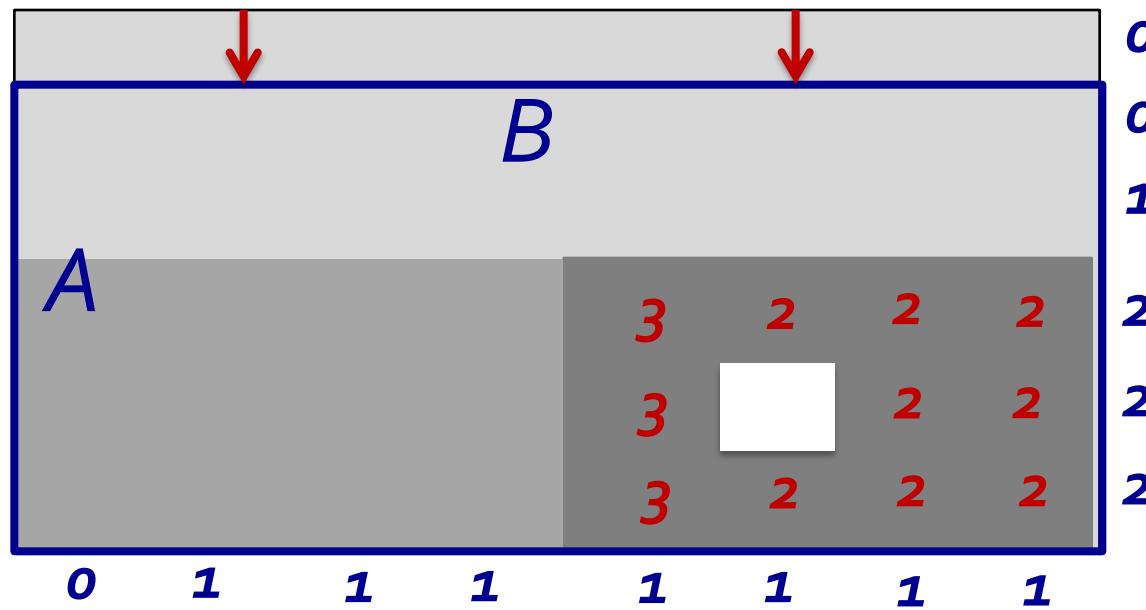
FRAUDAR: Greedy Algorithm (cont.)

- Delete rows / columns greedily to maximize g (average suspiciousness)



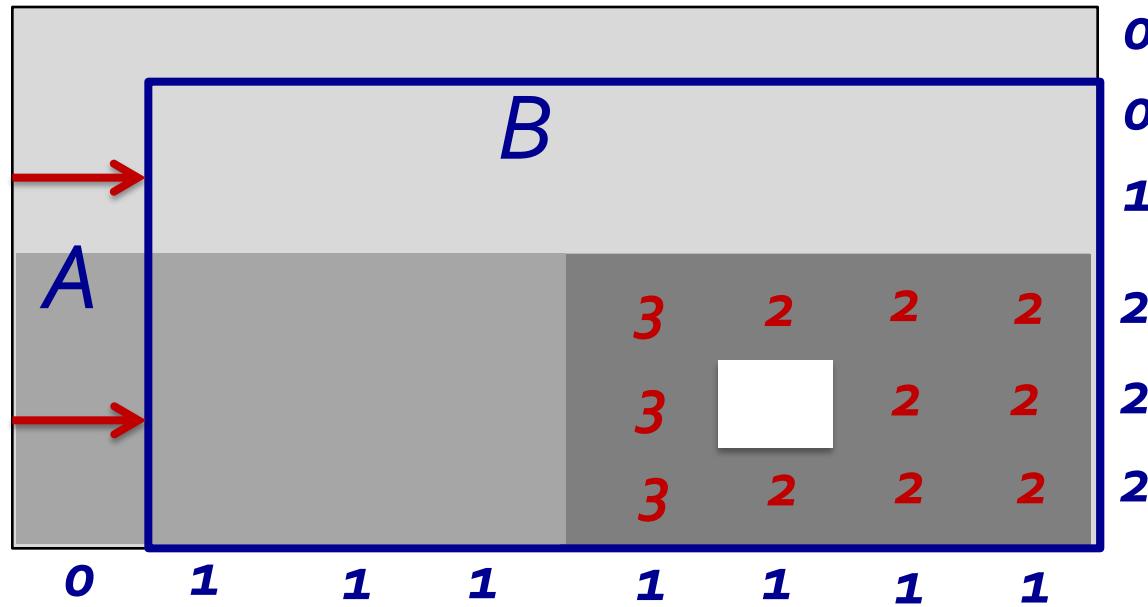
FRAUDAR: Greedy Algorithm (cont.)

- Delete rows / columns greedily to maximize g (average suspiciousness)



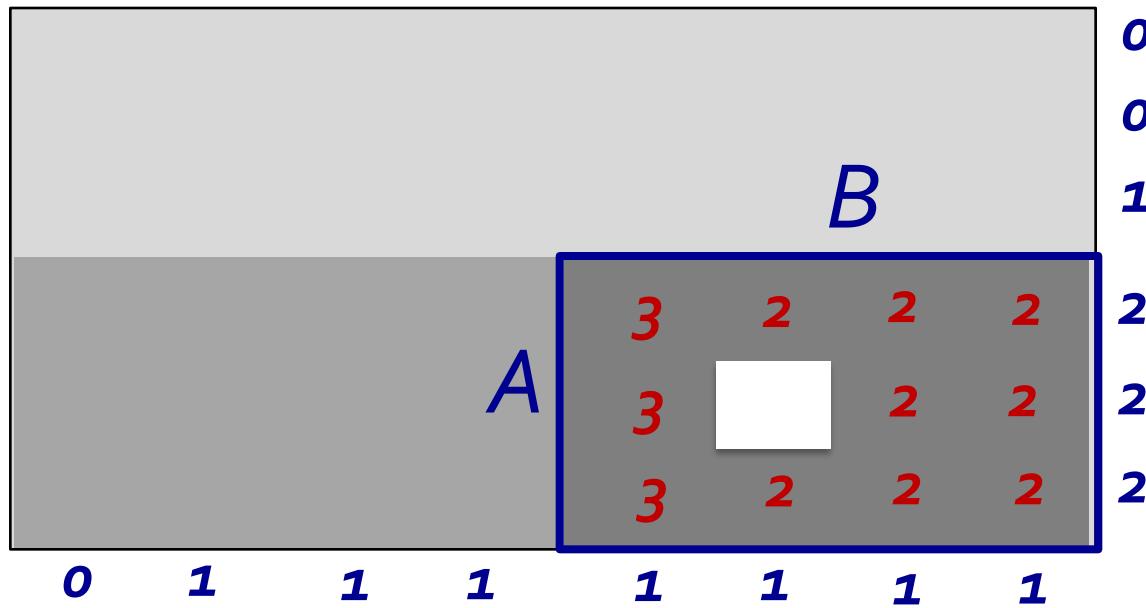
FRAUDAR: Greedy Algorithm (cont.)

- Delete rows / columns greedily to maximize g (average suspiciousness)



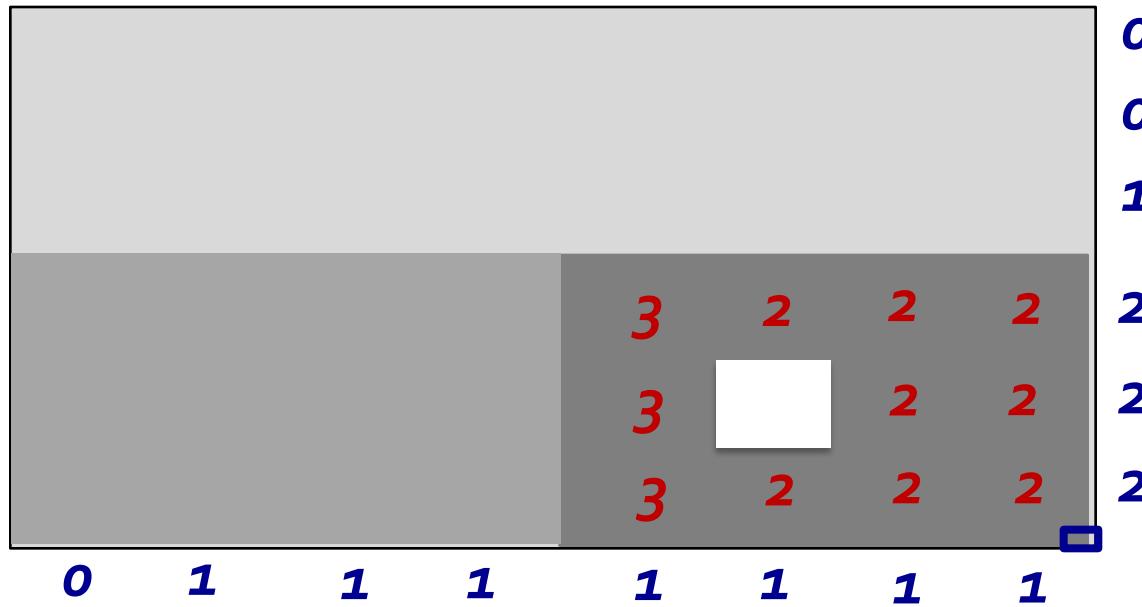
FRAUDAR: Greedy Algorithm (cont.)

- Delete rows / columns greedily to maximize g (average suspiciousness)



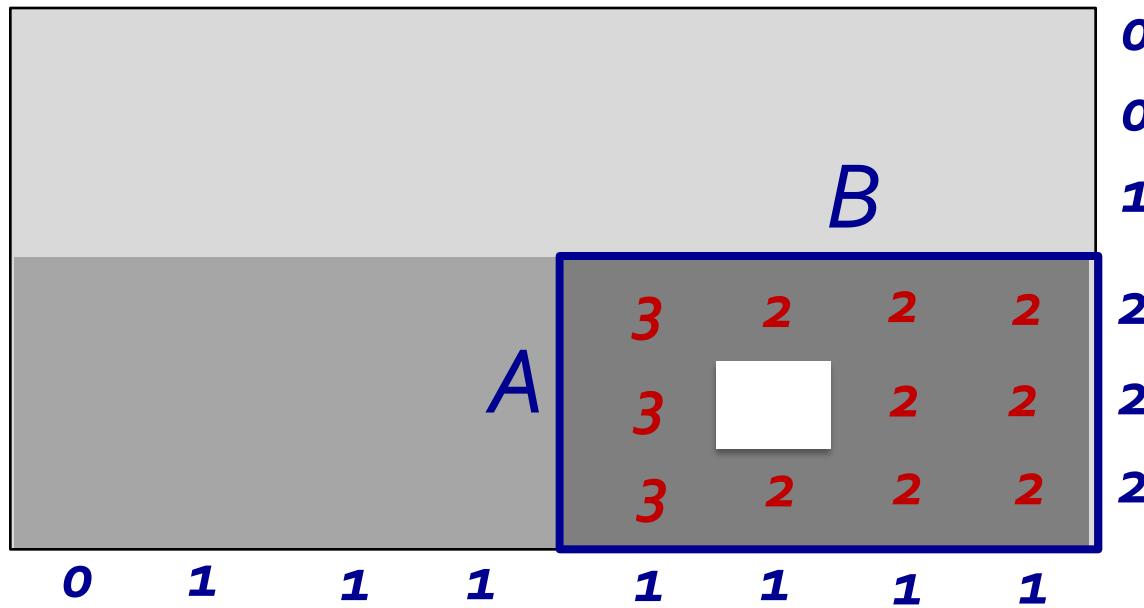
FRAUDAR: Greedy Algorithm (cont.)

- Continue until A and B are empty



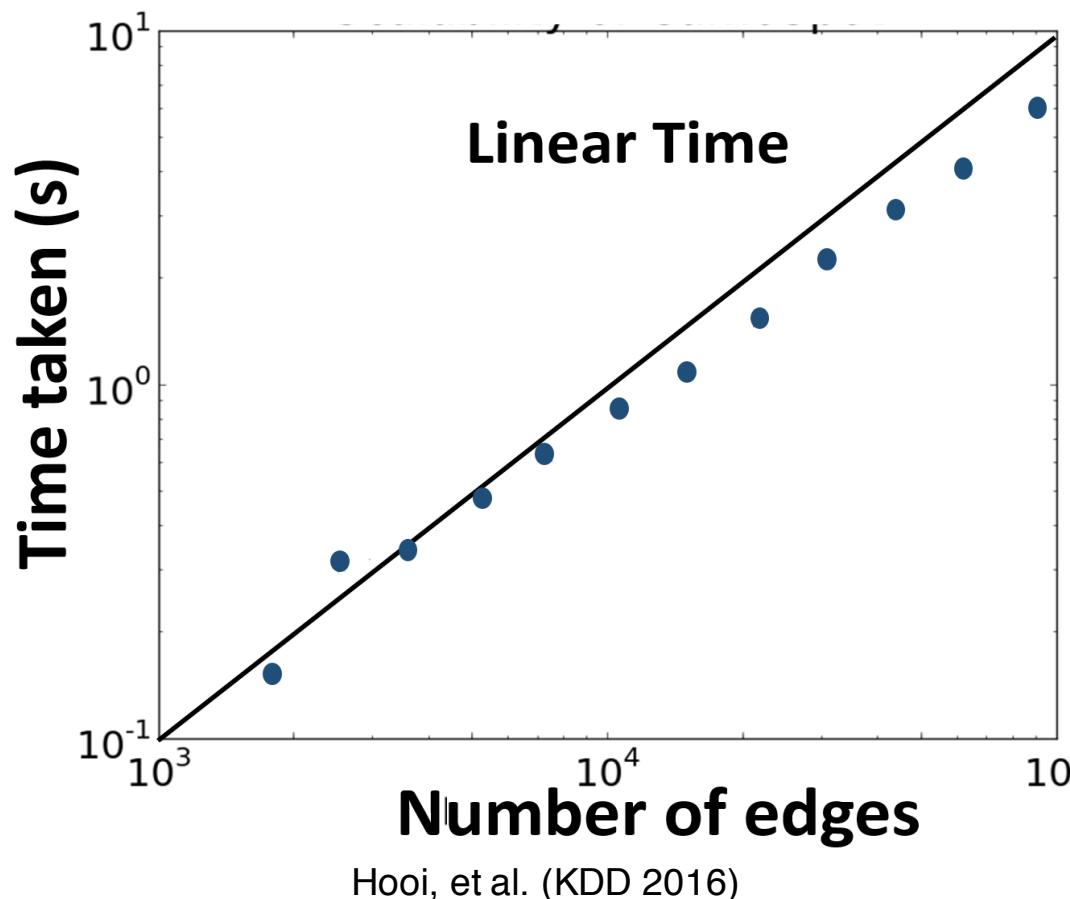
FRAUDAR: Greedy Algorithm (cont.)

- Return the best subsets A and B seen so far (based on g)



Computation Time

- $O(|E| \log(|V|))$: using appropriate data structures



Metric Properties

Average suspiciousness $g(A, B)$:

- Can be optimized in near-linear time
- Provable bounds**
- Camouflage-resistant
- Works in practice

Theoretical Guarantee

- **Theorem 1:** The subgraph (A, B) returned by FRAUDAR satisfies

$$g(\mathcal{A} \cup \mathcal{B}) \geq \frac{1}{2}g_{OPT}$$

FRAUDAR subgraph

Optimum value of g

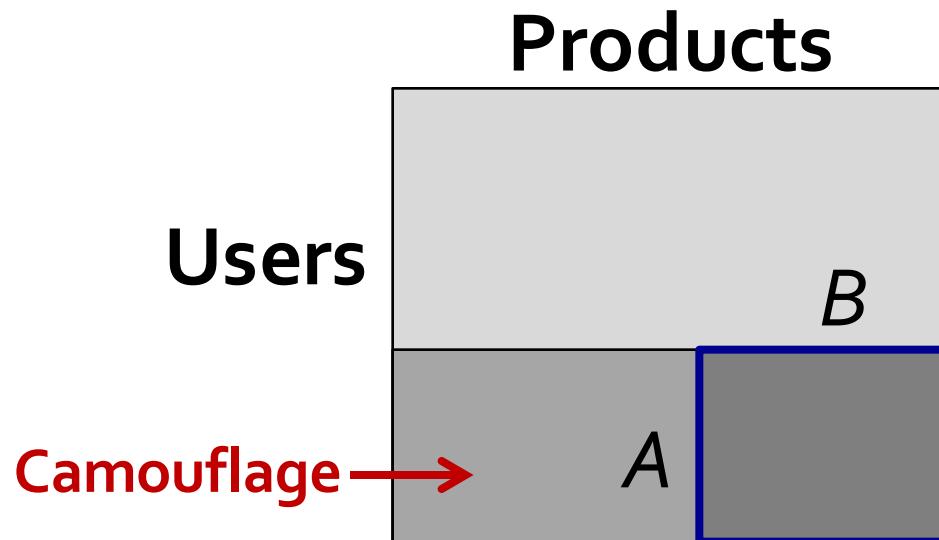
Metric Properties

Average suspiciousness $g(A, B)$:

- Can be optimized in near-linear time
- Provable bounds
- Camouflage-resistant**
- Works in practice

Camouflage Resistance

- **Theorem 2:** If c_{ij} is a *column weighting* (i.e. c_{ij} is any function of the j -th column), then **g is camouflage-resistant.**



$(c_{ij} = 1 / \log(\text{sum of } j\text{-th column})$ satisfies this)

Metric Properties

Average suspiciousness $g(A, B)$:

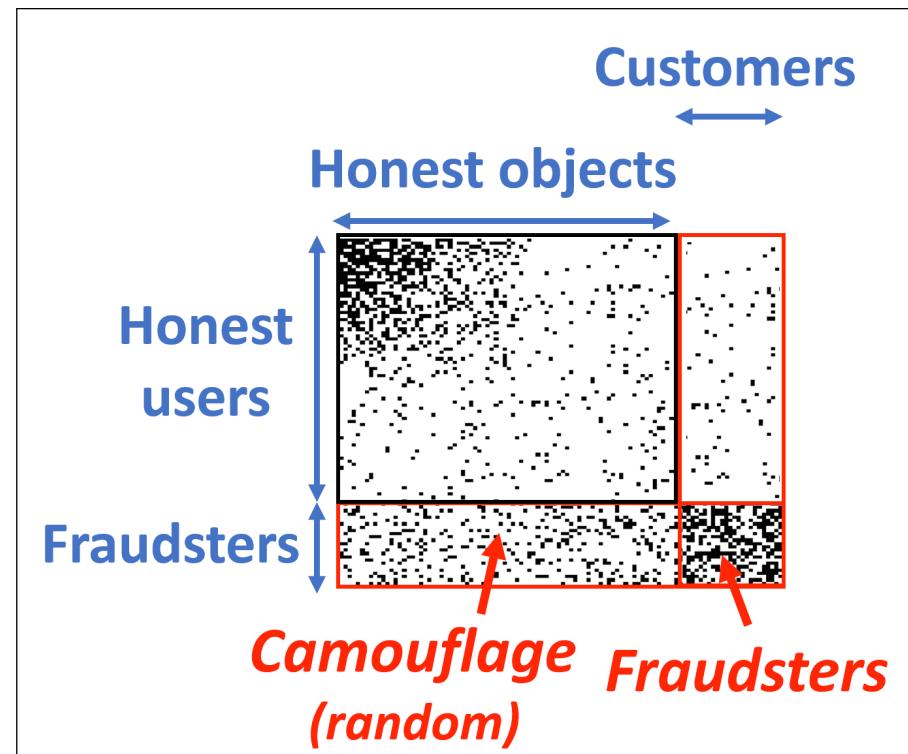
- Can be optimized in near-linear time
- Provable bounds
- Camouflage-resistant
- Works in practice**

Experiments: Detecting Injection of Various Types of “Camouflage”

- Amazon Review Graph:
24K users, 4K products
- Injected 200×200 blocks with various types of camouflage
 - None
 - Random camouflage
 - Biased camouflage
 - Hijacked accounts

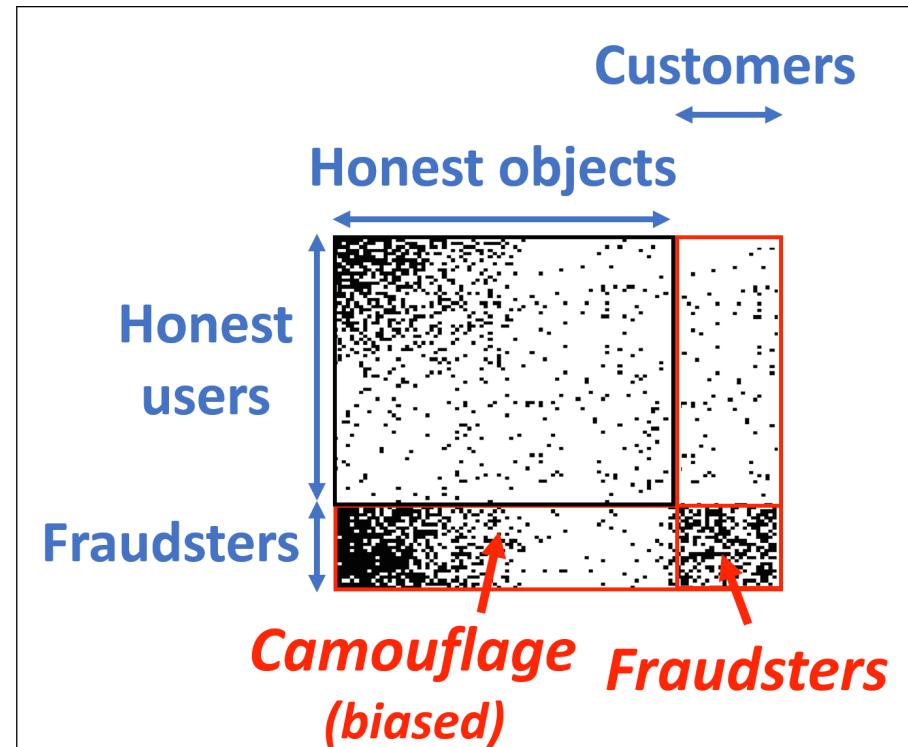
Experiments: Detecting Injection of Various Types of “Camouflage”

- Amazon Review Graph:
24K users, 4K products
- Injected 200×200 blocks with various types of camouflage
 - None
 - **Random camouflage**
 - Biased camouflage
 - Hijacked accounts



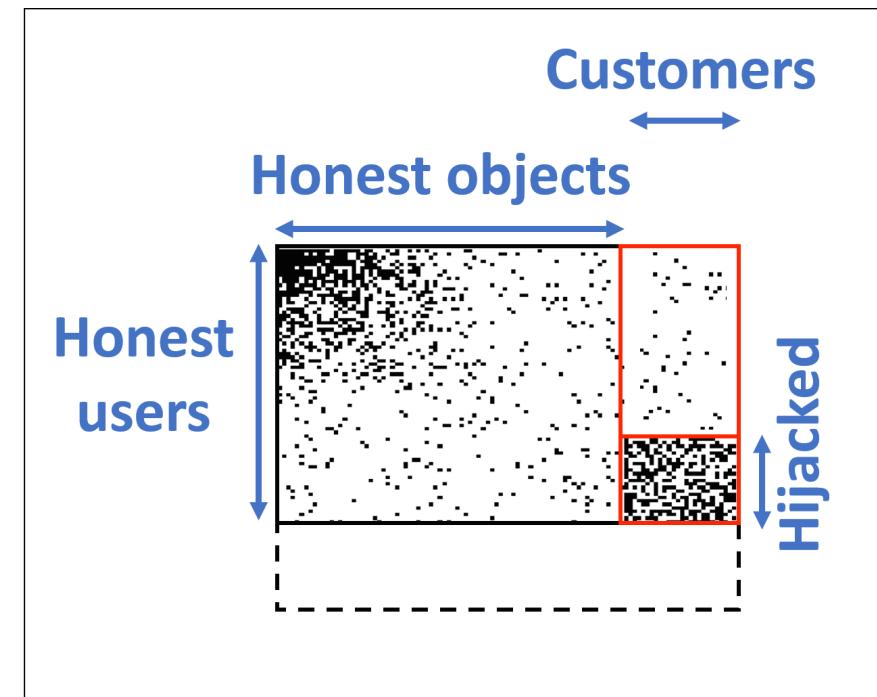
Experiments: Detecting Injection of Various Types of “Camouflage”

- Amazon Review Graph:
24K users, 4K products
- Injected 200×200 blocks with various types of camouflage
 - None
 - Random camouflage
 - **Biased camouflage**
 - Hijacked accounts

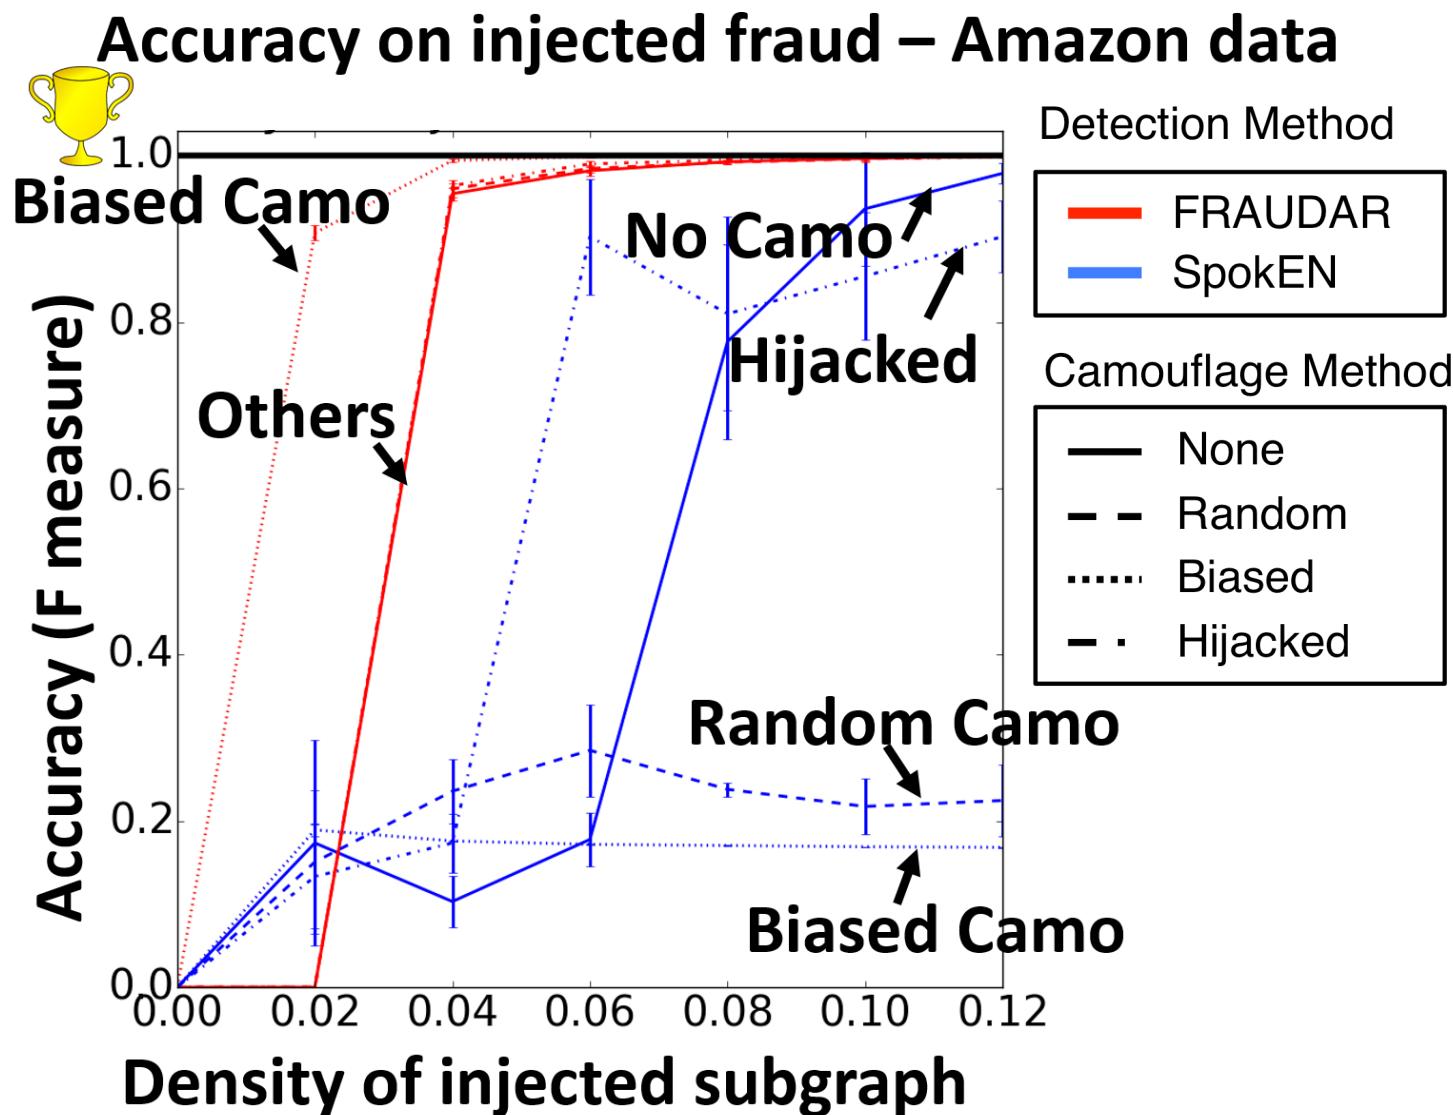


Experiments: Detecting Injection of Various Types of “Camouflage”

- Amazon Review Graph:
24K users, 4K products
- Injected 200×200 blocks with various types of camouflage
 - None
 - Random camouflage
 - Biased camouflage
 - **Hijacked accounts**

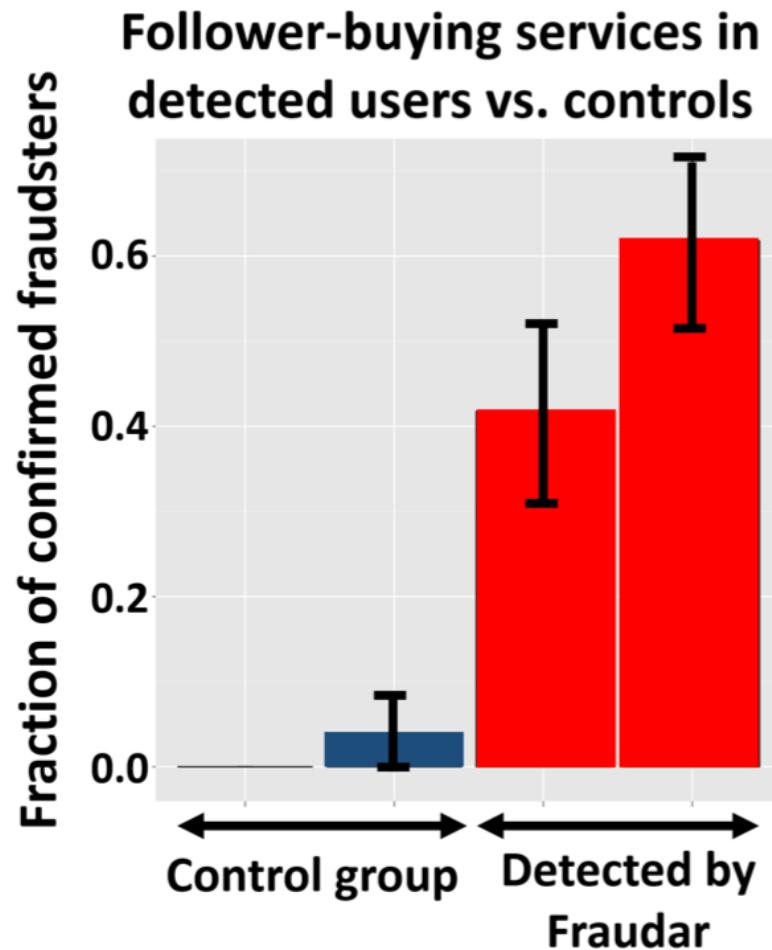


Accuracy on Detecting Injected Fraud



Accuracy on Detecting Fraud in Real Twitter Data

- Found 4031x4313 size block of followers-followees with 68% density
- Users detected as fraudulent by Fraudar are more likely to be deleted, suspended, use Twitter user buying services.



Summary

- Spectral methods
 - Spectral clustering and community detection
- Spectral subspaces and spectral subspace plots
- EigenSpokes (singular vectors and “spokes”)
- LockInfer (“**camouflage**”, “fame”, “pearls”, “staircase”, etc.)
- fBox (small-scale, stealthy attacks; reconstructed degrees)
- FRAUDAR (theoretical guarantees for bounding graph fraud in the face of **camouflage**)
- Applications: Mobile calls, Twitter social network, “user-product” reviews

References

- Andrew Y. Ng, Michael I. Jordan, Yair Weiss. “On Spectral Clustering: Analysis and an algorithm”, NIPS 2001.
- Scott White and Padhraic Smyth. “A spectral clustering approach to finding communities in graphs”, SDM 2005.
- M. E. J. Newman. “Finding community structure in networks using the eigenvectors of matrices”, Physical Review E 2006.
- Aditya Prakash, Mukund Seshadri, Ashwin Sridharan, Sridhar Machiraju, Christos Faloutsos. “EigenSpokes: Surprising Patterns and Scalable Community Chipping in Large Graphs”, PAKDD 2010.
- Meng Jiang, Peng Cui, Alex Beutel, Christos Faloutsos, Shiqiang Yang. “Inferring lockstep behavior from connectivity pattern in large graphs”, KAIS 2016.
- Neil Shah, Alex Beutel, Brian Gallagher, Christos Faloutsos. “Spotting suspicious link behavior with fBox: An adversarial perspective”, ICDM 2014.
- Bryan Hooi, Hyun Ah Song, Alex Beutel, Neil Shah, Kijung Shin, Christos Faloutsos. “FRAUDAR: Bounding Graph Fraud in the Face of Camouflage”, KDD 2016 Best Research Paper Award.



KDD 2017

Halifax, Nova Scotia - Canada
August 13 - 17, 2017

Tutorial: Data-Driven Approaches towards Malicious Behavior Modeling



Meng Jiang
University of Notre Dame



Srijan Kumar
Stanford University



Christos Faloutsos
Carnegie Mellon University



V.S. Subrahmanian
University of Maryland, College Park

Tutorial link: <http://bit.ly/kdd2017>